

# Deep Learning: Variational Autoencoders

**Ozan Özdenizci**

Institute of Theoretical Computer Science

[ozan.ozdenizci@igi.tugraz.at](mailto:ozan.ozdenizci@igi.tugraz.at)

Deep Learning VO - WS 23/24

Lecture 9 - December 4th, 2023

# Today

---

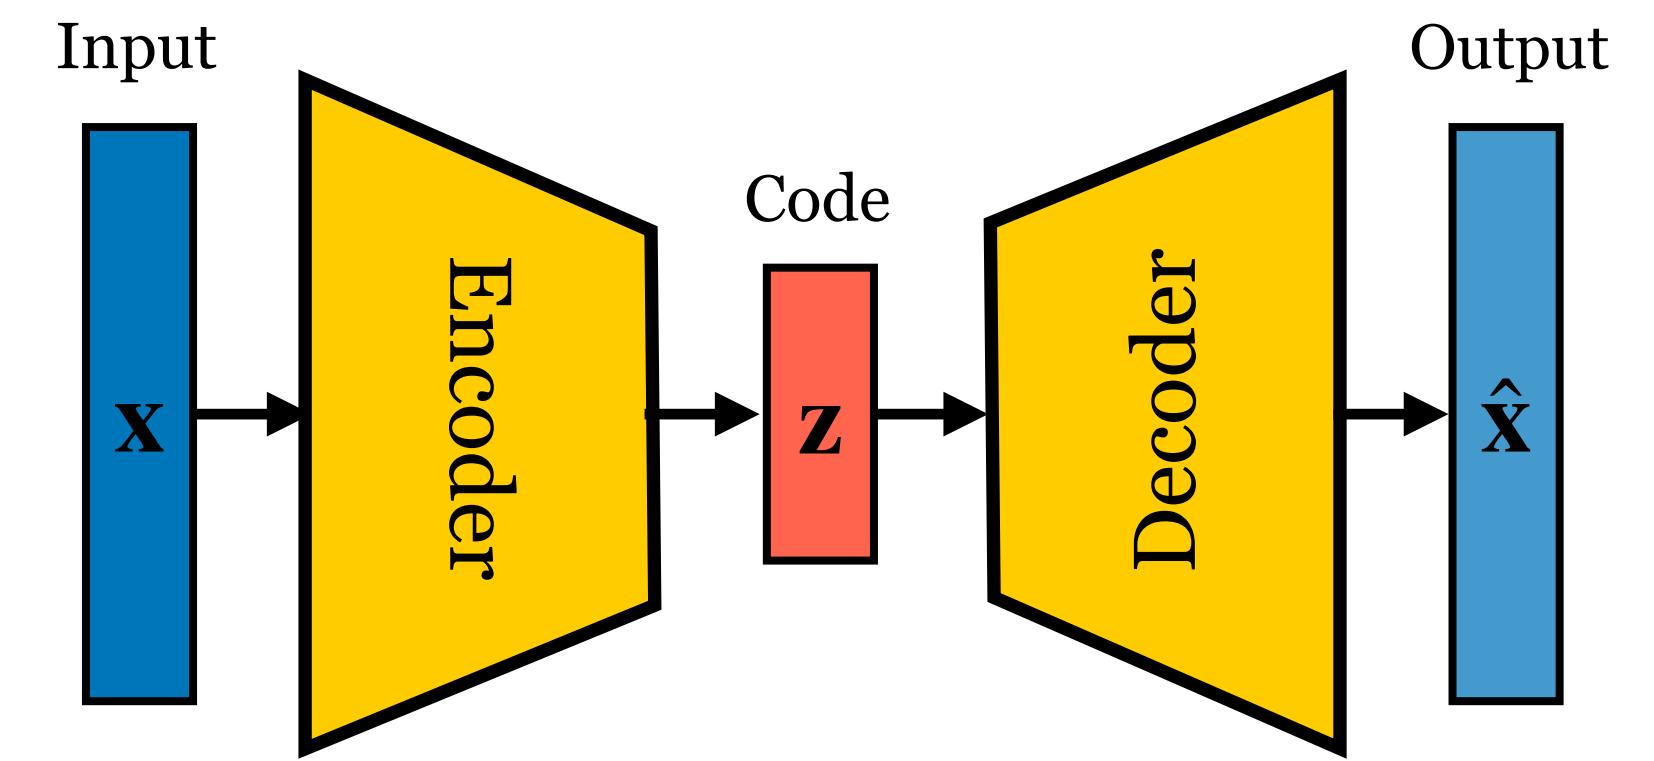
- ❑ Autoencoders
- ❑ Deep Generative Models
  - ❑ Variational Autoencoders
    - ❑ Architecture & Training
    - ❑ Variational Lower Bound
    - ❑ Experiments & Quantifying Performance

# Autoencoders

An **autoencoder** is a neural network that gets as input some  $\mathbf{x} \in \mathbb{R}^D$  and outputs a reconstruction  $\hat{\mathbf{x}} \in \mathbb{R}^D$ .

The neural network consists of an **encoder**  $f_{\text{enc}} : \mathbb{R}^D \rightarrow \mathbb{R}^{d_z}$ , and a **decoder**  $f_{\text{dec}} : \mathbb{R}^{d_z} \rightarrow \mathbb{R}^D$ , with a bottleneck layer with activations  $\mathbf{z} \in \mathbb{R}^{d_z}$  where the information about the input is compressed.

The network is trained with gradient descent to minimize the reconstruction error (e.g., mean-squared error).



# Autoencoders

An **autoencoder** is a neural network that gets as input some  $\mathbf{x} \in \mathbb{R}^D$  and outputs a reconstruction  $\hat{\mathbf{x}} \in \mathbb{R}^D$ .

The neural network consists of an **encoder**  $f_{\text{enc}} : \mathbb{R}^D \rightarrow \mathbb{R}^{d_z}$ , and a **decoder**  $f_{\text{dec}} : \mathbb{R}^{d_z} \rightarrow \mathbb{R}^D$ , with a bottleneck layer with activations  $\mathbf{z} \in \mathbb{R}^{d_z}$  where the information about the input is compressed.

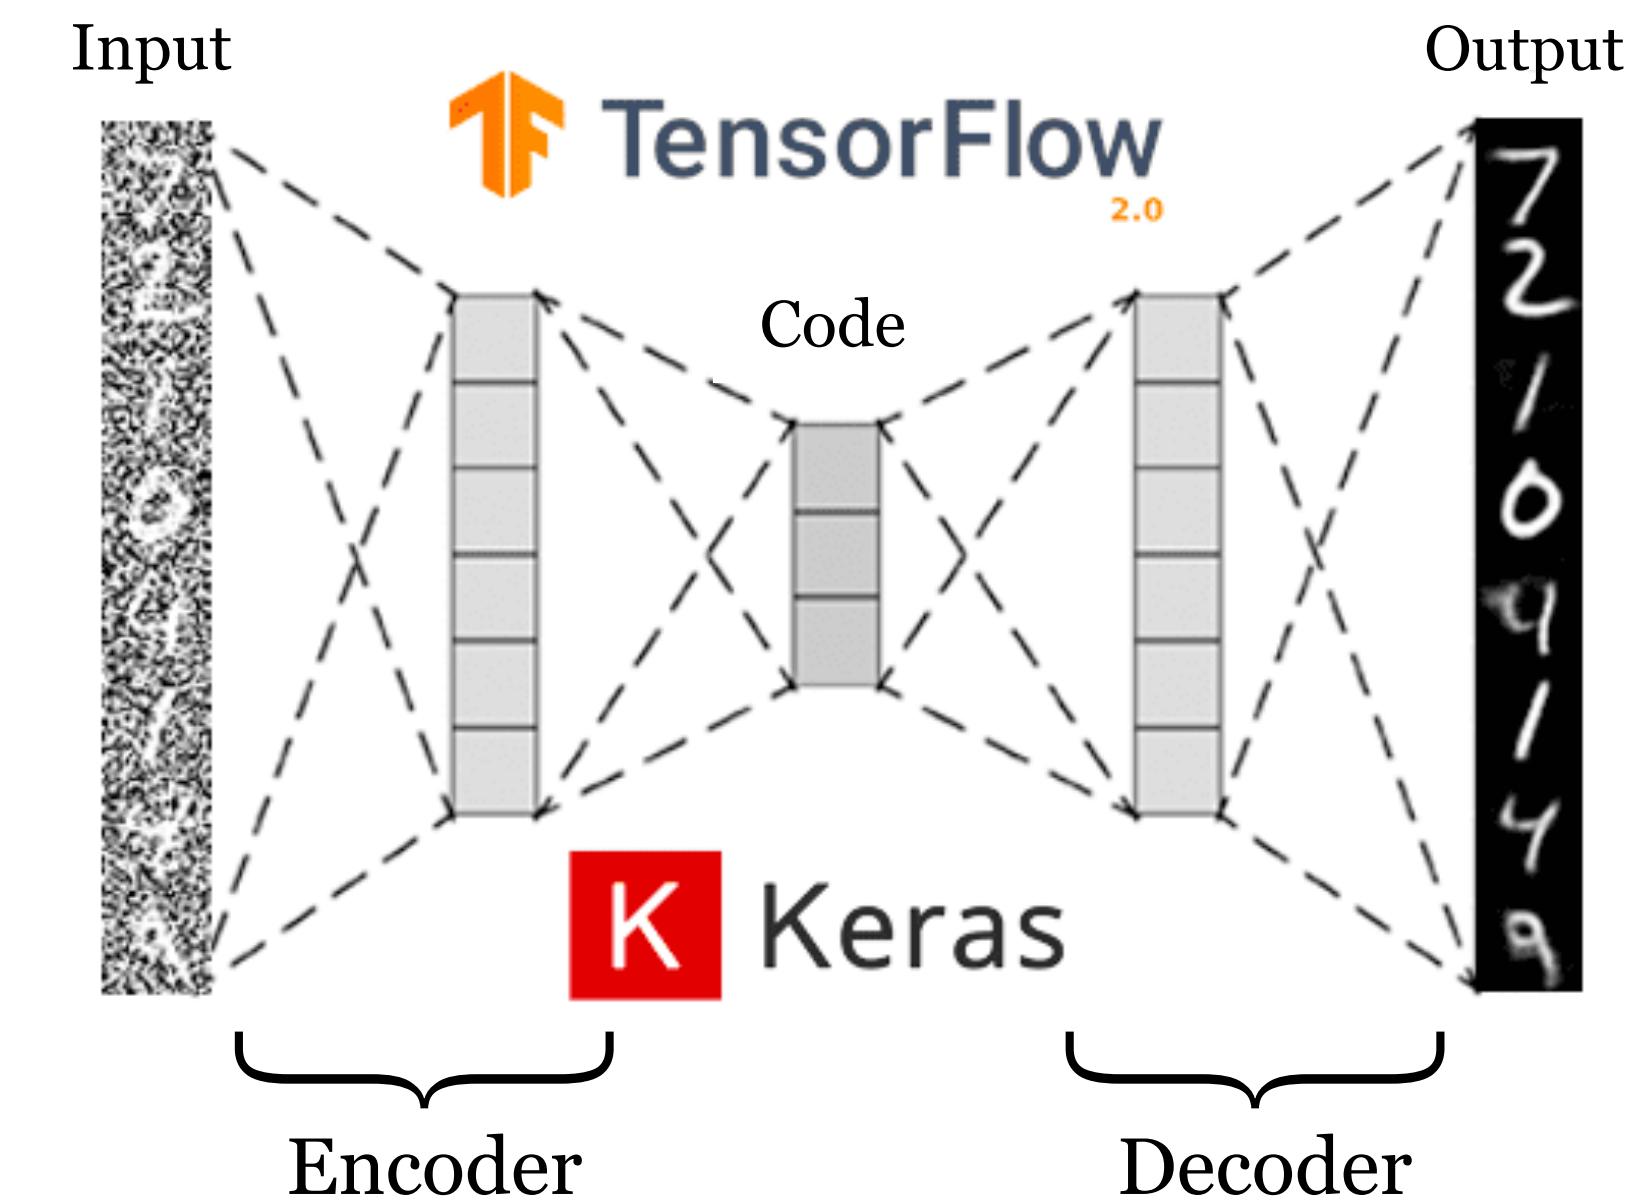
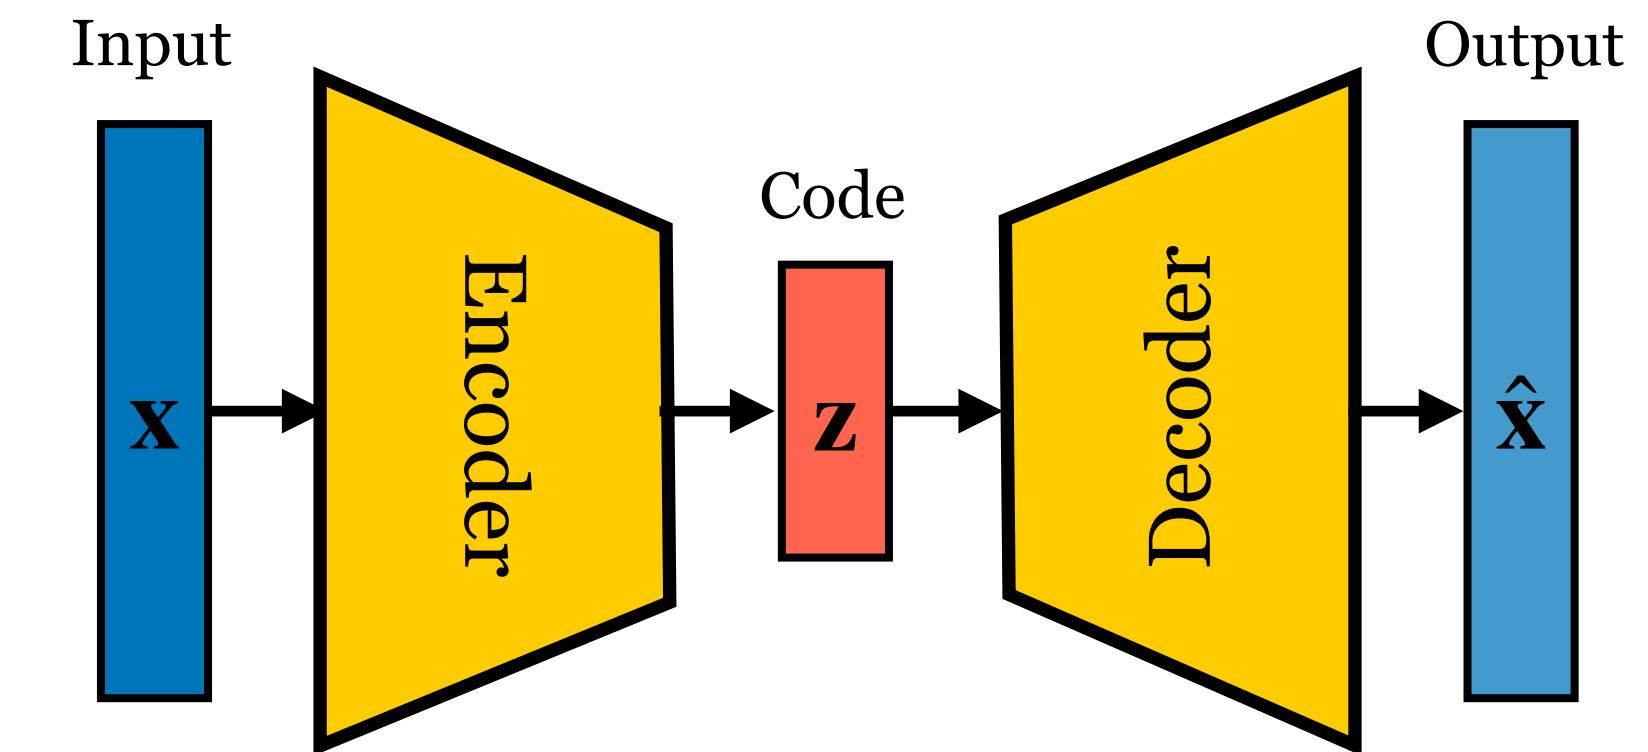
The network is trained with gradient descent to minimize the reconstruction error (e.g., mean-squared error).

## Applications:

- Dimensionality reduction
- Feature extraction
- Denoising autoencoders

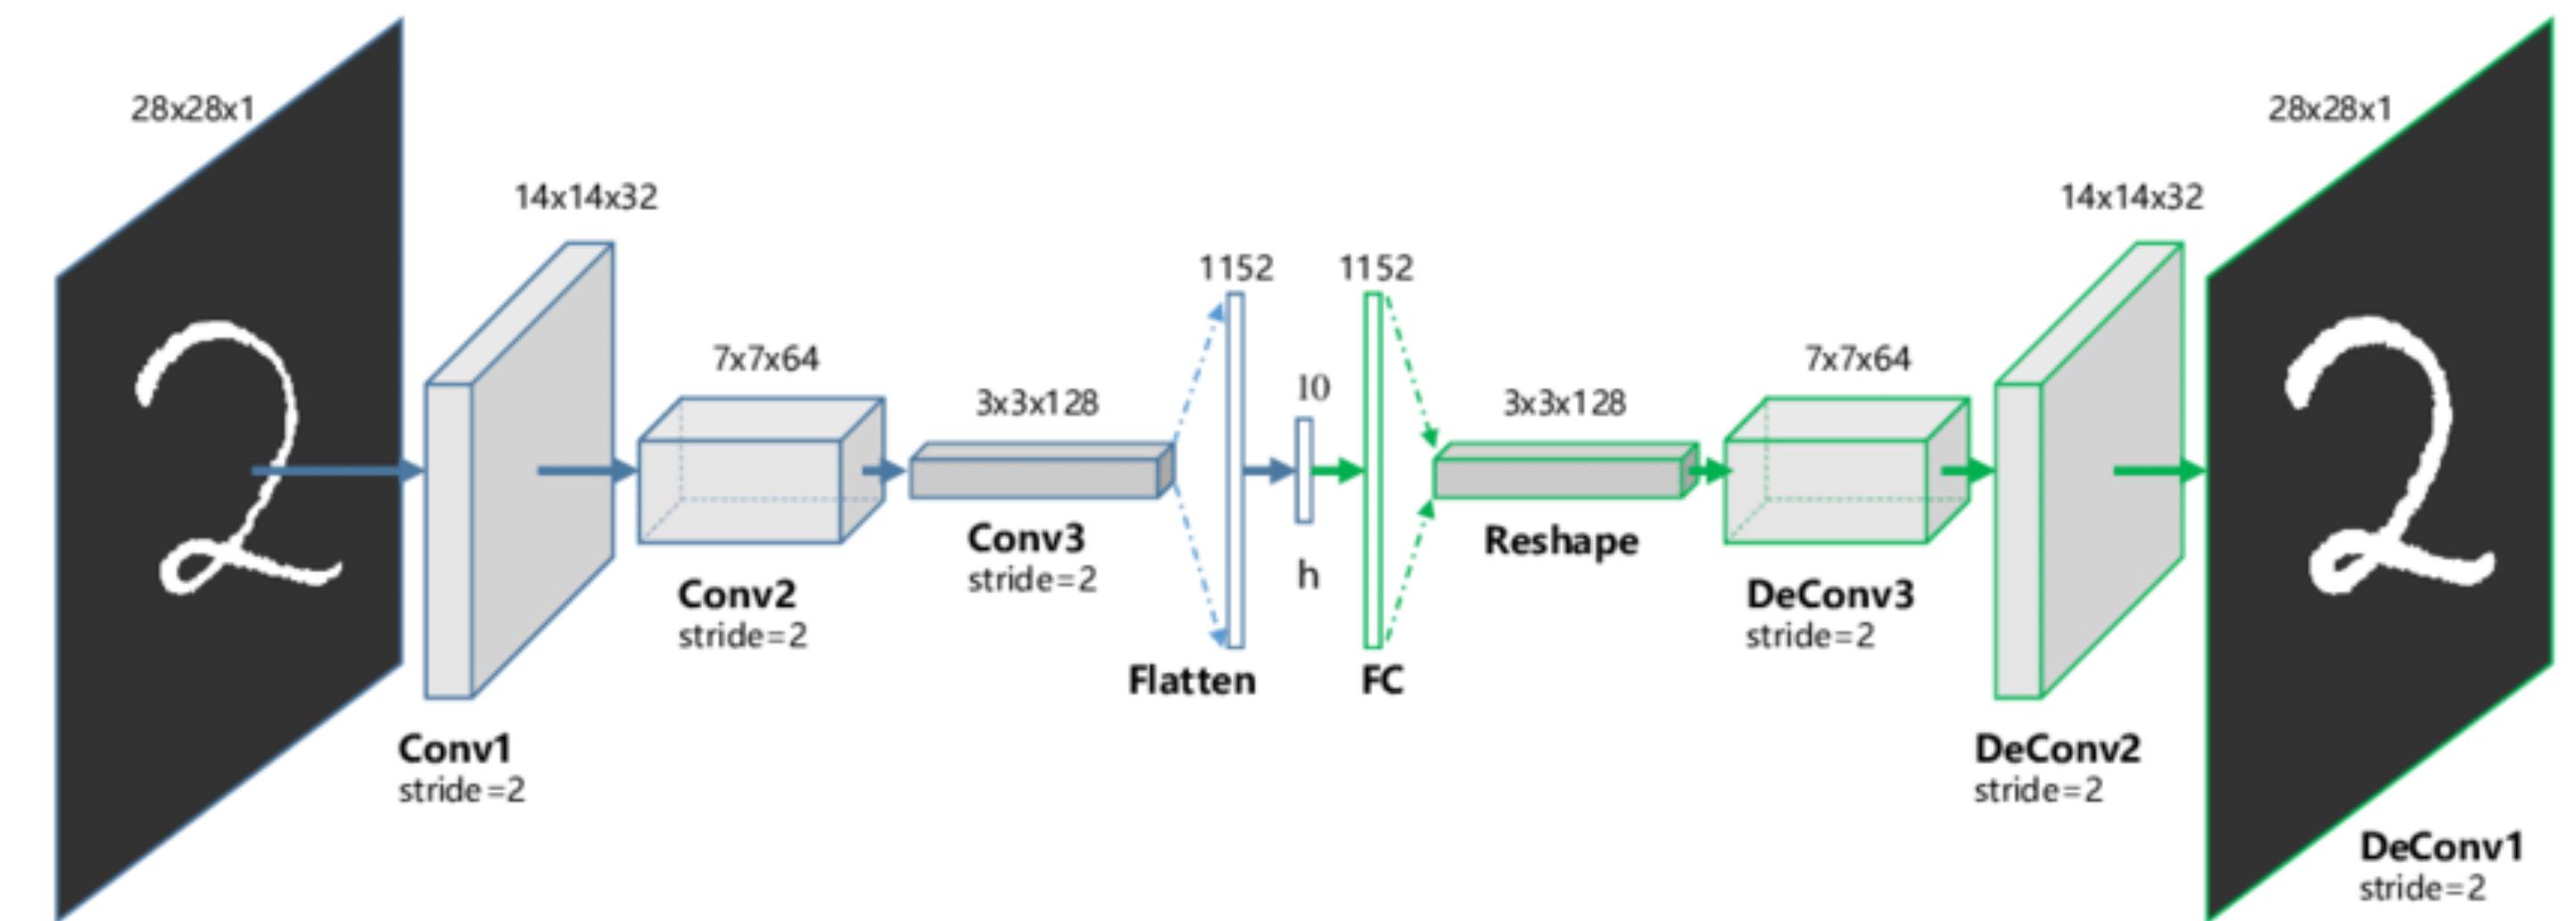
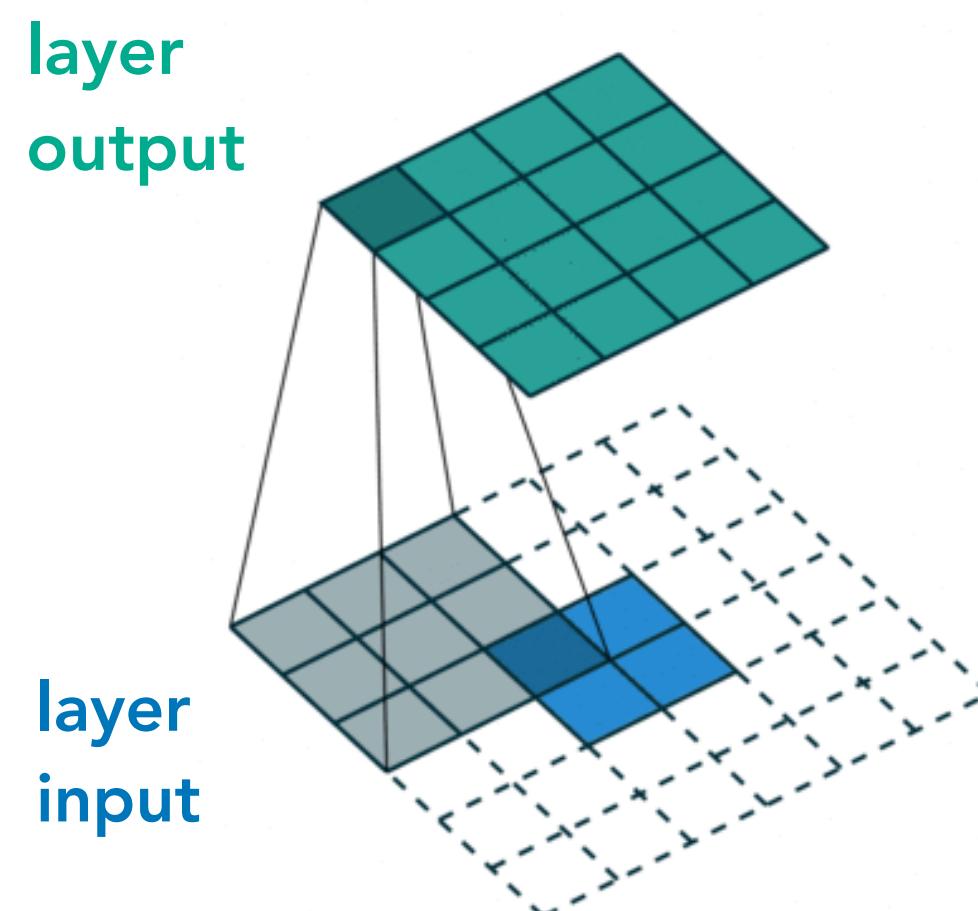
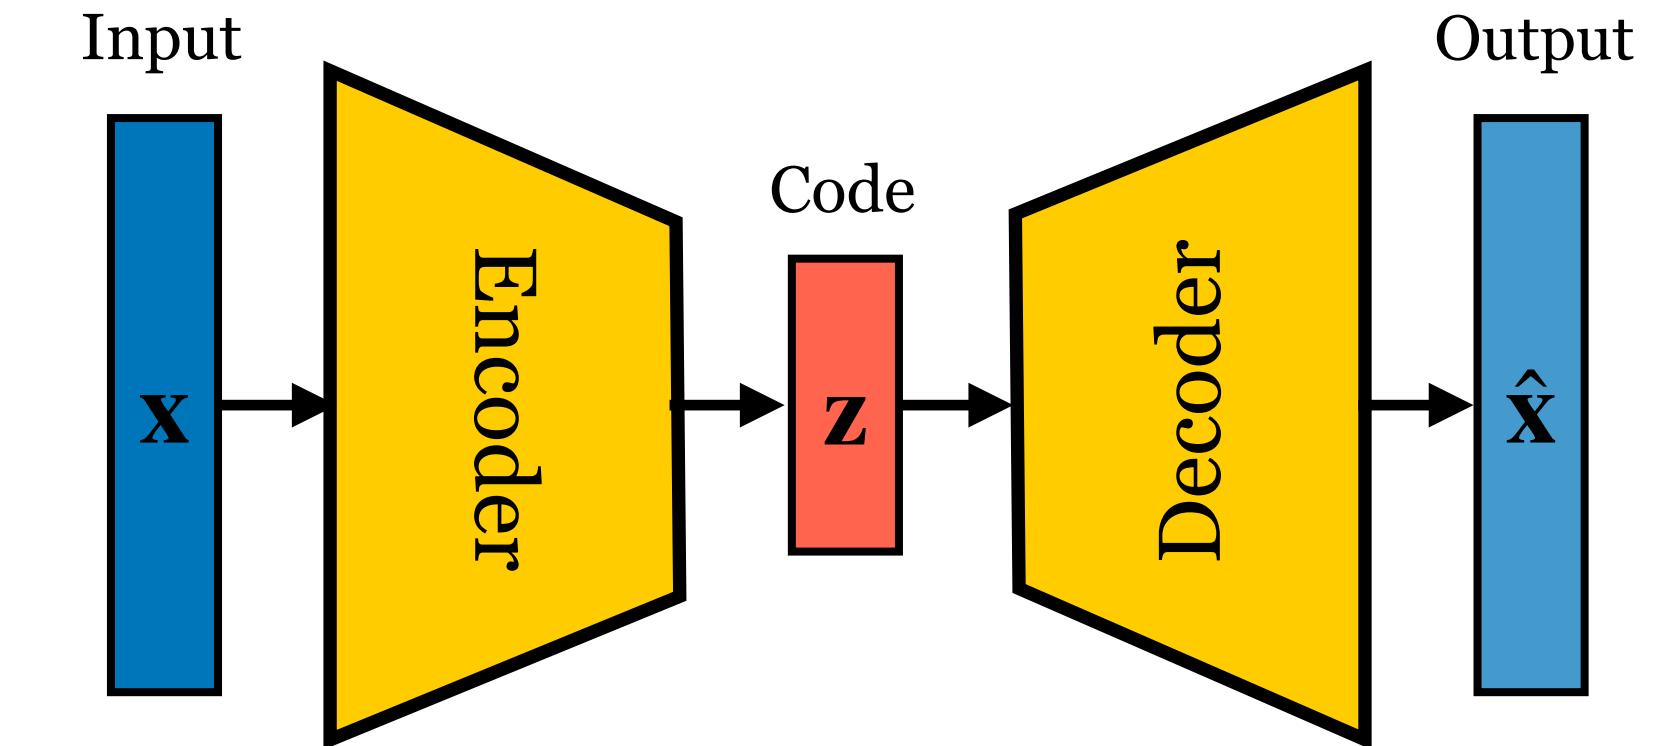


Training an autoencoder to reconstruct the original data from corrupted/noisy versions.



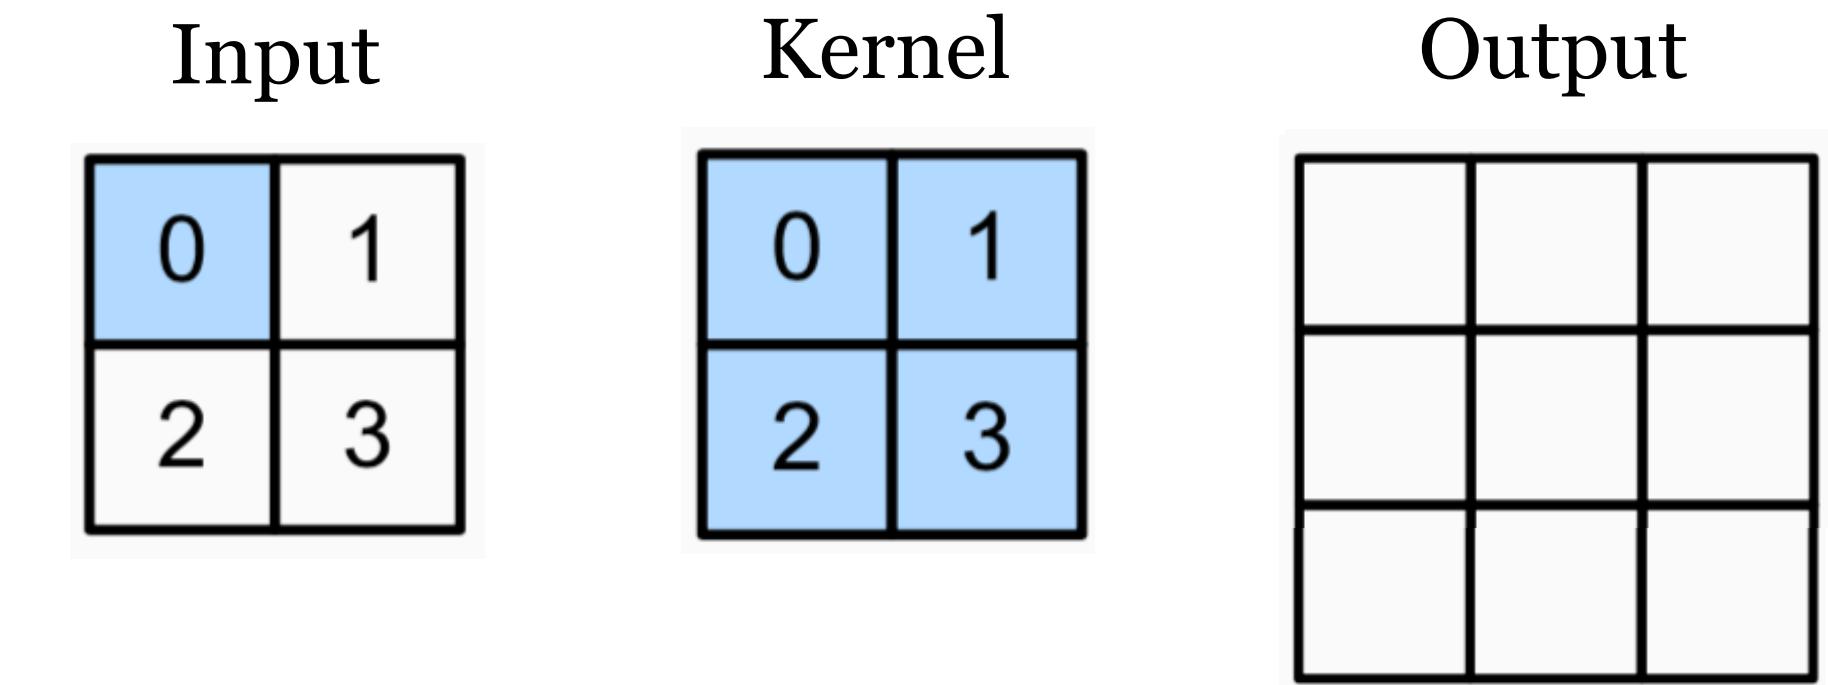
# Encoder-Decoder Architectures

- **Encoder** maps the input to a low-dimensional latent representation by using (often) a CNN.
- In that case it makes sense to use as **decoder** something like the inverse of a CNN.
- This can be done with **transposed convolutions** (i.e., de-convolution) in the decoder.



# Learnable Upsampling: Transposed Convolution

- Consider a 2x2 feature map as **input**.
- We want to expand it to a 3x3 feature map **output**,
- By using a 2x2 **kernel** for transposed convolution.



Each element of the input is multiplied with the kernel and added to the output at the corresponding position.

**Example:** Transposed convolution with stride 1:

$$\begin{array}{cc} \text{Input} & \text{Kernel} \\ \begin{array}{|c|c|}\hline 0 & 1 \\ \hline 2 & 3 \\ \hline \end{array} & \begin{array}{|c|c|}\hline 0 & 1 \\ \hline 2 & 3 \\ \hline \end{array} \end{array} = \begin{array}{c} \begin{array}{|c|c|c|}\hline 0 & 0 & \\ \hline 0 & 0 & \\ \hline & & \\ \hline \end{array} + \begin{array}{|c|c|c|}\hline & 0 & 1 \\ \hline & 2 & 3 \\ \hline & & \\ \hline \end{array} + \begin{array}{|c|c|c|}\hline & & \\ \hline 0 & 2 & \\ \hline 4 & 6 & \\ \hline \end{array} + \begin{array}{|c|c|c|}\hline & & \\ \hline & 0 & 3 \\ \hline & 6 & 9 \\ \hline \end{array} \end{array} = \begin{array}{|c|c|c|}\hline 0 & 0 & 1 \\ \hline 0 & 4 & 6 \\ \hline 4 & 12 & 9 \\ \hline \end{array}$$

# Today

---

Autoencoders

Deep Generative Models

Variational Autoencoders

Architecture & Training

Variational Lower Bound

Experiments & Quantifying Performance

Diederik P. Kingma & Max Welling, "Auto-Encoding Variational Bayes",  
arXiv preprint arXiv:1312.6114 (2013).

# Generative Models

- Given a list of data points  $X = \langle \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \rangle$  coming from a data distribution  $p(\mathbf{x})$ .
- **Generative model:** a model for this distribution, from which we can sample.
- Notation: typically the model distribution is parameterized by  $\theta$  and we write  $p_\theta(\mathbf{x})$ .

## Simple example from previous lectures:

- Assume that data comes from a Gaussian distribution.  
Our model could then be:  $p_\theta(\mathbf{x}) = N(\mu, \Sigma)$ ,
- ... where the parameters  $\theta = (\mu, \Sigma)$  are the mean and covariance matrix of the Gaussian.
- We have already seen how we can learn  $\theta$  with maximum likelihood.



# Variational Autoencoders (VAEs)

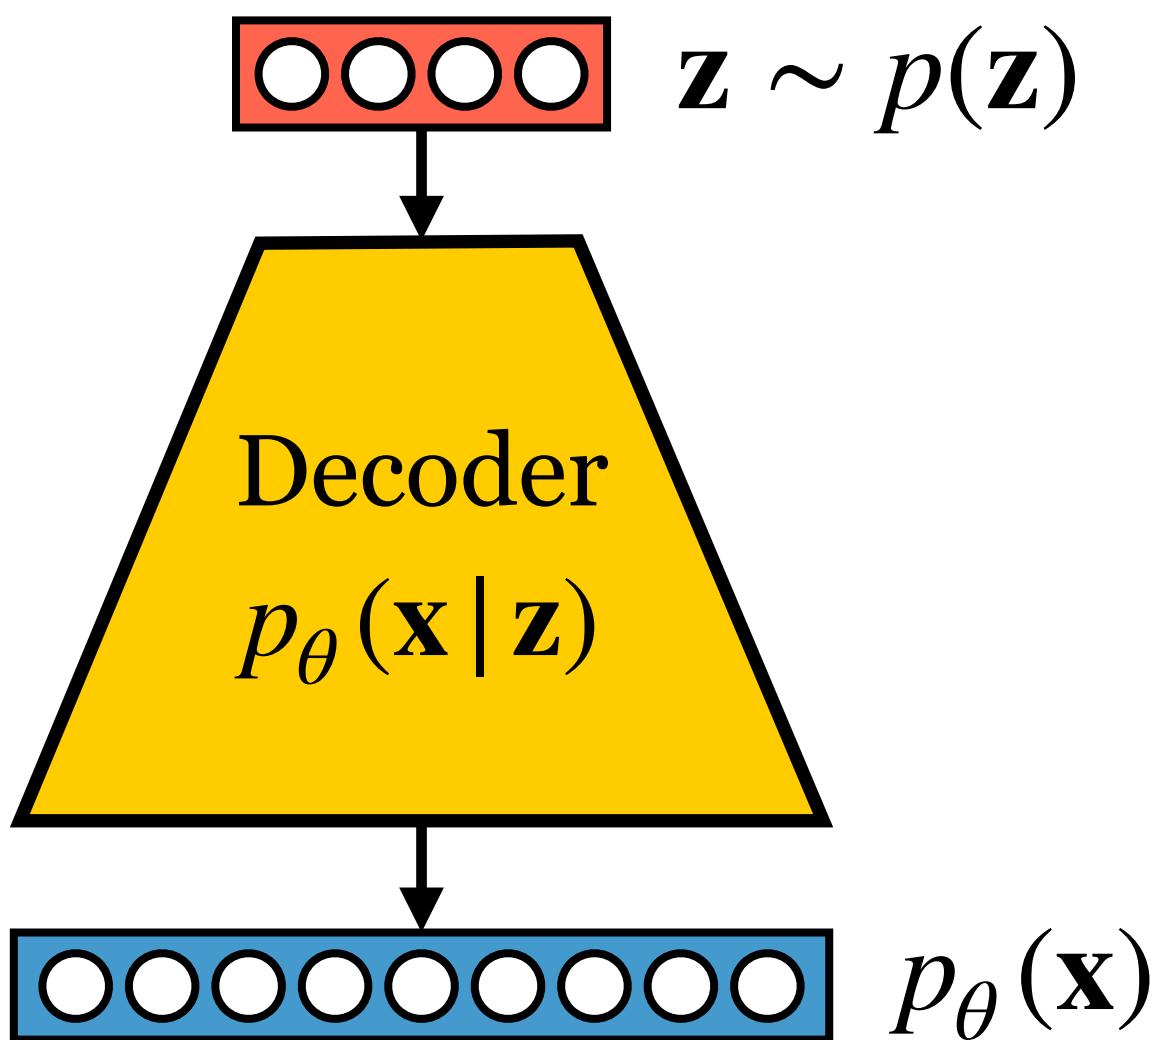
## Deep Generative Models

There are many ways to model more complex distributions.

### Simple Idea:

- Draw latent variables  $\mathbf{z}$  from a simple distribution  $p(\mathbf{z})$ .
- Use a model with parameters  $\theta$  to transform  $p(\mathbf{z})$  to a complex distribution:

$$p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$

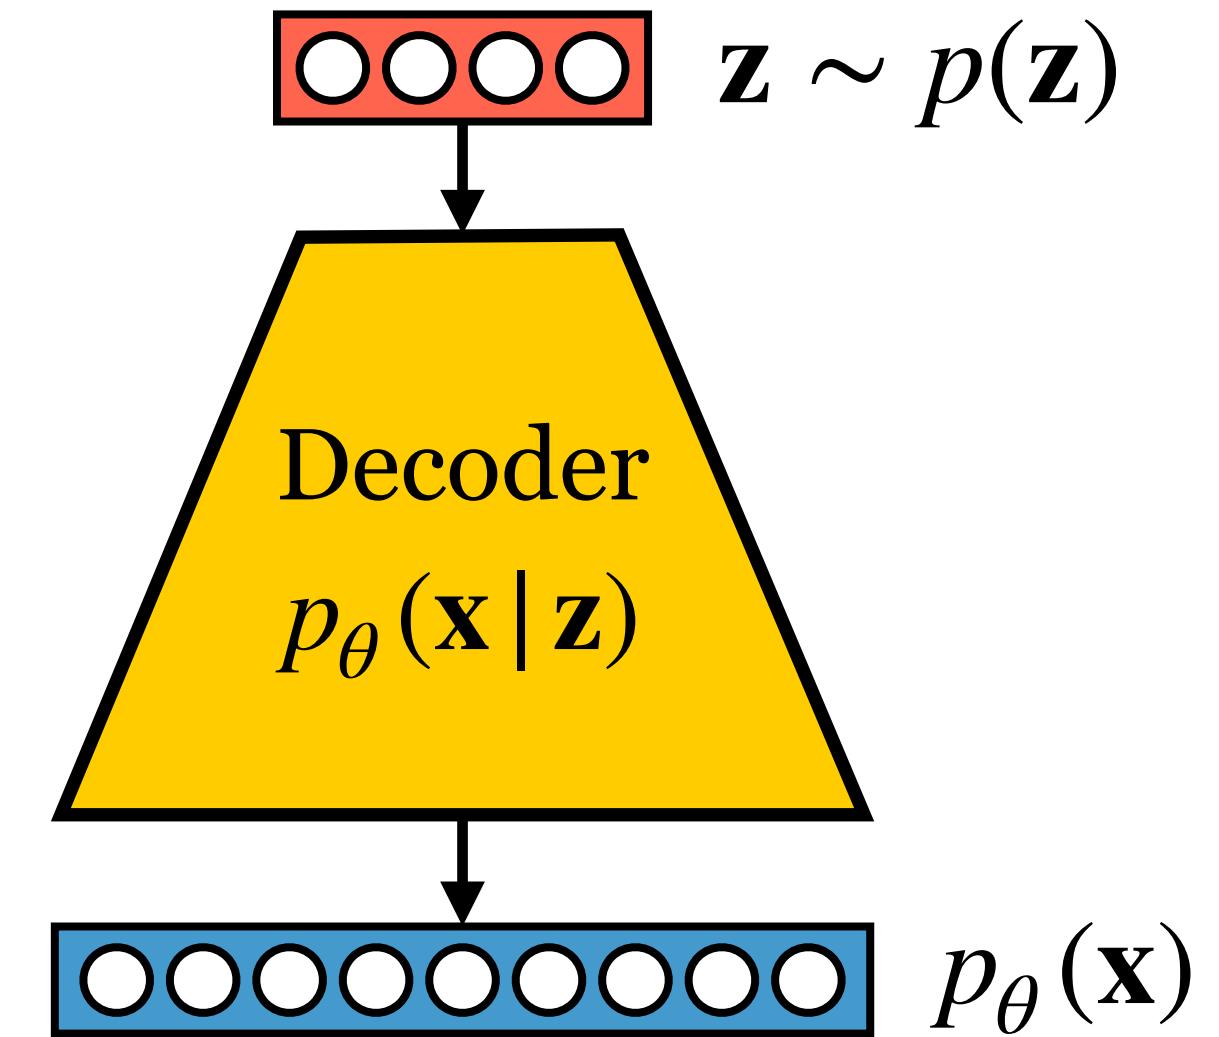


- To generate an  $\mathbf{x}$  :
  - We first sample  $\mathbf{z} \sim p(\mathbf{z})$ , then compute neural network output with  $\mathbf{z}$  as input.
  - In VAEs, we usually use as priors independent standard normal variables (zero mean, unit variance).

# Variational Autoencoders: Training

## How to train VAEs?

- For a dataset  $X = \langle \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \rangle$  we want to maximize the likelihood that it is generated.
- If we would knew  $\mathbf{z}$  for a given  $\mathbf{x}$ , we could simply optimize  $p_\theta(\mathbf{x} | \mathbf{z})$ , but  $\mathbf{z}$  is latent.
- So we would need to compute  $p_\theta(\mathbf{z} | \mathbf{x})$ , which is intractable.



# Variational Autoencoders: Training

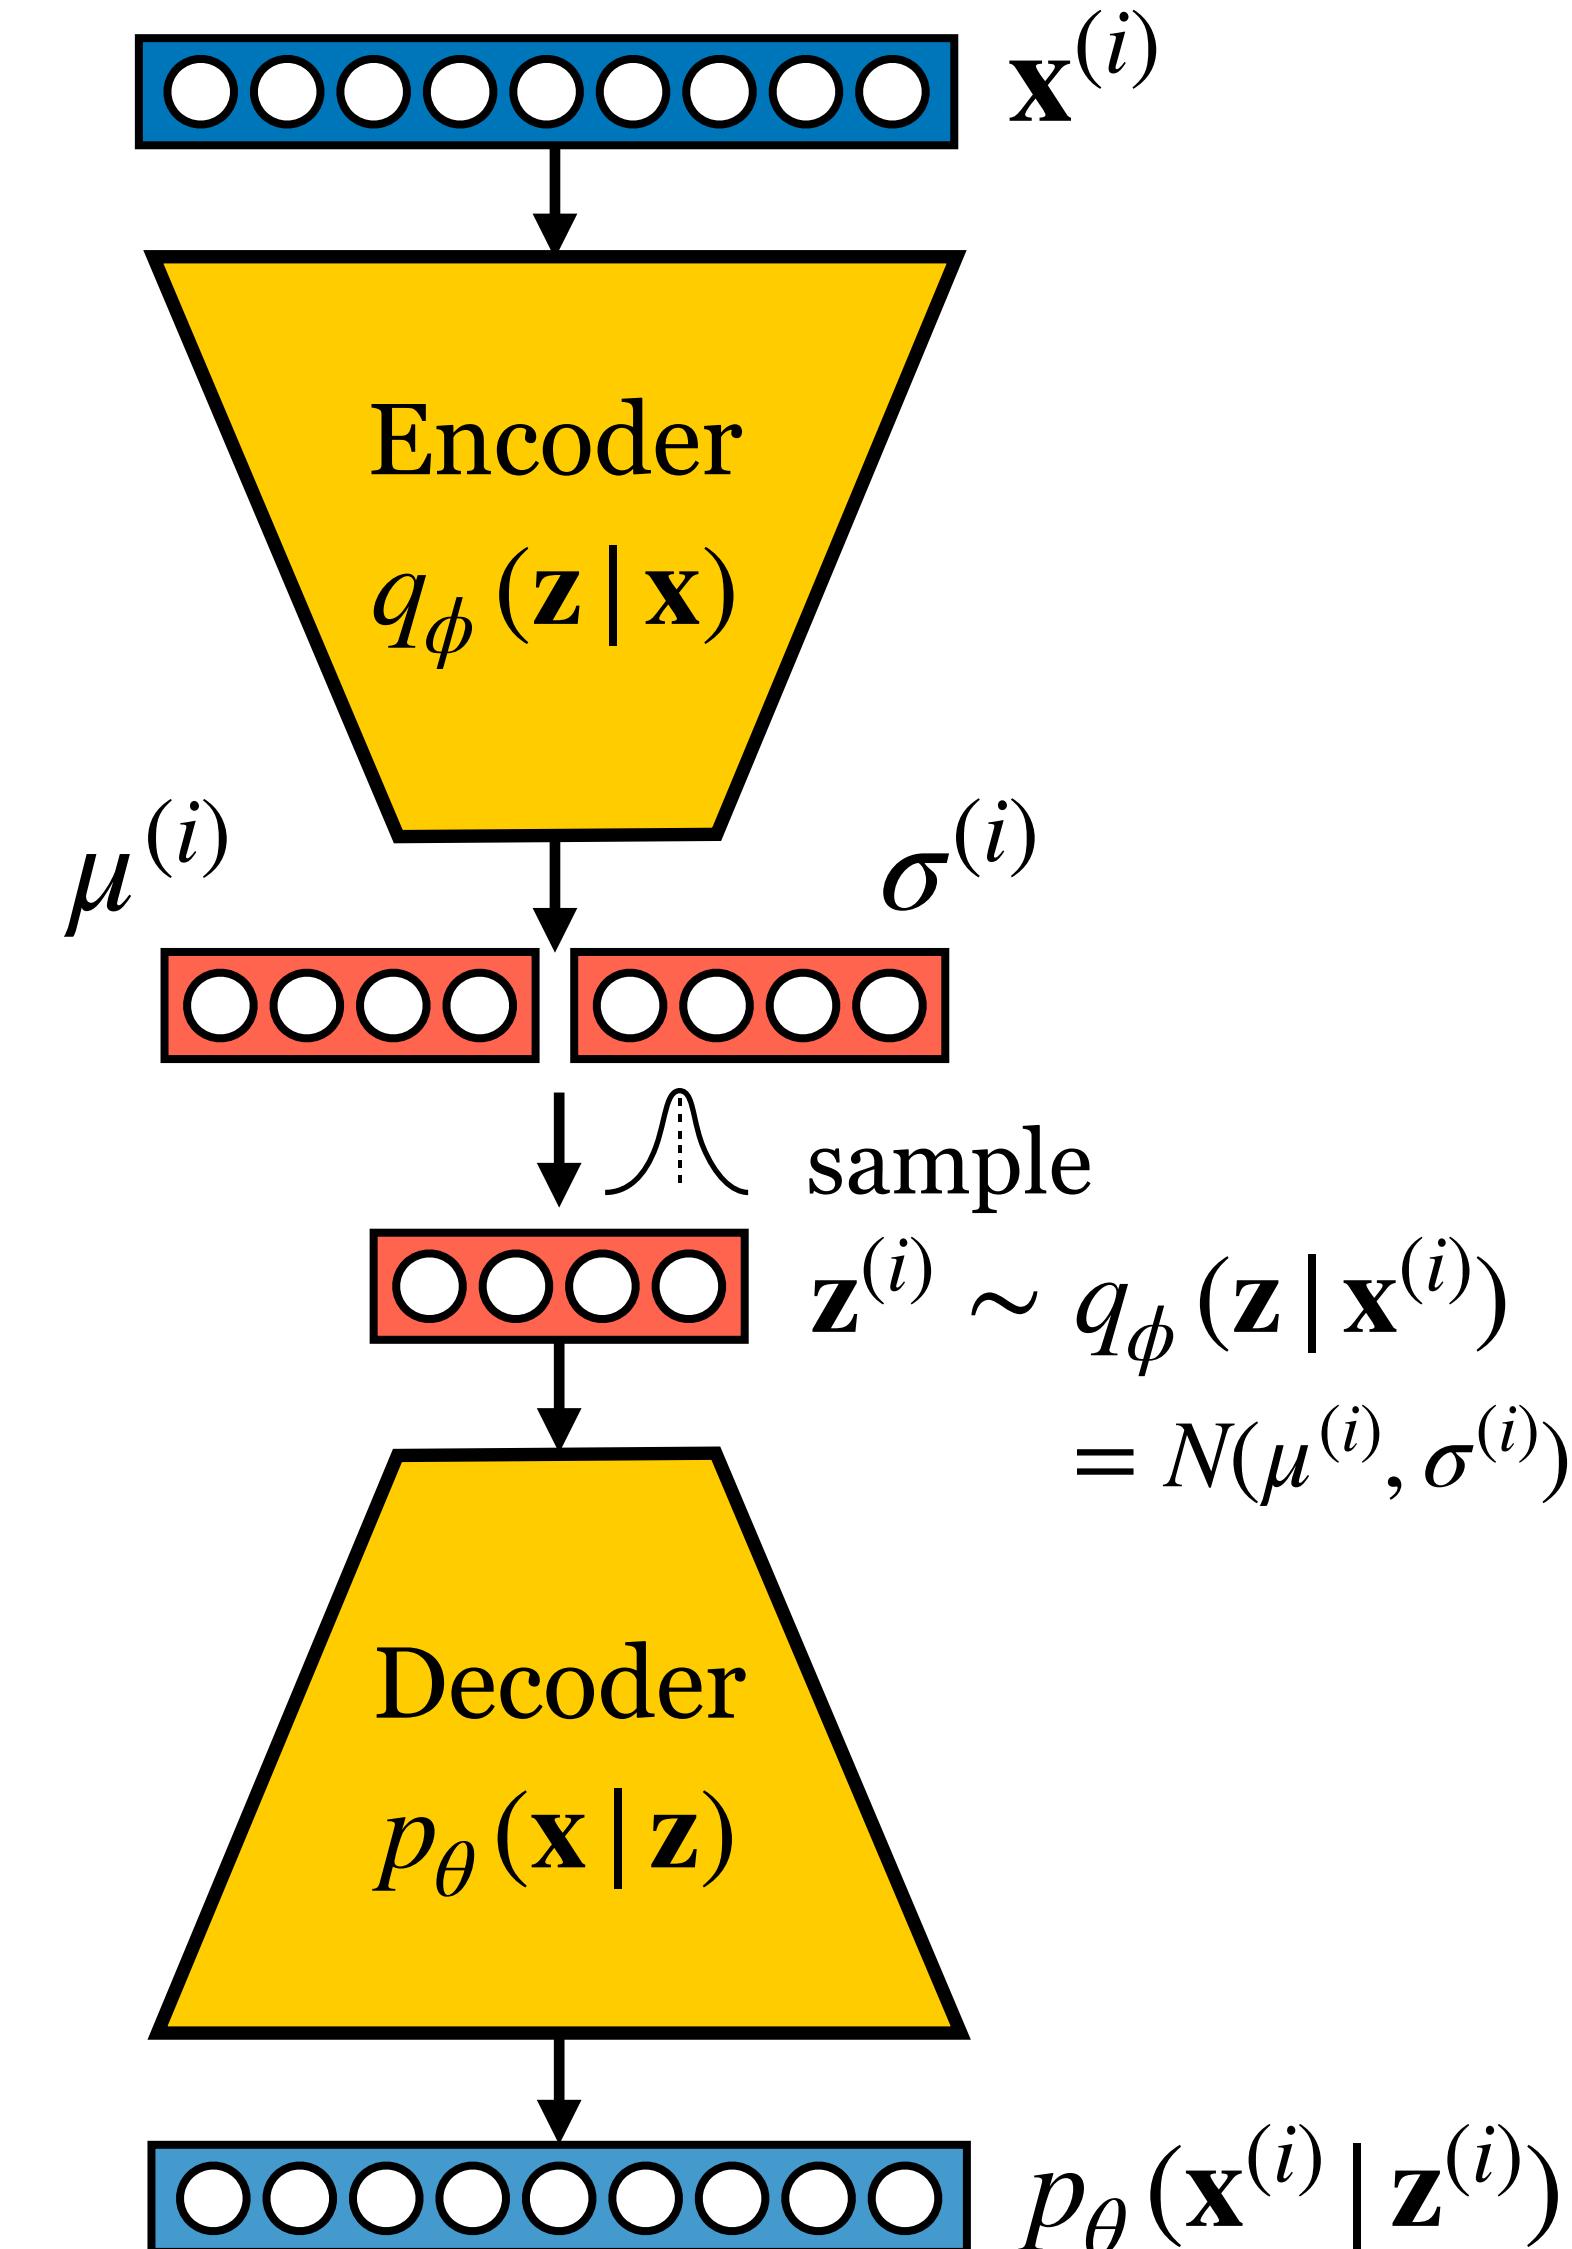
## Solution of VAEs:

- Predict latent variables  $\mathbf{z}$  using a neural network with parameters  $\phi$  that computes the approximate posterior  $q_\phi(\mathbf{z} | \mathbf{x}) \approx p_\theta(\mathbf{z} | \mathbf{x})$ .

For a given  $\mathbf{x}$ ,

- the **encoder** samples  $\mathbf{z} \sim q_\phi(\mathbf{z} | \mathbf{x})$ ,
- the **decoder** takes  $\mathbf{z}$  and computes  $p_\theta(\mathbf{x} | \mathbf{z})$ .

Now we have  $\mathbf{z}$  and can optimize both  $\phi$  and  $\theta$ .

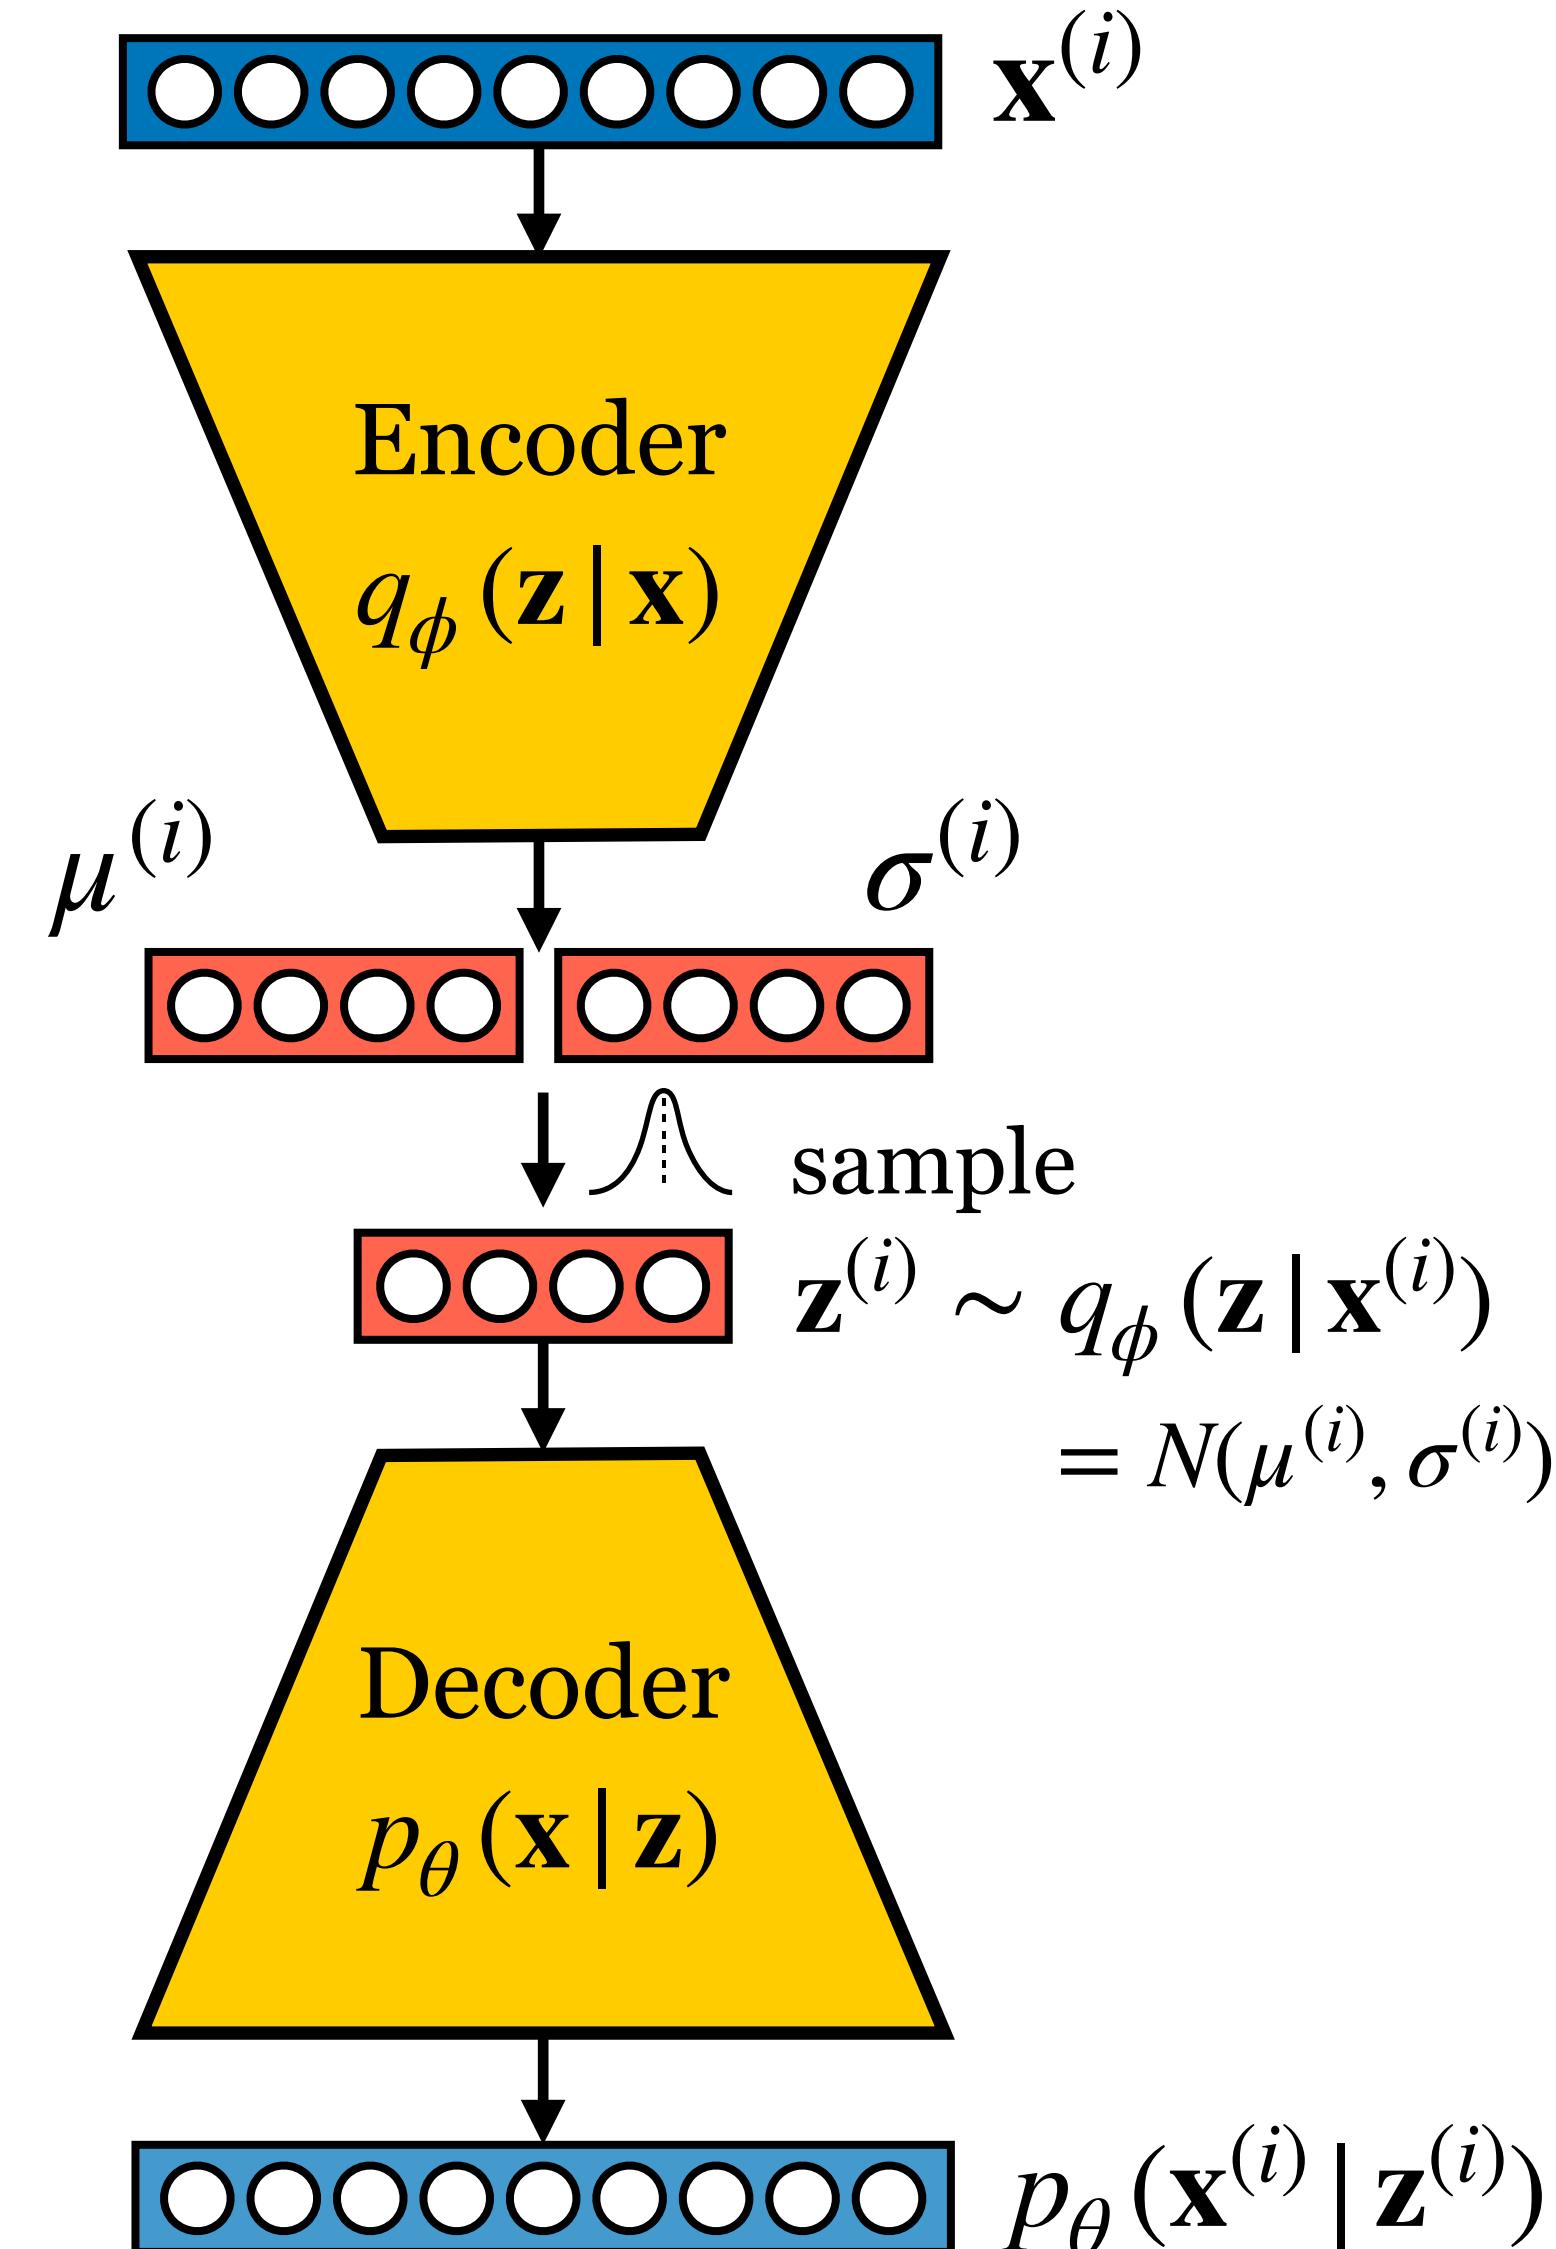


# Variational Autoencoders: Training

## Training objective

We want that:

- 1) The distribution  $q_\phi(\mathbf{z} | \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ .
- 2) The decoder computes a good reconstruction of  $\mathbf{x}^{(i)}$ .



# Variational Autoencoders: Training

---

1) The distribution  $q_\phi(\mathbf{z} \mid \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ :

- We achieve this by minimizing the KL divergence between these distributions.

$$D_{KL}(q \parallel p) = \int q(x) \log \frac{q(x)}{p(x)} dx$$

- We will use a Gaussian prior with zero-mean independent variables of unit variance:  $p(\mathbf{z}) = N(\mathbf{0}, \mathbf{I})$ , i.e.,  $p(z_m) = N(0, 1)$ .
- For each latent variable  $z_m$ , the encoder produces a Gaussian distribution with mean  $\mu_m$  and standard deviation  $\sigma_m$ , i.e.,

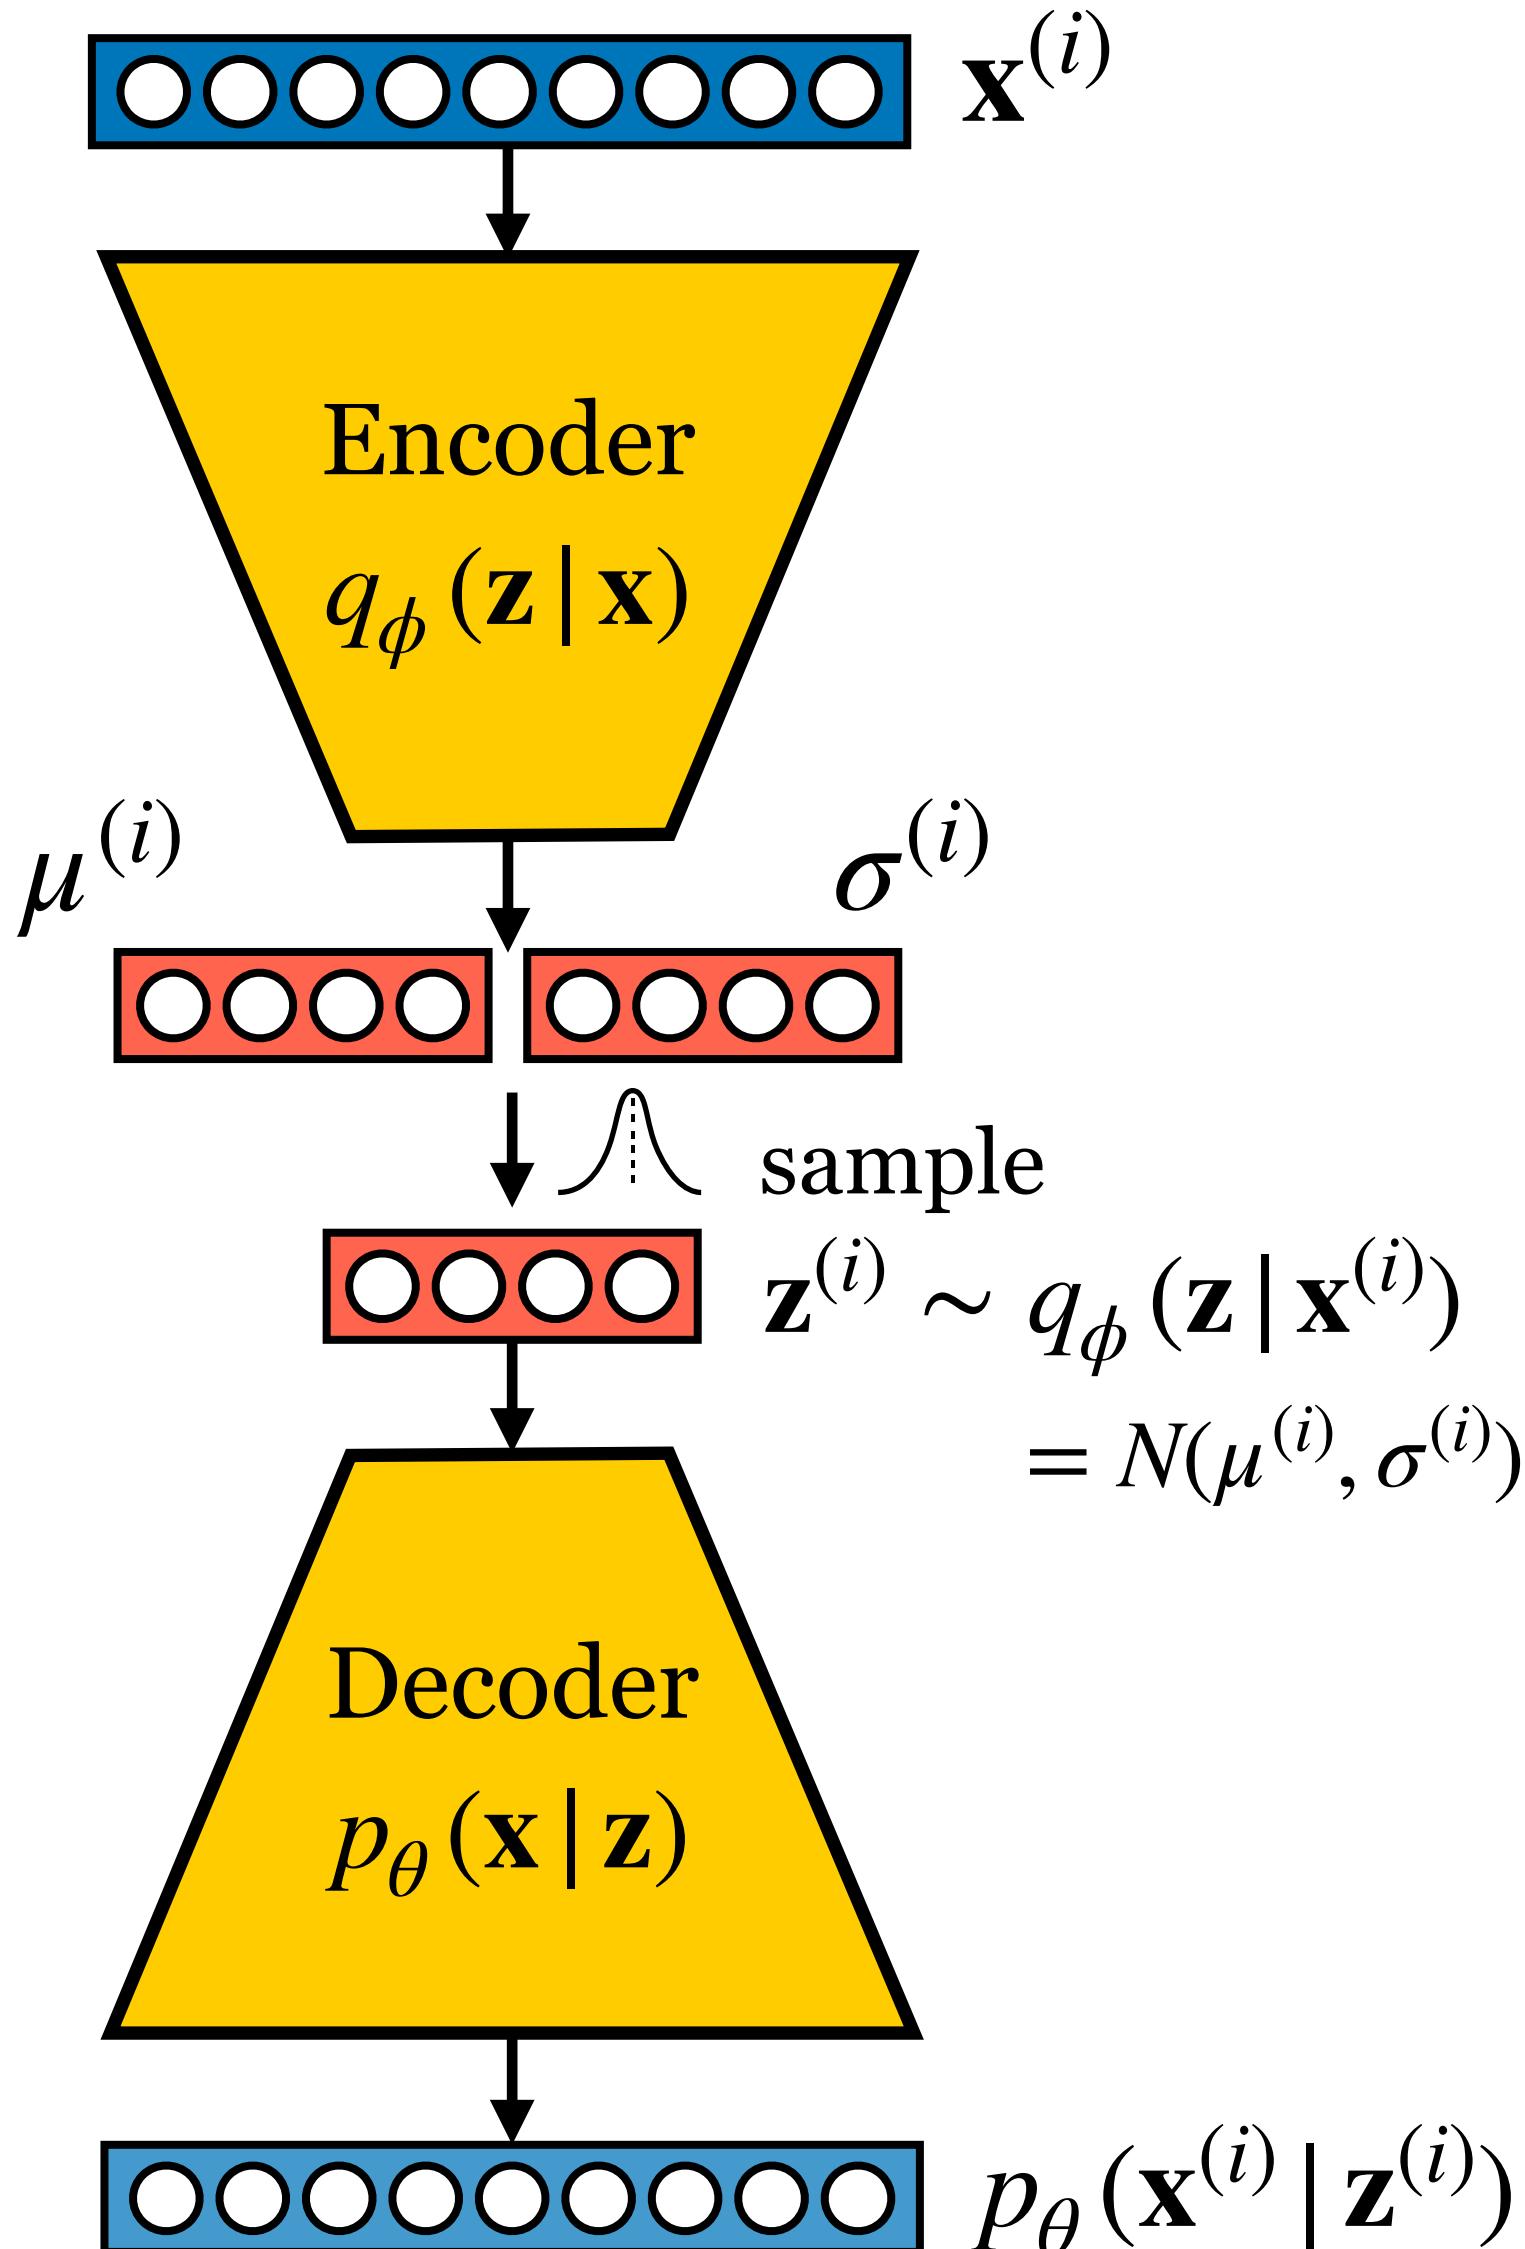
$$q_\phi(z_m \mid \mathbf{x}^{(i)}) = N(\mu_m(\mathbf{x}^{(i)}), \sigma_m(\mathbf{x}^{(i)}))$$

- The KL divergence in this case is:  $D_{KL}(q_\phi \parallel p) = -\frac{1}{2} \sum_m \left( 1 + \log (\sigma_m(\mathbf{x}^{(i)})^2) - \mu_m(\mathbf{x}^{(i)})^2 - \sigma_m(\mathbf{x}^{(i)})^2 \right)$

# Variational Autoencoders: Training

- 1) The distribution  $q_{\phi}(\mathbf{z} | \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ :
- 2) The decoder computes a good reconstruction of  $\mathbf{x}^{(i)}$ :
  - We achieve this by maximizing the likelihood on  $L$  drawn samples  $\mathbf{z}^{(i,1)}, \dots, \mathbf{z}^{(i,L)}$ .

$$\frac{1}{L} \sum_{l=1}^L \log p_{\theta} (\mathbf{x}^{(i)} | \mathbf{z}^{(i,l)})$$



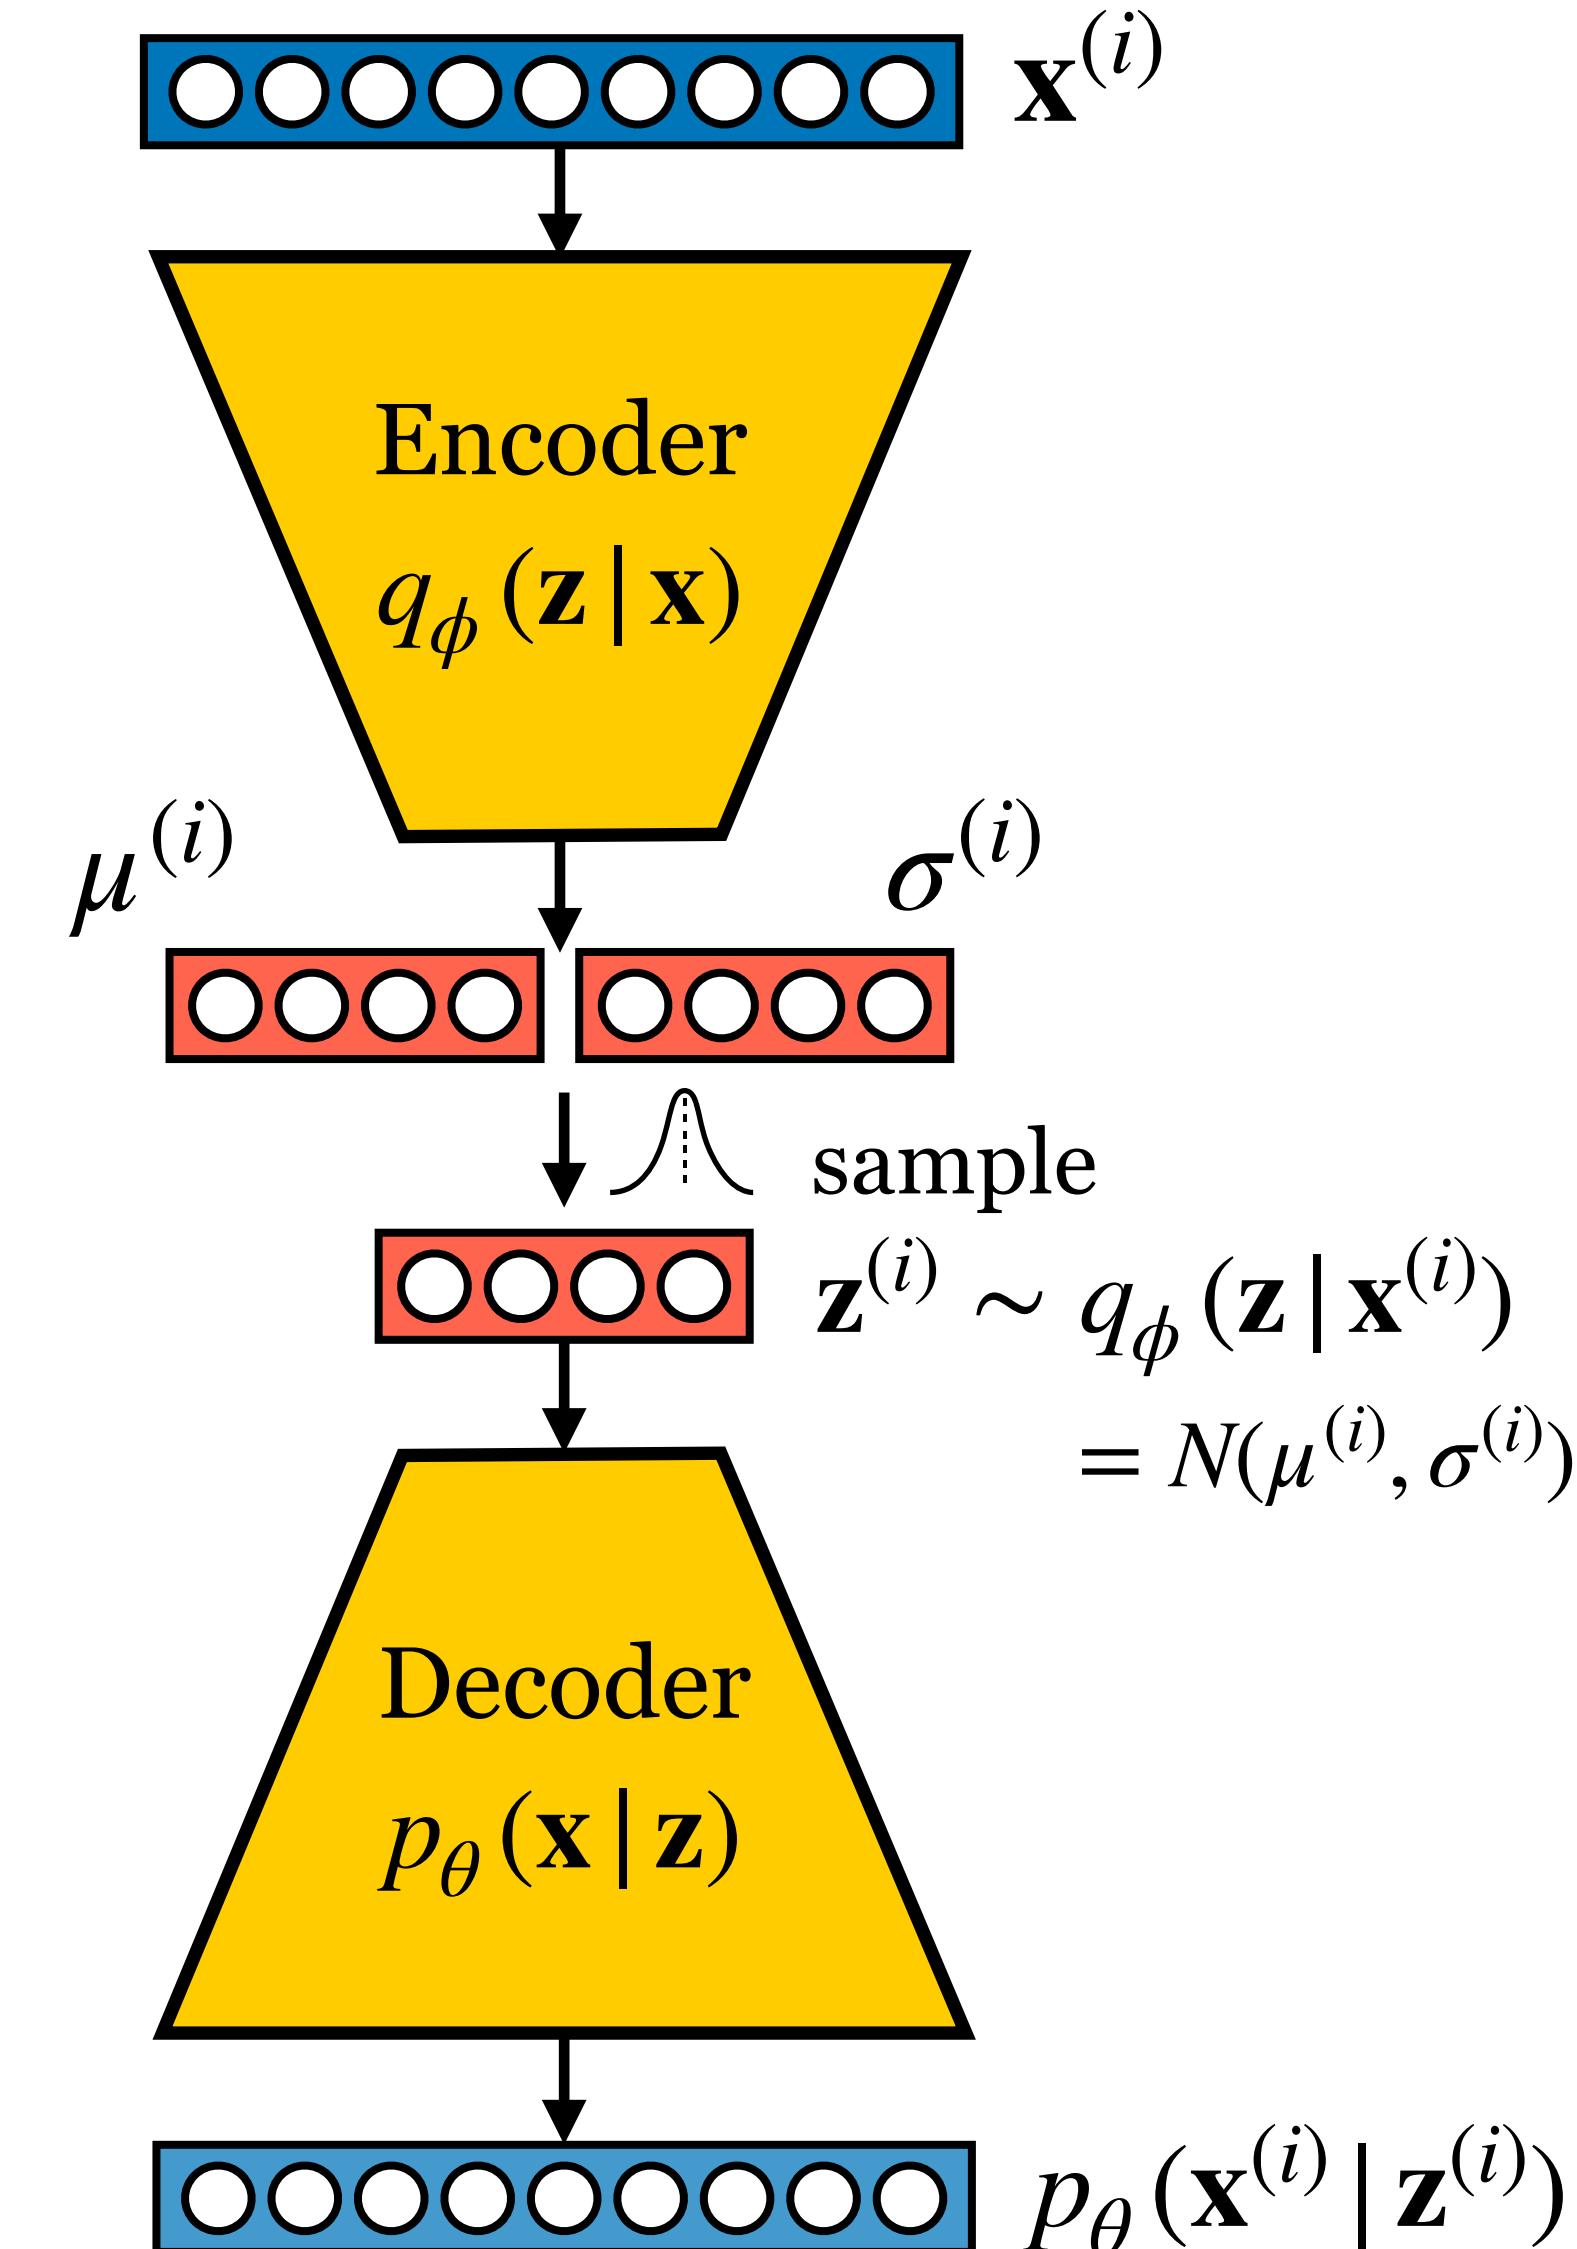
# Variational Autoencoders: Training

- 1) The distribution  $q_\phi(\mathbf{z} | \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ :
- 2) The decoder computes a good reconstruction of  $\mathbf{x}^{(i)}$ :

► Both parts together give rise to the loss to be minimized:

$$\mathcal{L}(\theta, \phi, X) = D_{KL}\left(q_\phi(\mathbf{z} | \mathbf{x}^{(i)}) \| p(\mathbf{z})\right) - \frac{1}{L} \sum_{l=1}^L \log p_\theta\left(\mathbf{x}^{(i)} | \mathbf{z}^{(i,l)}\right)$$

approx. posterior  $\approx$  prior                          "reconstruction" term

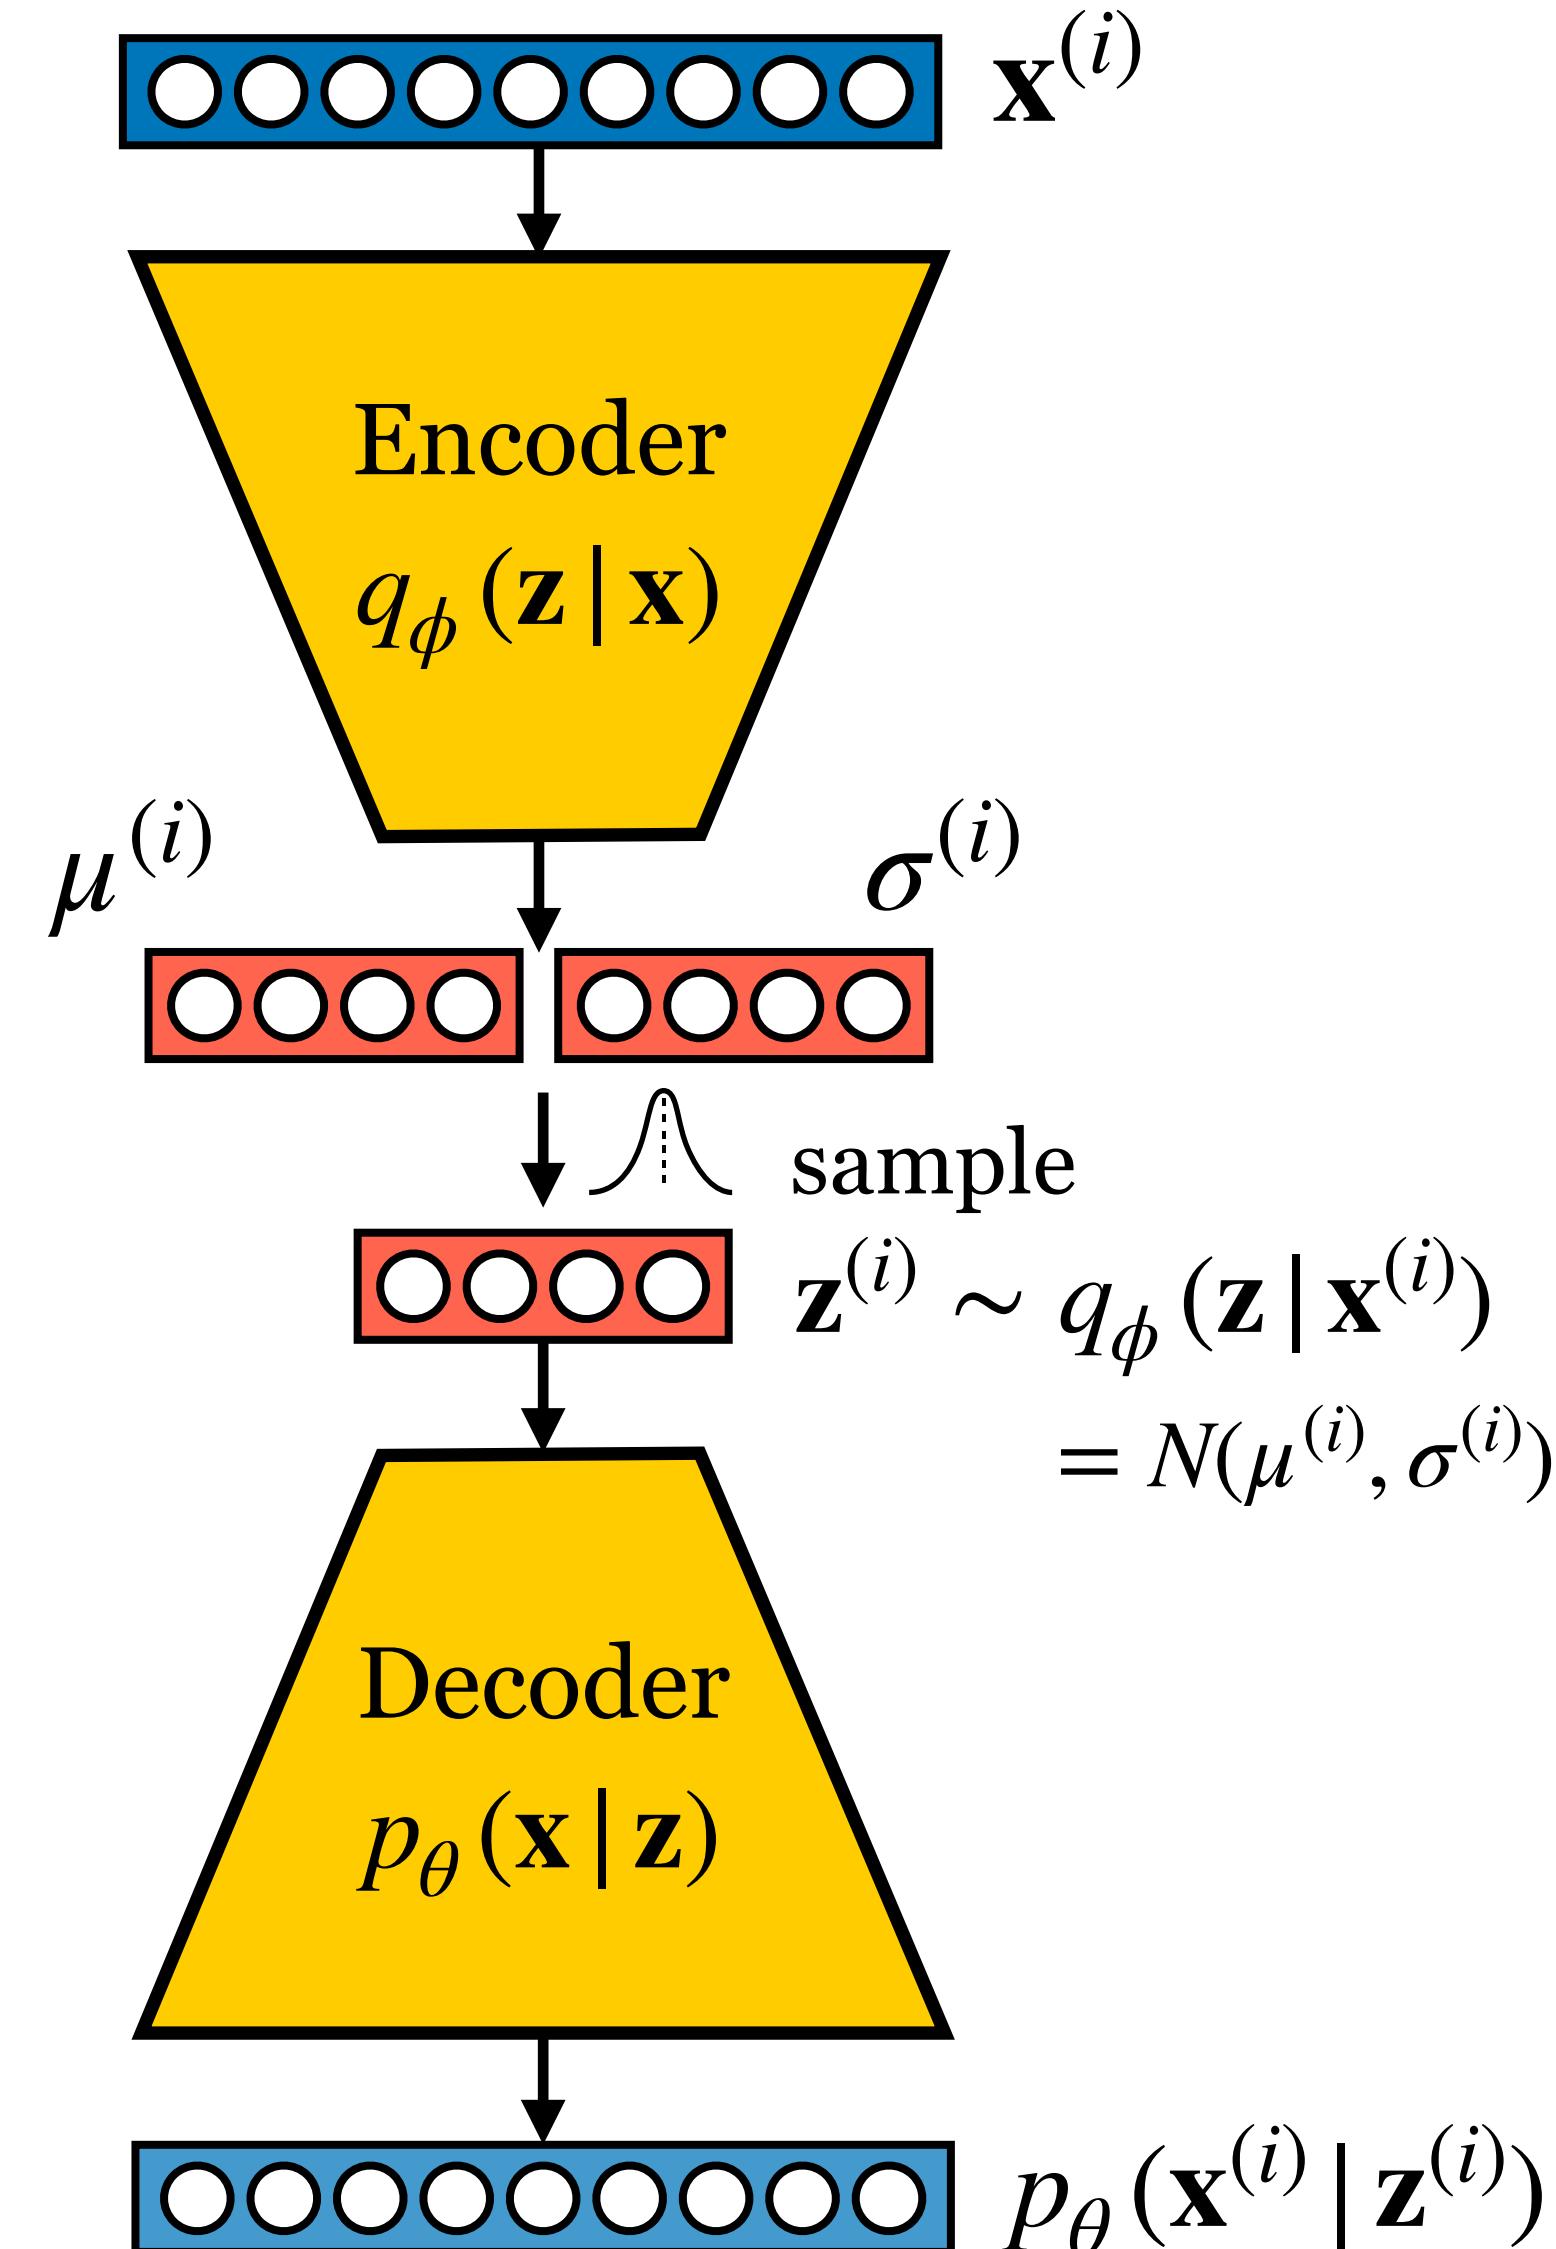


# Variational Autoencoders: Training

- 1) The distribution  $q_{\phi}(\mathbf{z} | \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ :
- 2) The decoder computes a good reconstruction of  $\mathbf{x}^{(i)}$ :

► **Note:** Once the model is trained,

- The encoder is not needed for generation.
  - We just need to sample from the latent prior  $p(\mathbf{z})$ , and then use the decoder.
- The encoder can be used to extract the latent variables.
  - Low-dimensional representations of data.
  - For this we don't need the decoder.



# Variational Autoencoders: Training

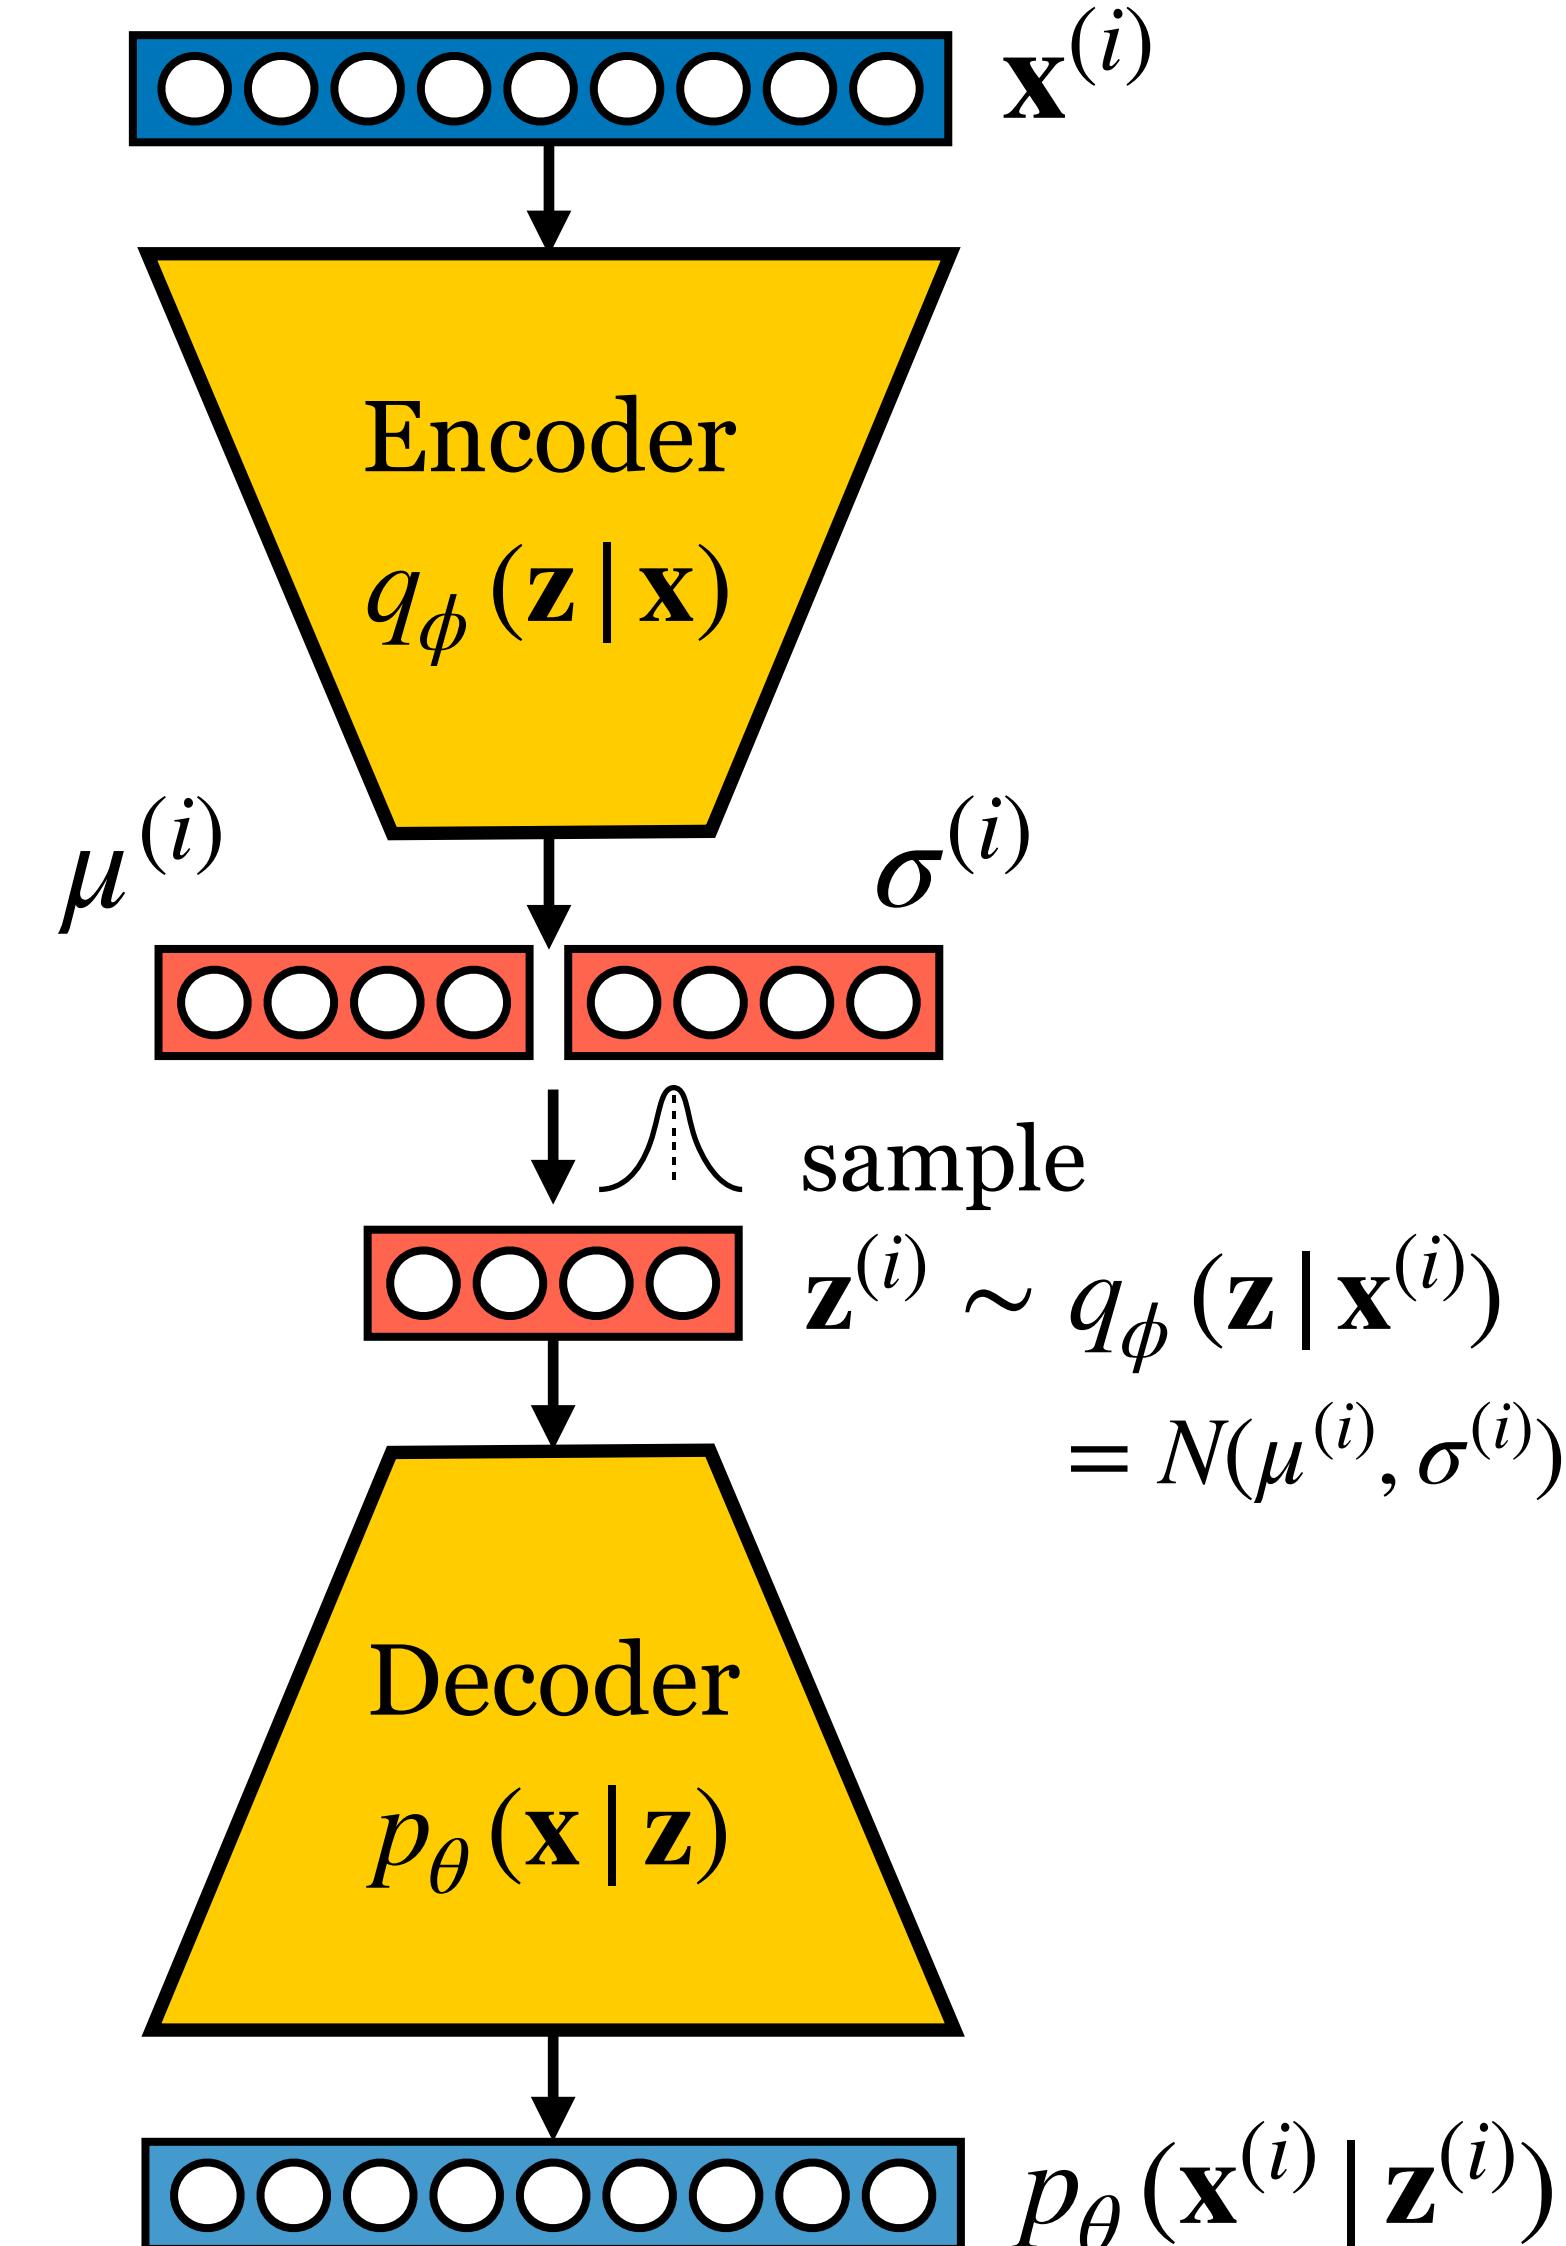
- 1) The distribution  $q_\phi(\mathbf{z} | \mathbf{x}^{(i)})$  is close to our prior distribution  $p(\mathbf{z})$ :
- 2) The decoder computes a good reconstruction of  $\mathbf{x}^{(i)}$ :

► Both parts together give rise to the loss to be minimized:

$$\mathcal{L}(\theta, \phi, X) = D_{KL}\left(q_\phi(\mathbf{z} | \mathbf{x}^{(i)}) \| p(\mathbf{z})\right) - \frac{1}{L} \sum_{l=1}^L \log p_\theta\left(\mathbf{x}^{(i)} | \mathbf{z}^{(i,l)}\right)$$

approx. posterior  $\approx$  prior                          "reconstruction" term

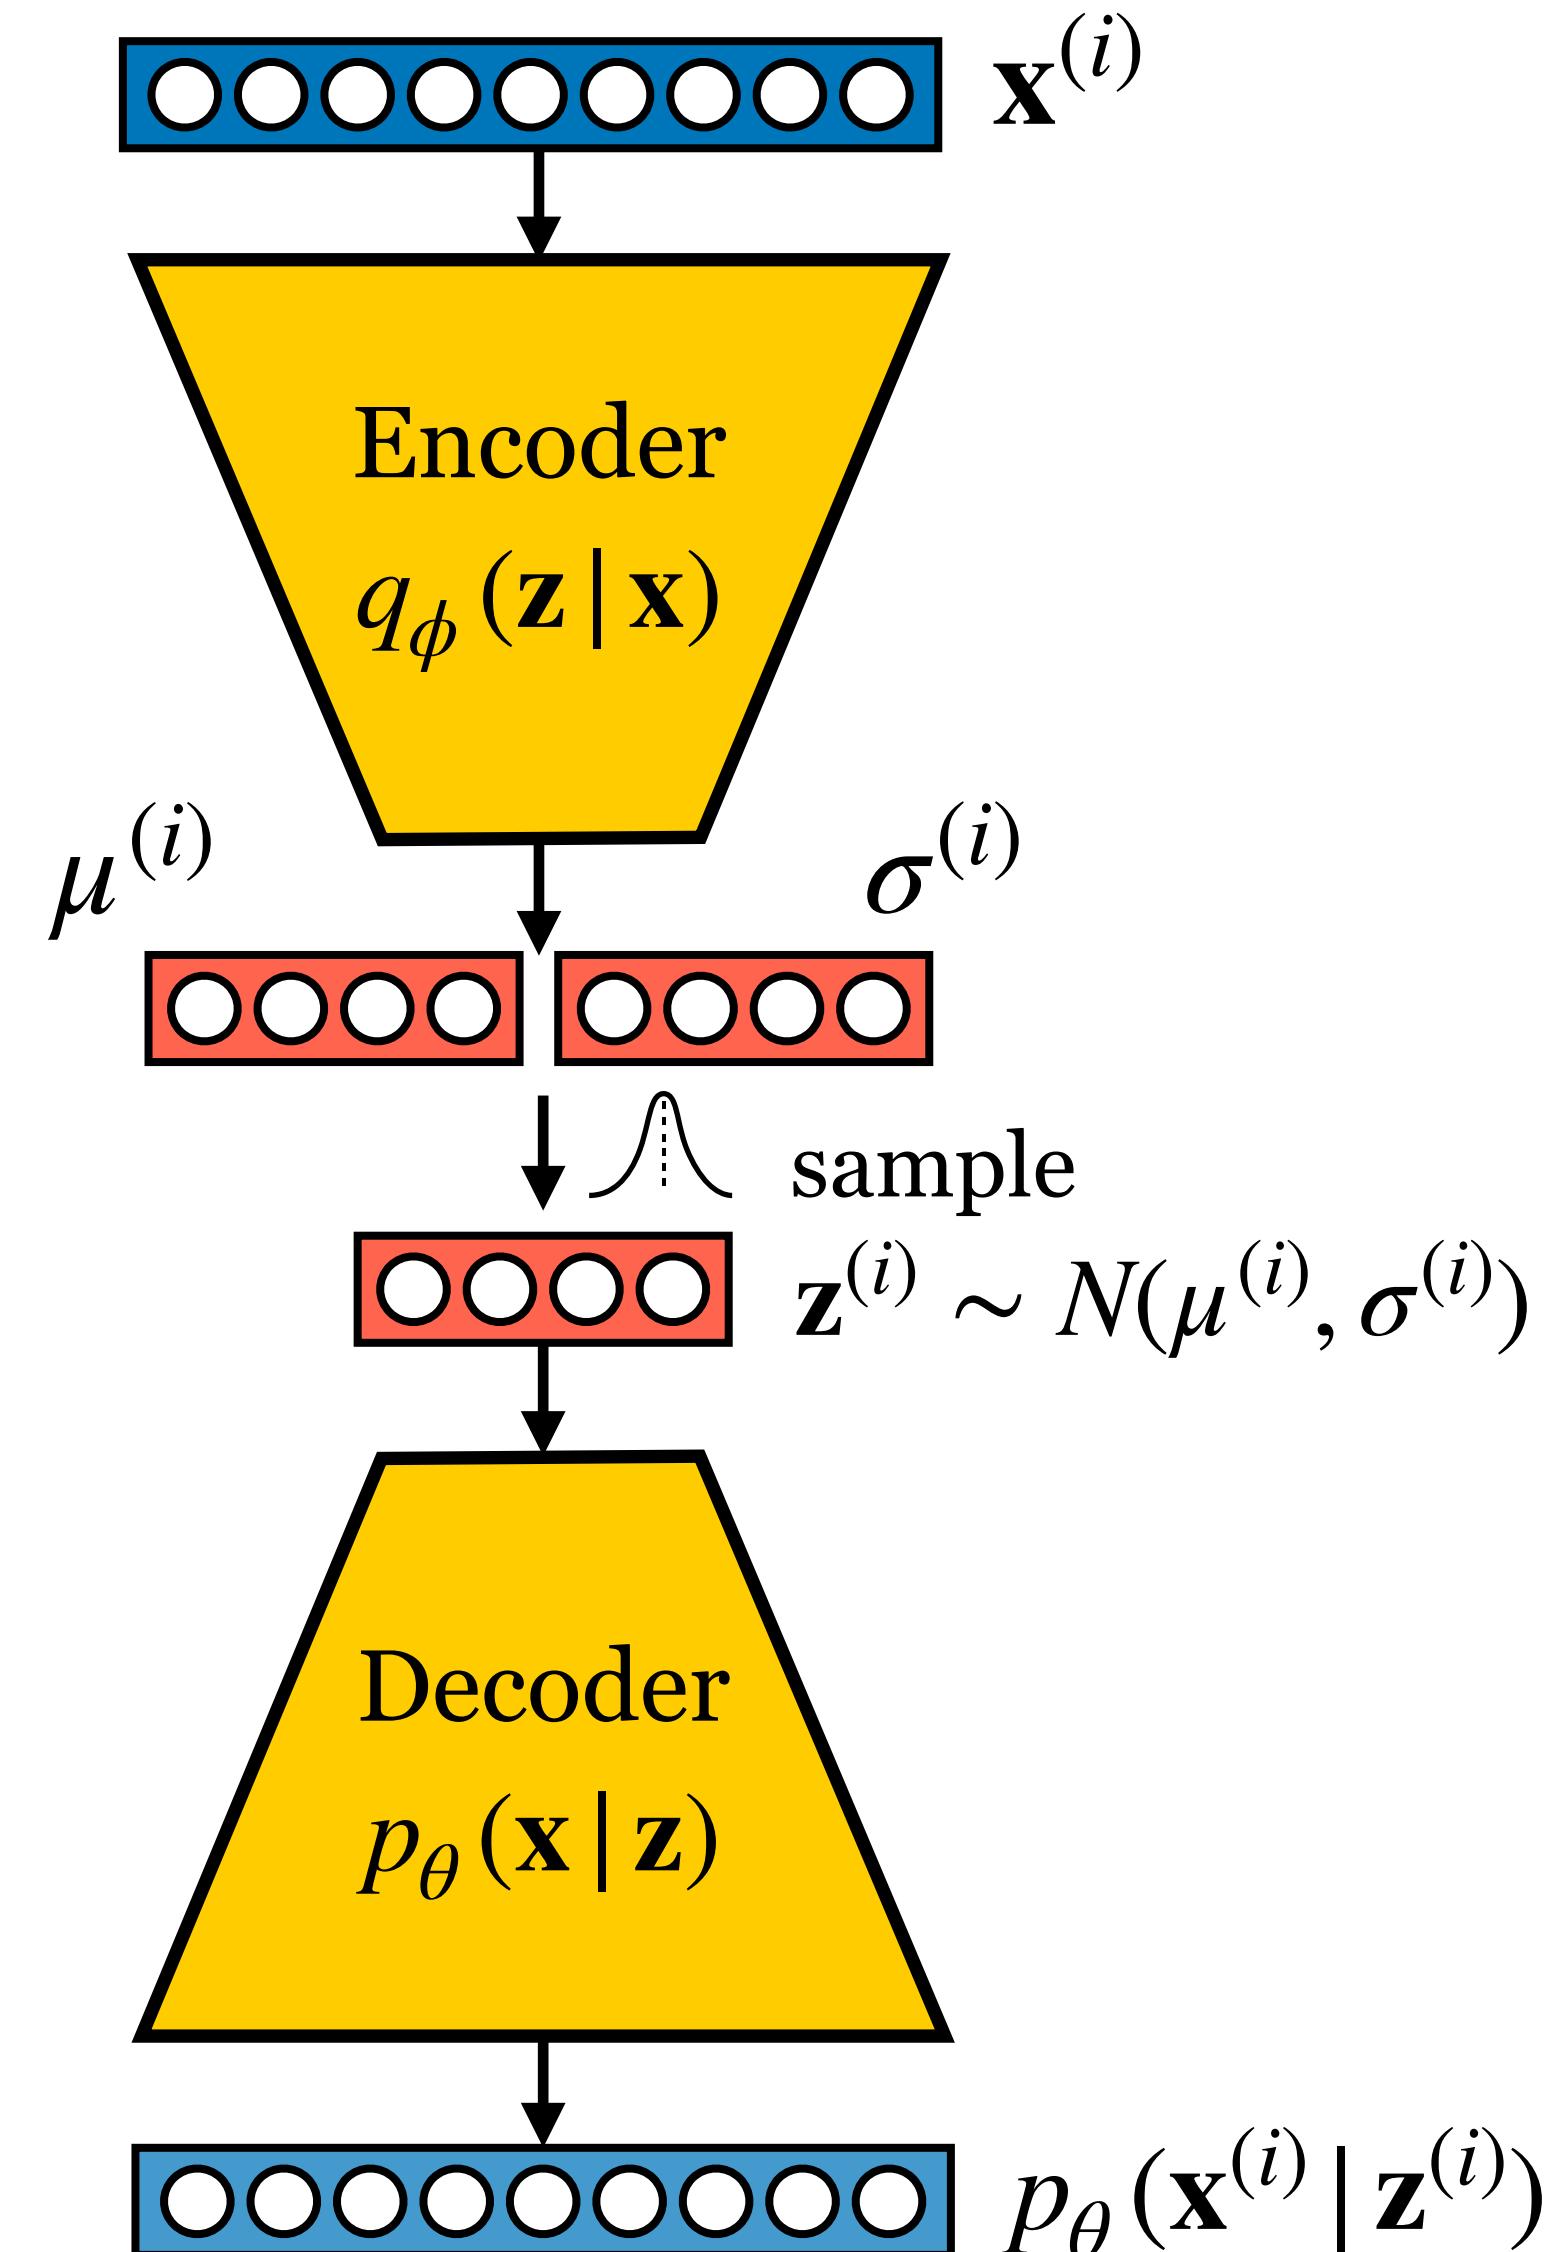
► We want to use gradient descent on mini-batches,  
**backpropagating** loss gradients through both networks.



# Reparametrization Trick

**Problem:**  $\mathbf{z}$  is a sampled variable.

- We can compute  $\frac{\partial \mathcal{L}}{\partial \mathbf{z}}$ , but not  $\frac{\partial \mathbf{z}}{\partial \mu}$  and  $\frac{\partial \mathbf{z}}{\partial \sigma}$  because the sampler is not differentiable.



# Reparametrization Trick

**Problem:**  $\mathbf{z}$  is a sampled variable.

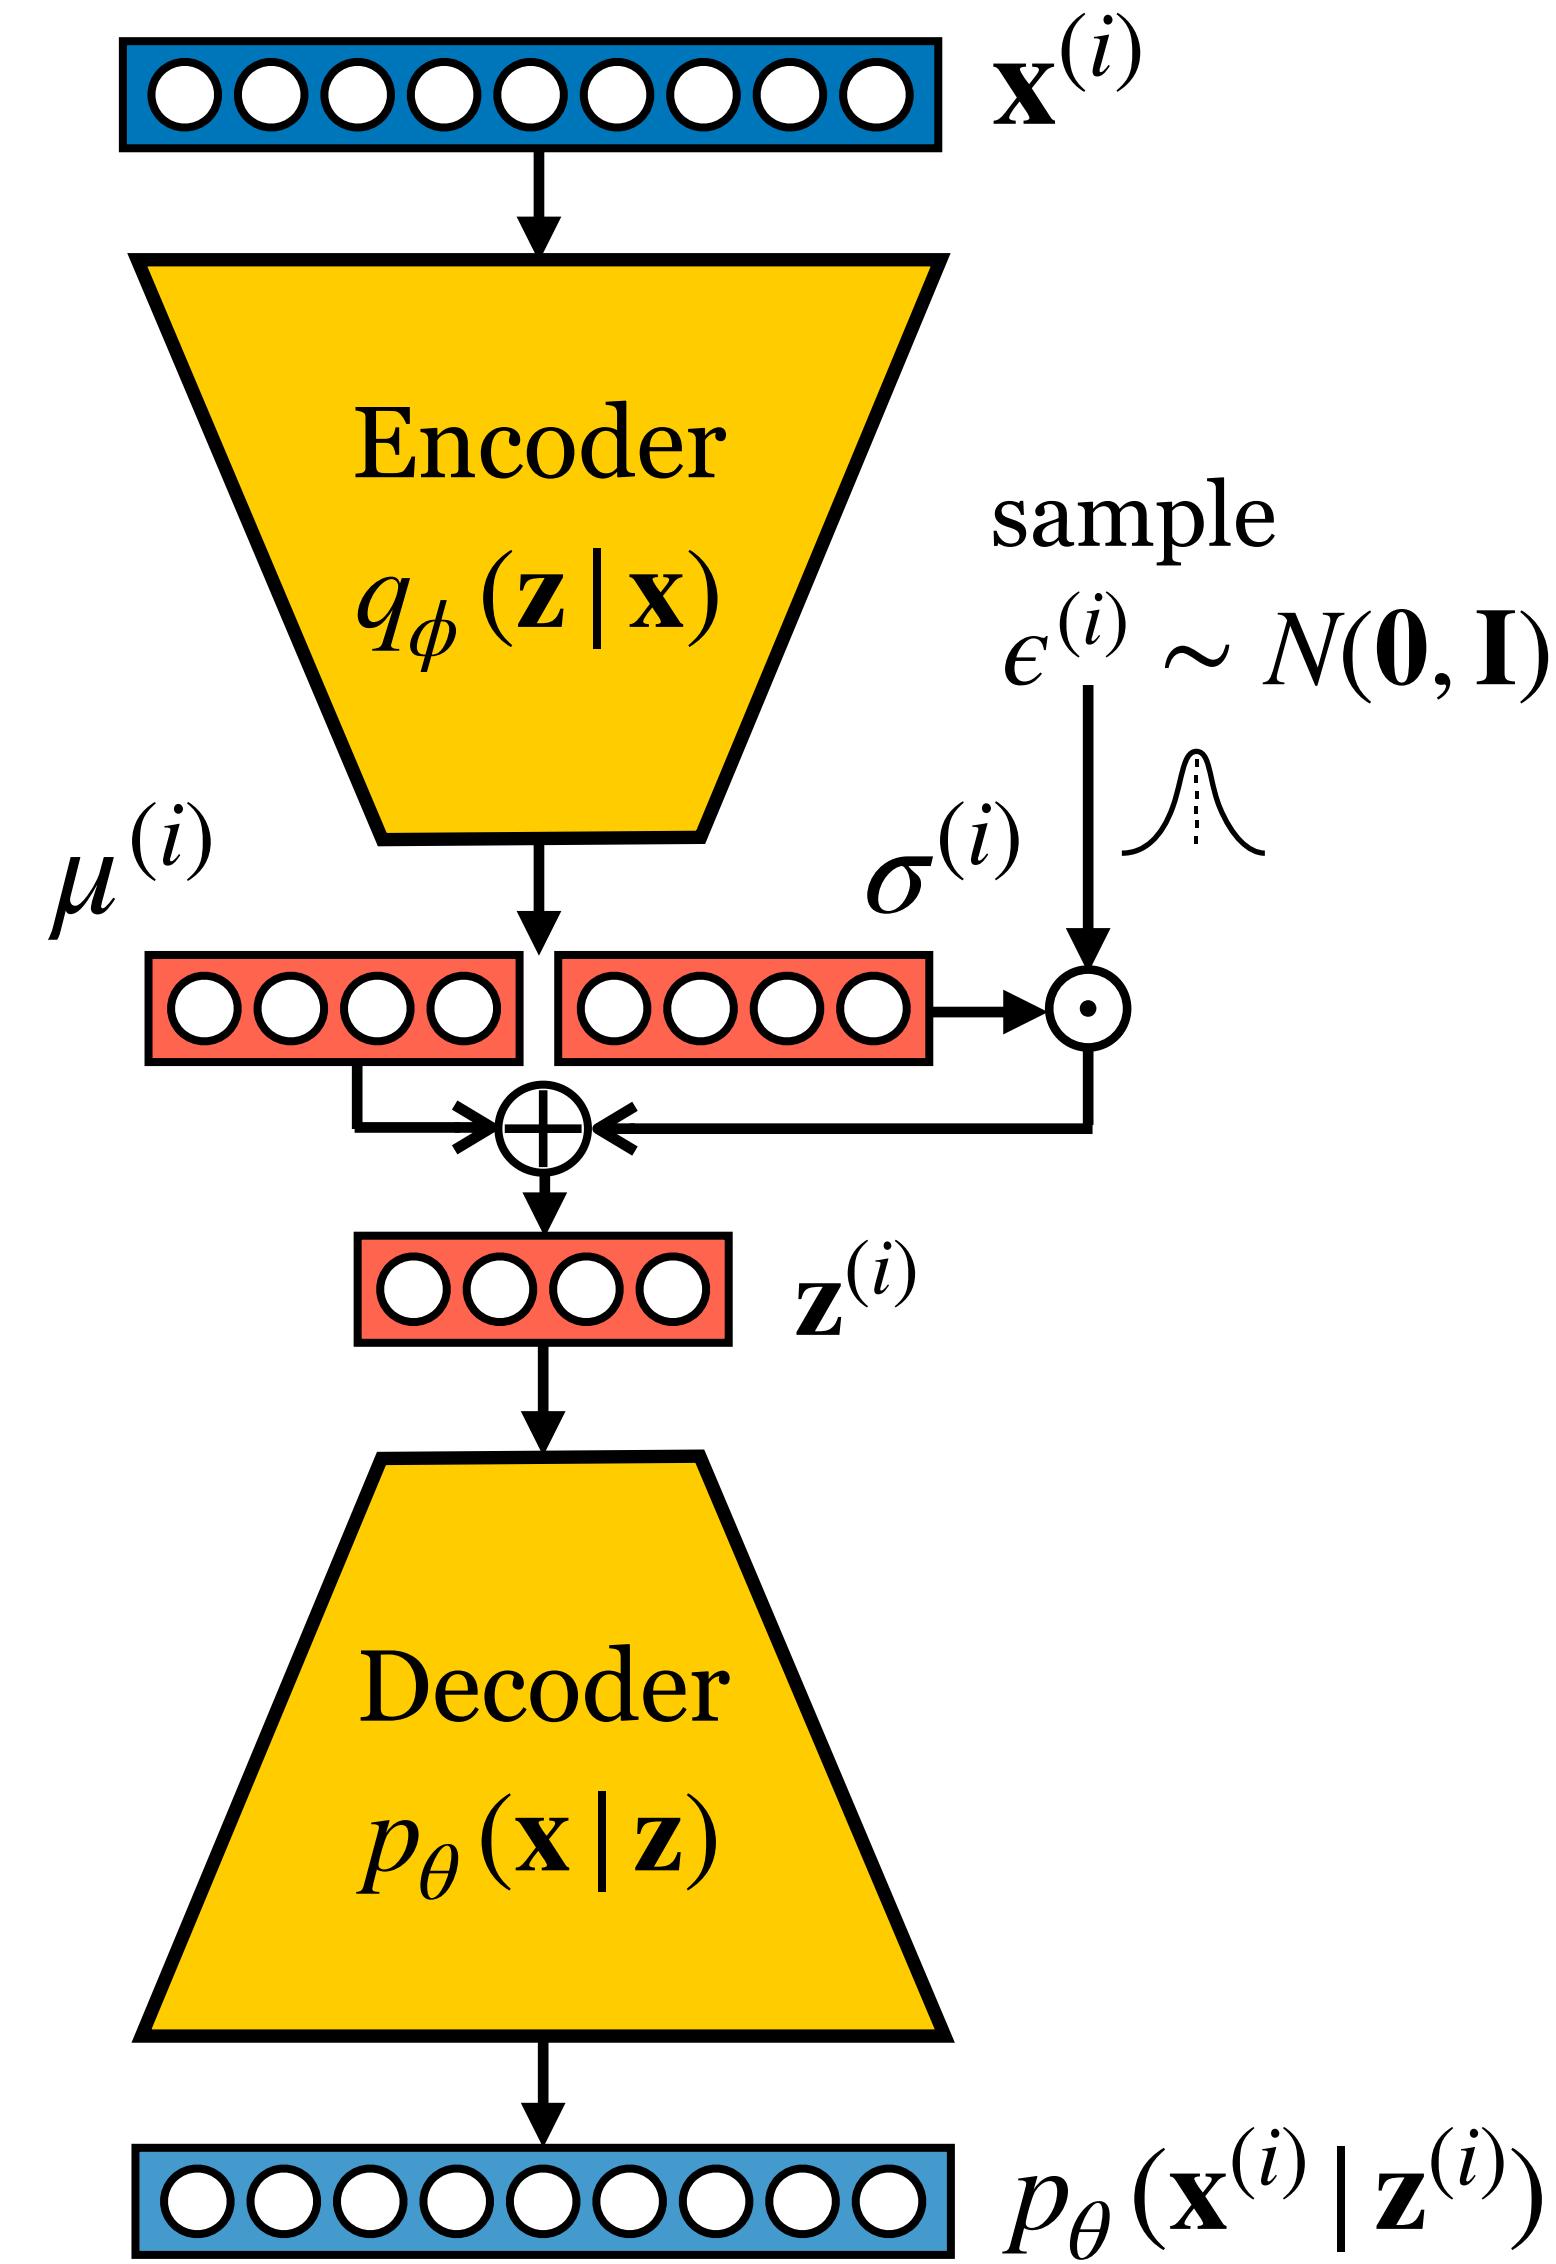
- We can compute  $\frac{\partial \mathcal{L}}{\partial \mathbf{z}}$ , but not  $\frac{\partial \mathbf{z}}{\partial \mu}$  and  $\frac{\partial \mathbf{z}}{\partial \sigma}$  because the sampler is not differentiable.

**Reparametrization trick:**

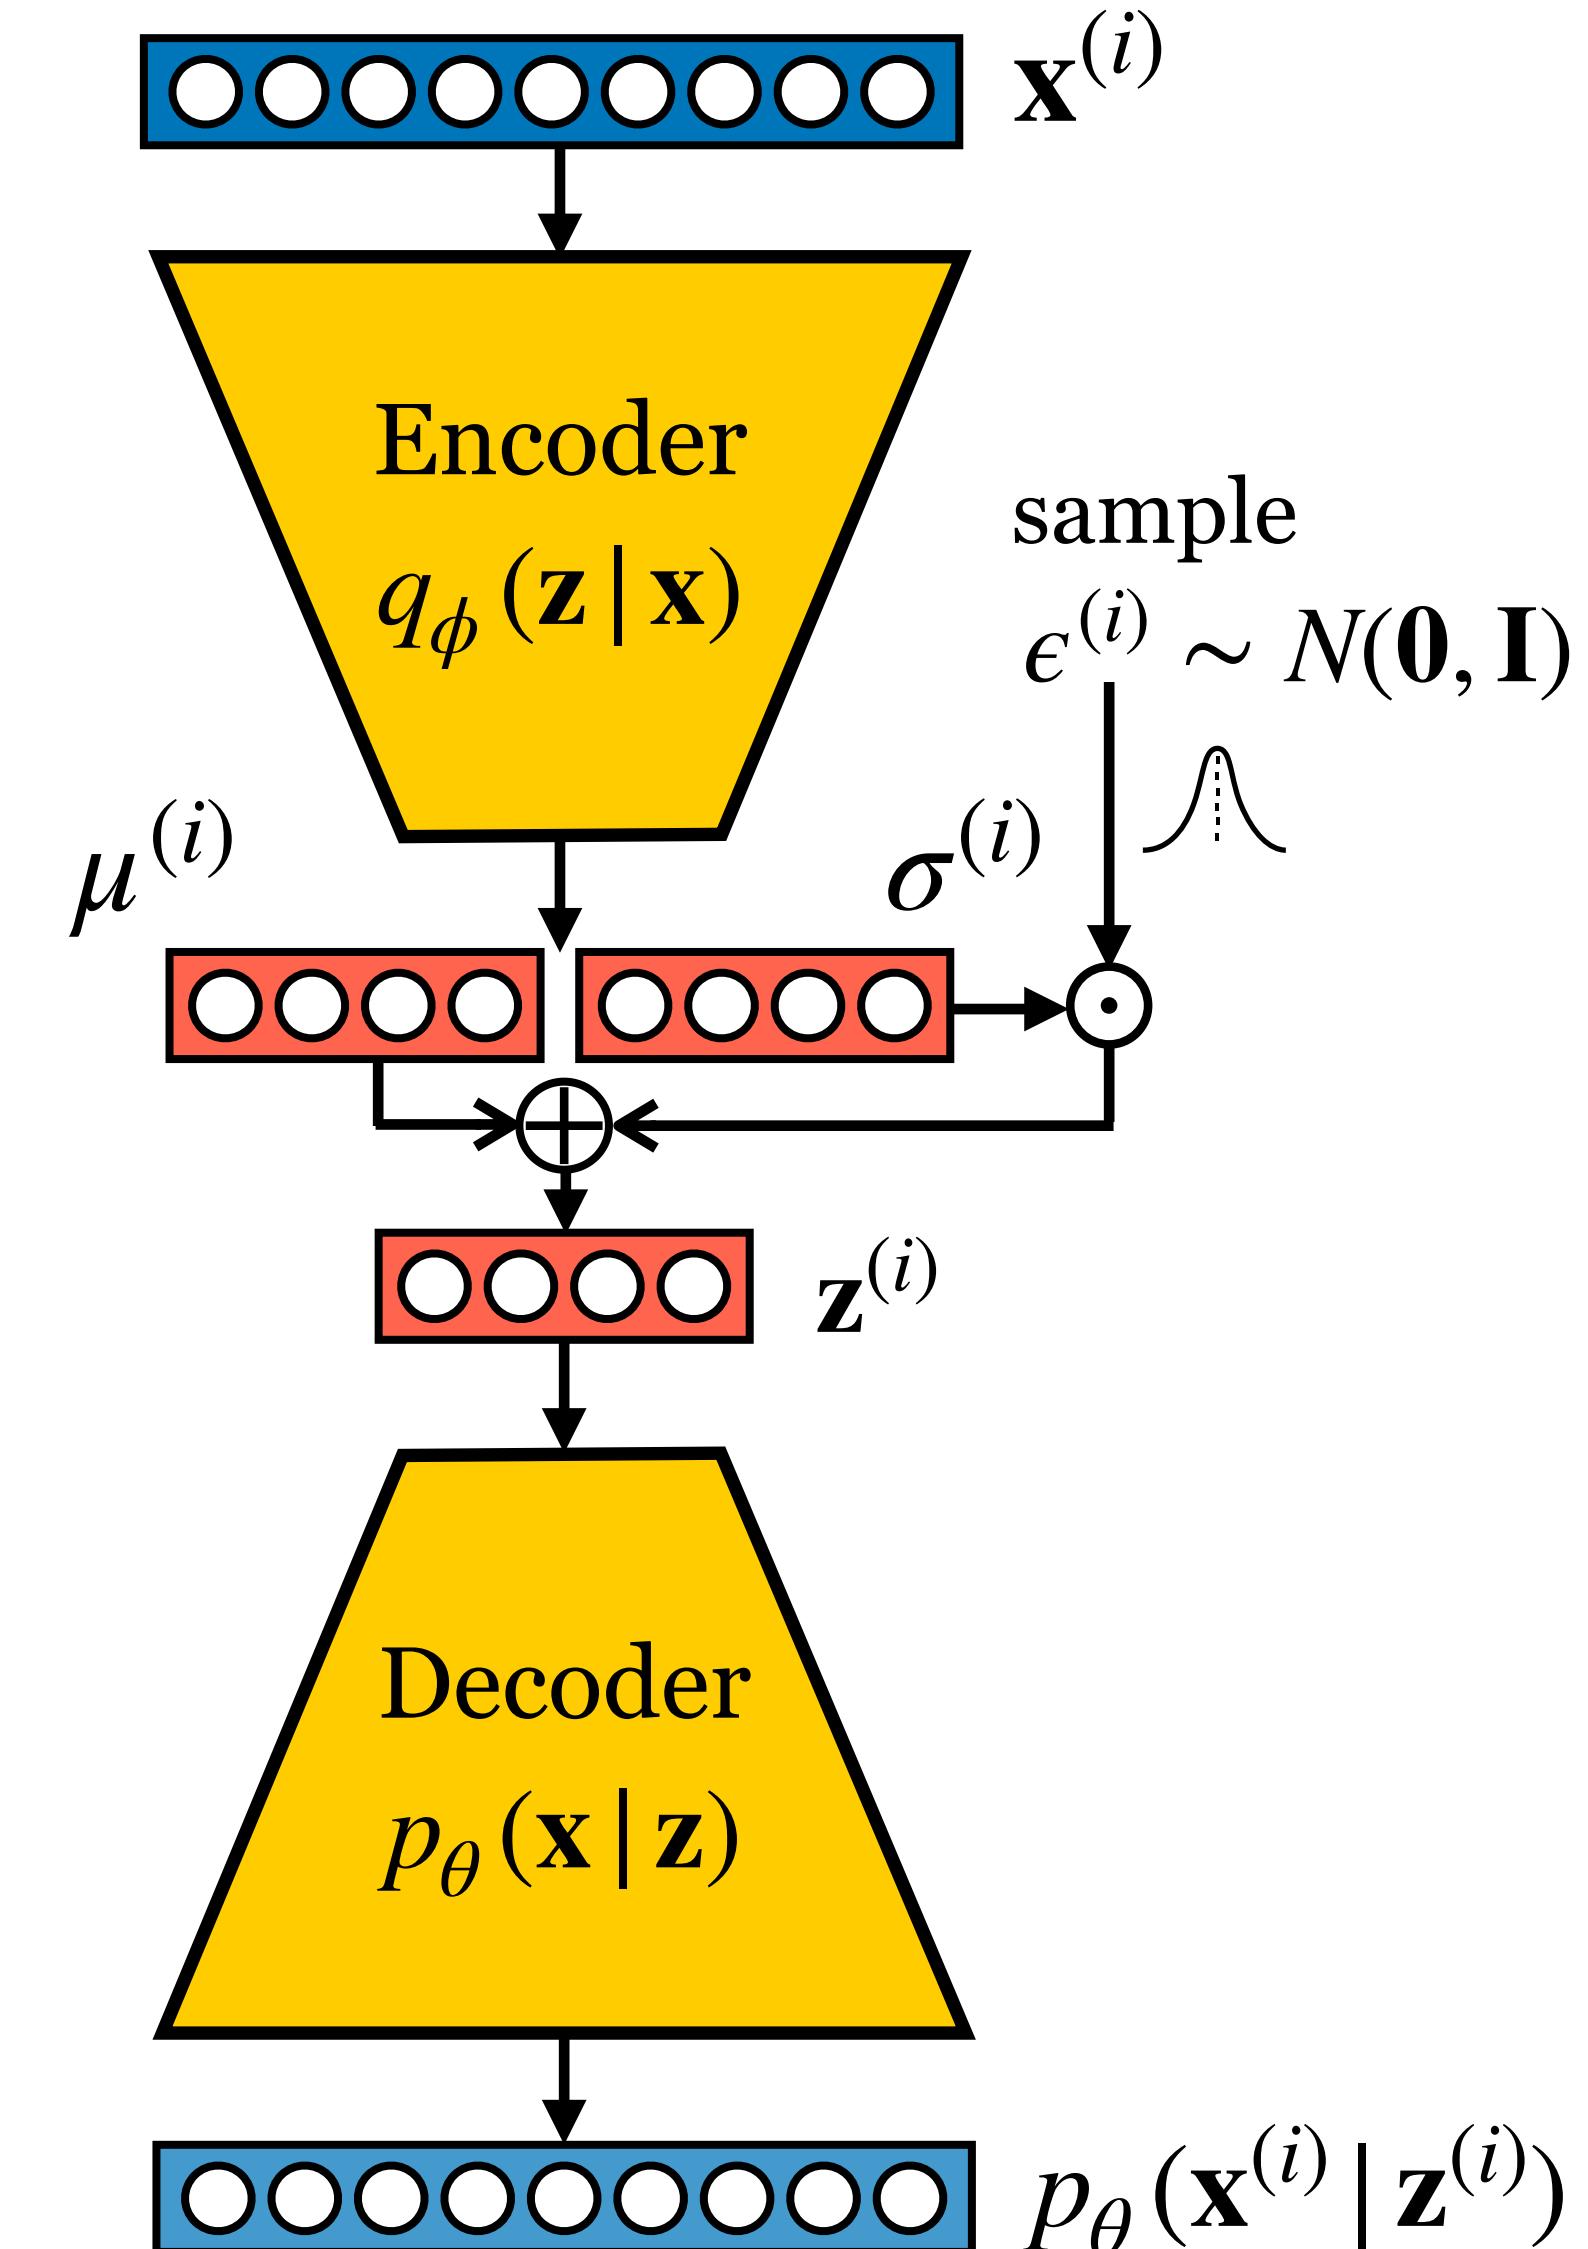
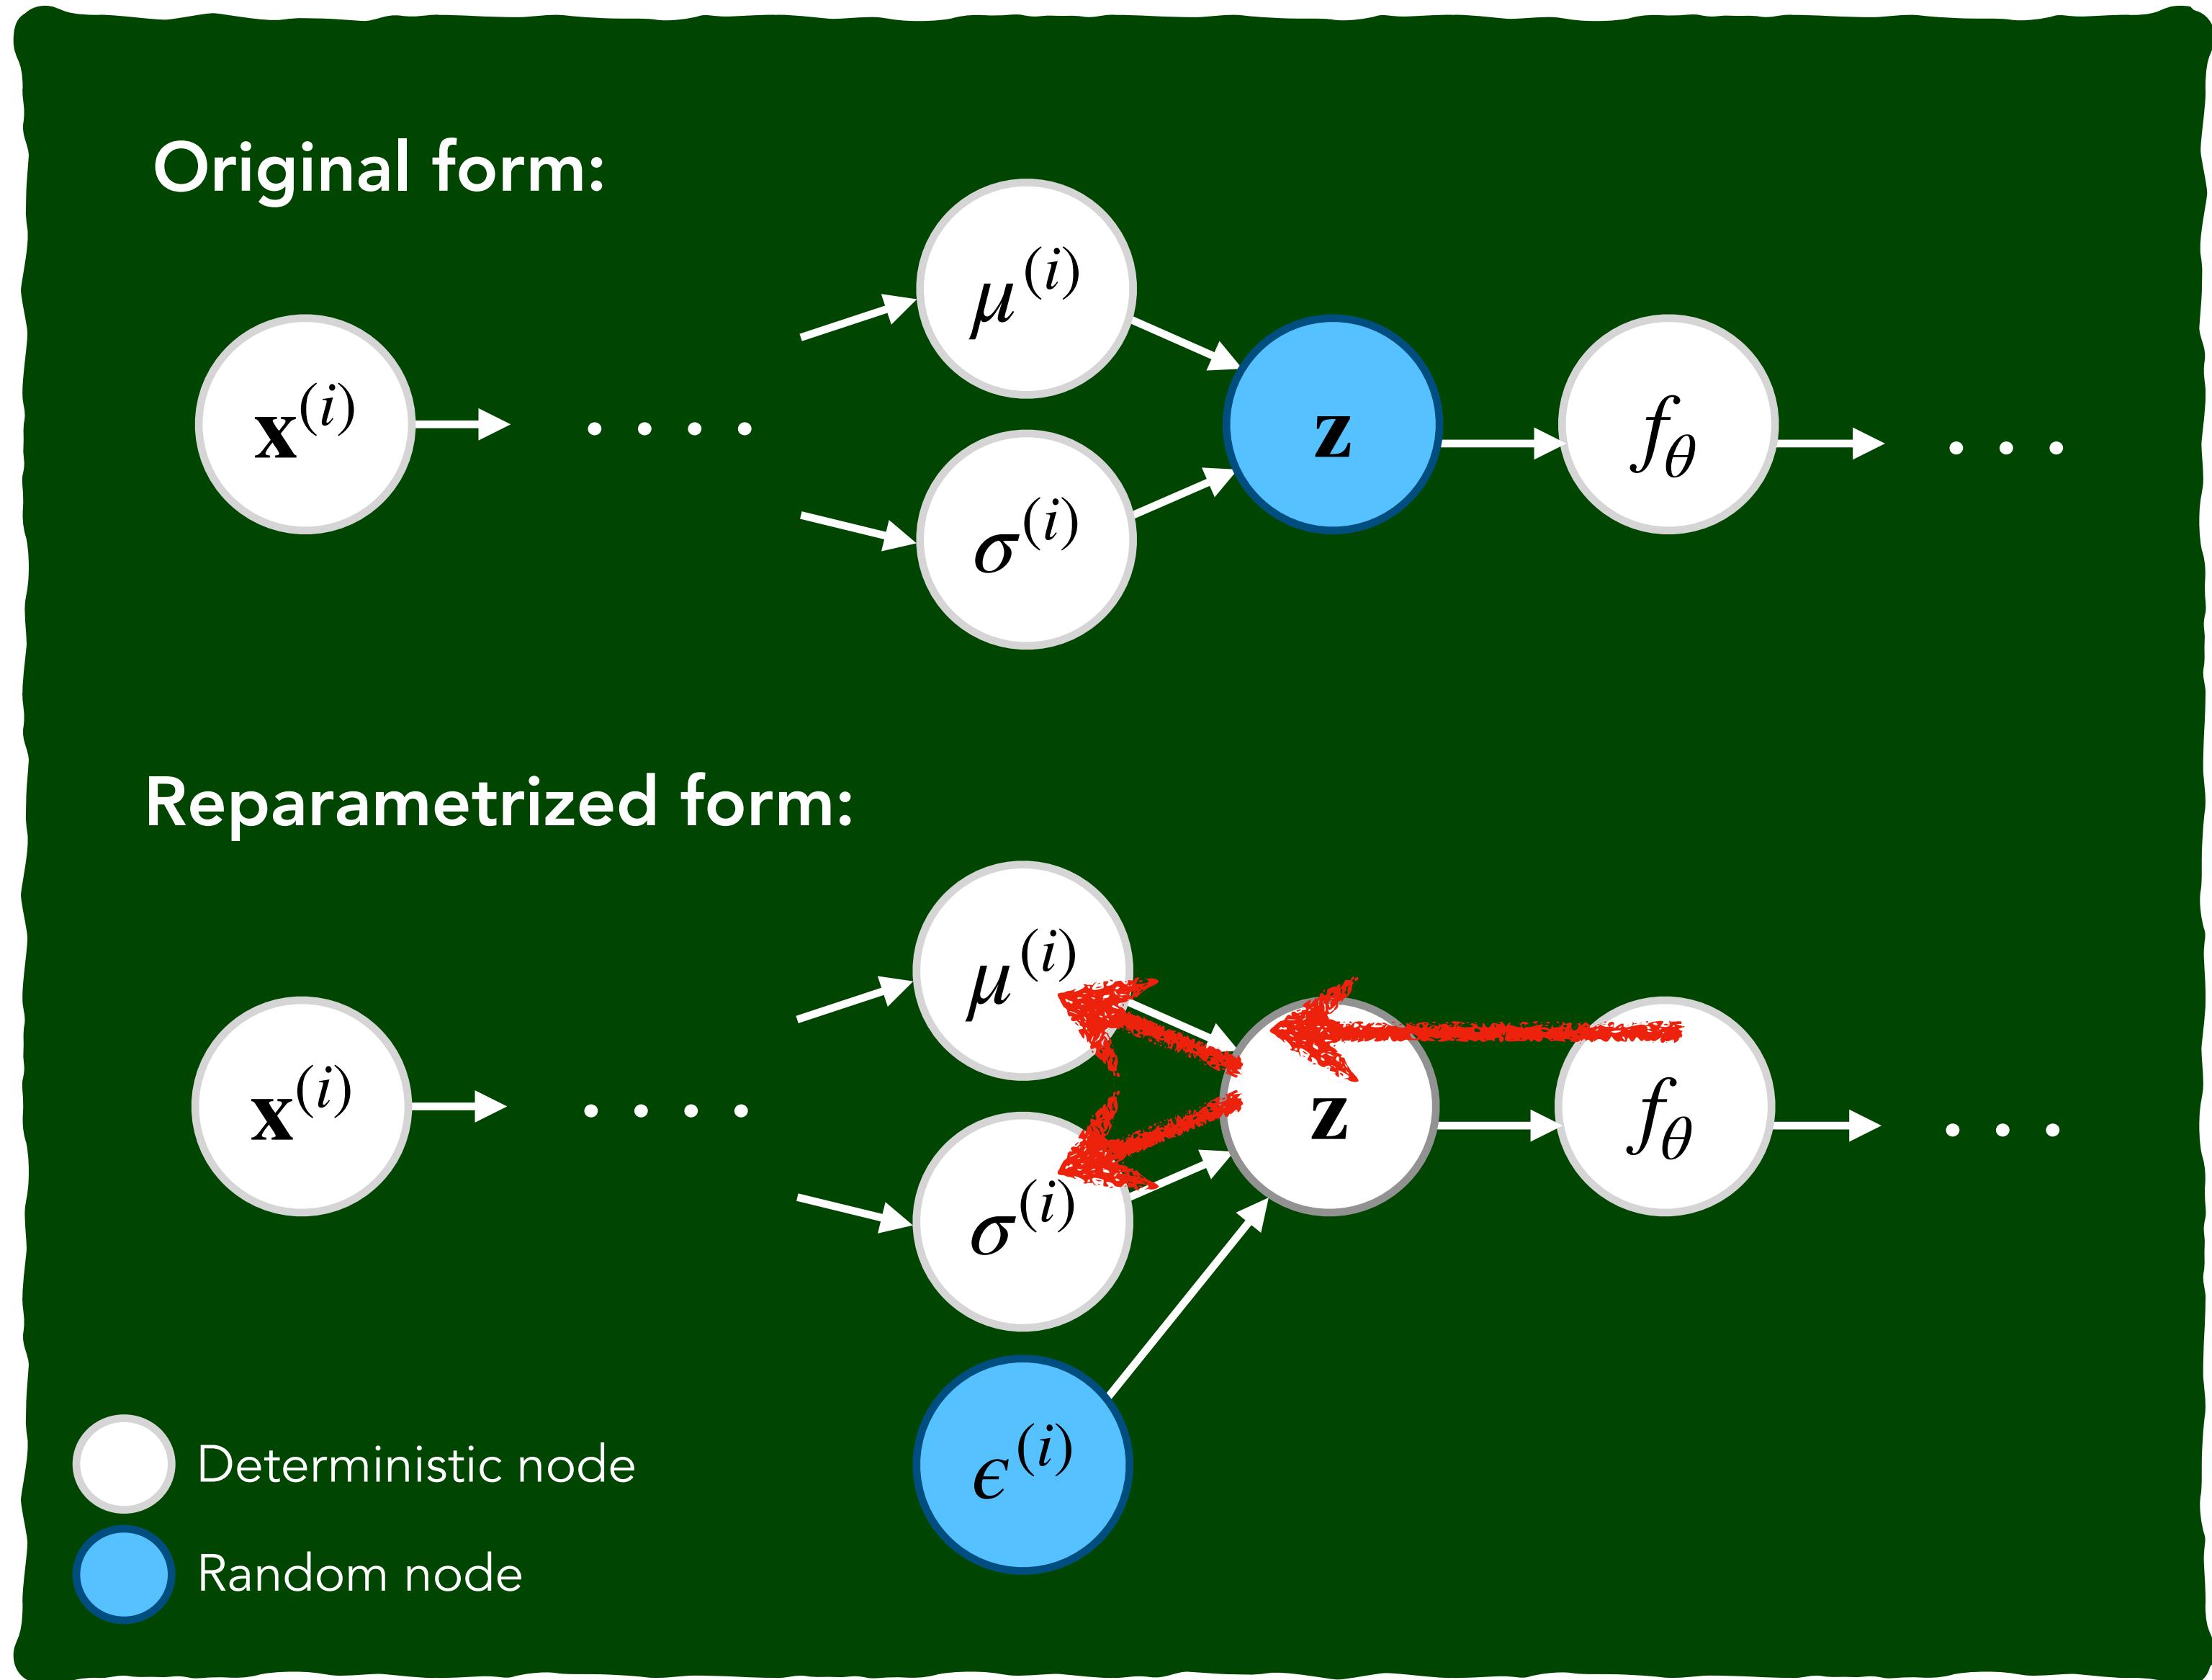
- For many distributions (e.g., for independent Gaussians) we can obtain the same distribution over  $\mathbf{z}$  by including an auxiliary random variable  $\epsilon$ .
- For univariate Gaussian:

$$z^{(i,l)} = \mu(\mathbf{x}^{(i)}) + \sigma(\mathbf{x}^{(i)}) \epsilon^{(i,l)} \quad \text{with} \quad \epsilon^{(i,l)} \sim N(0, 1)$$

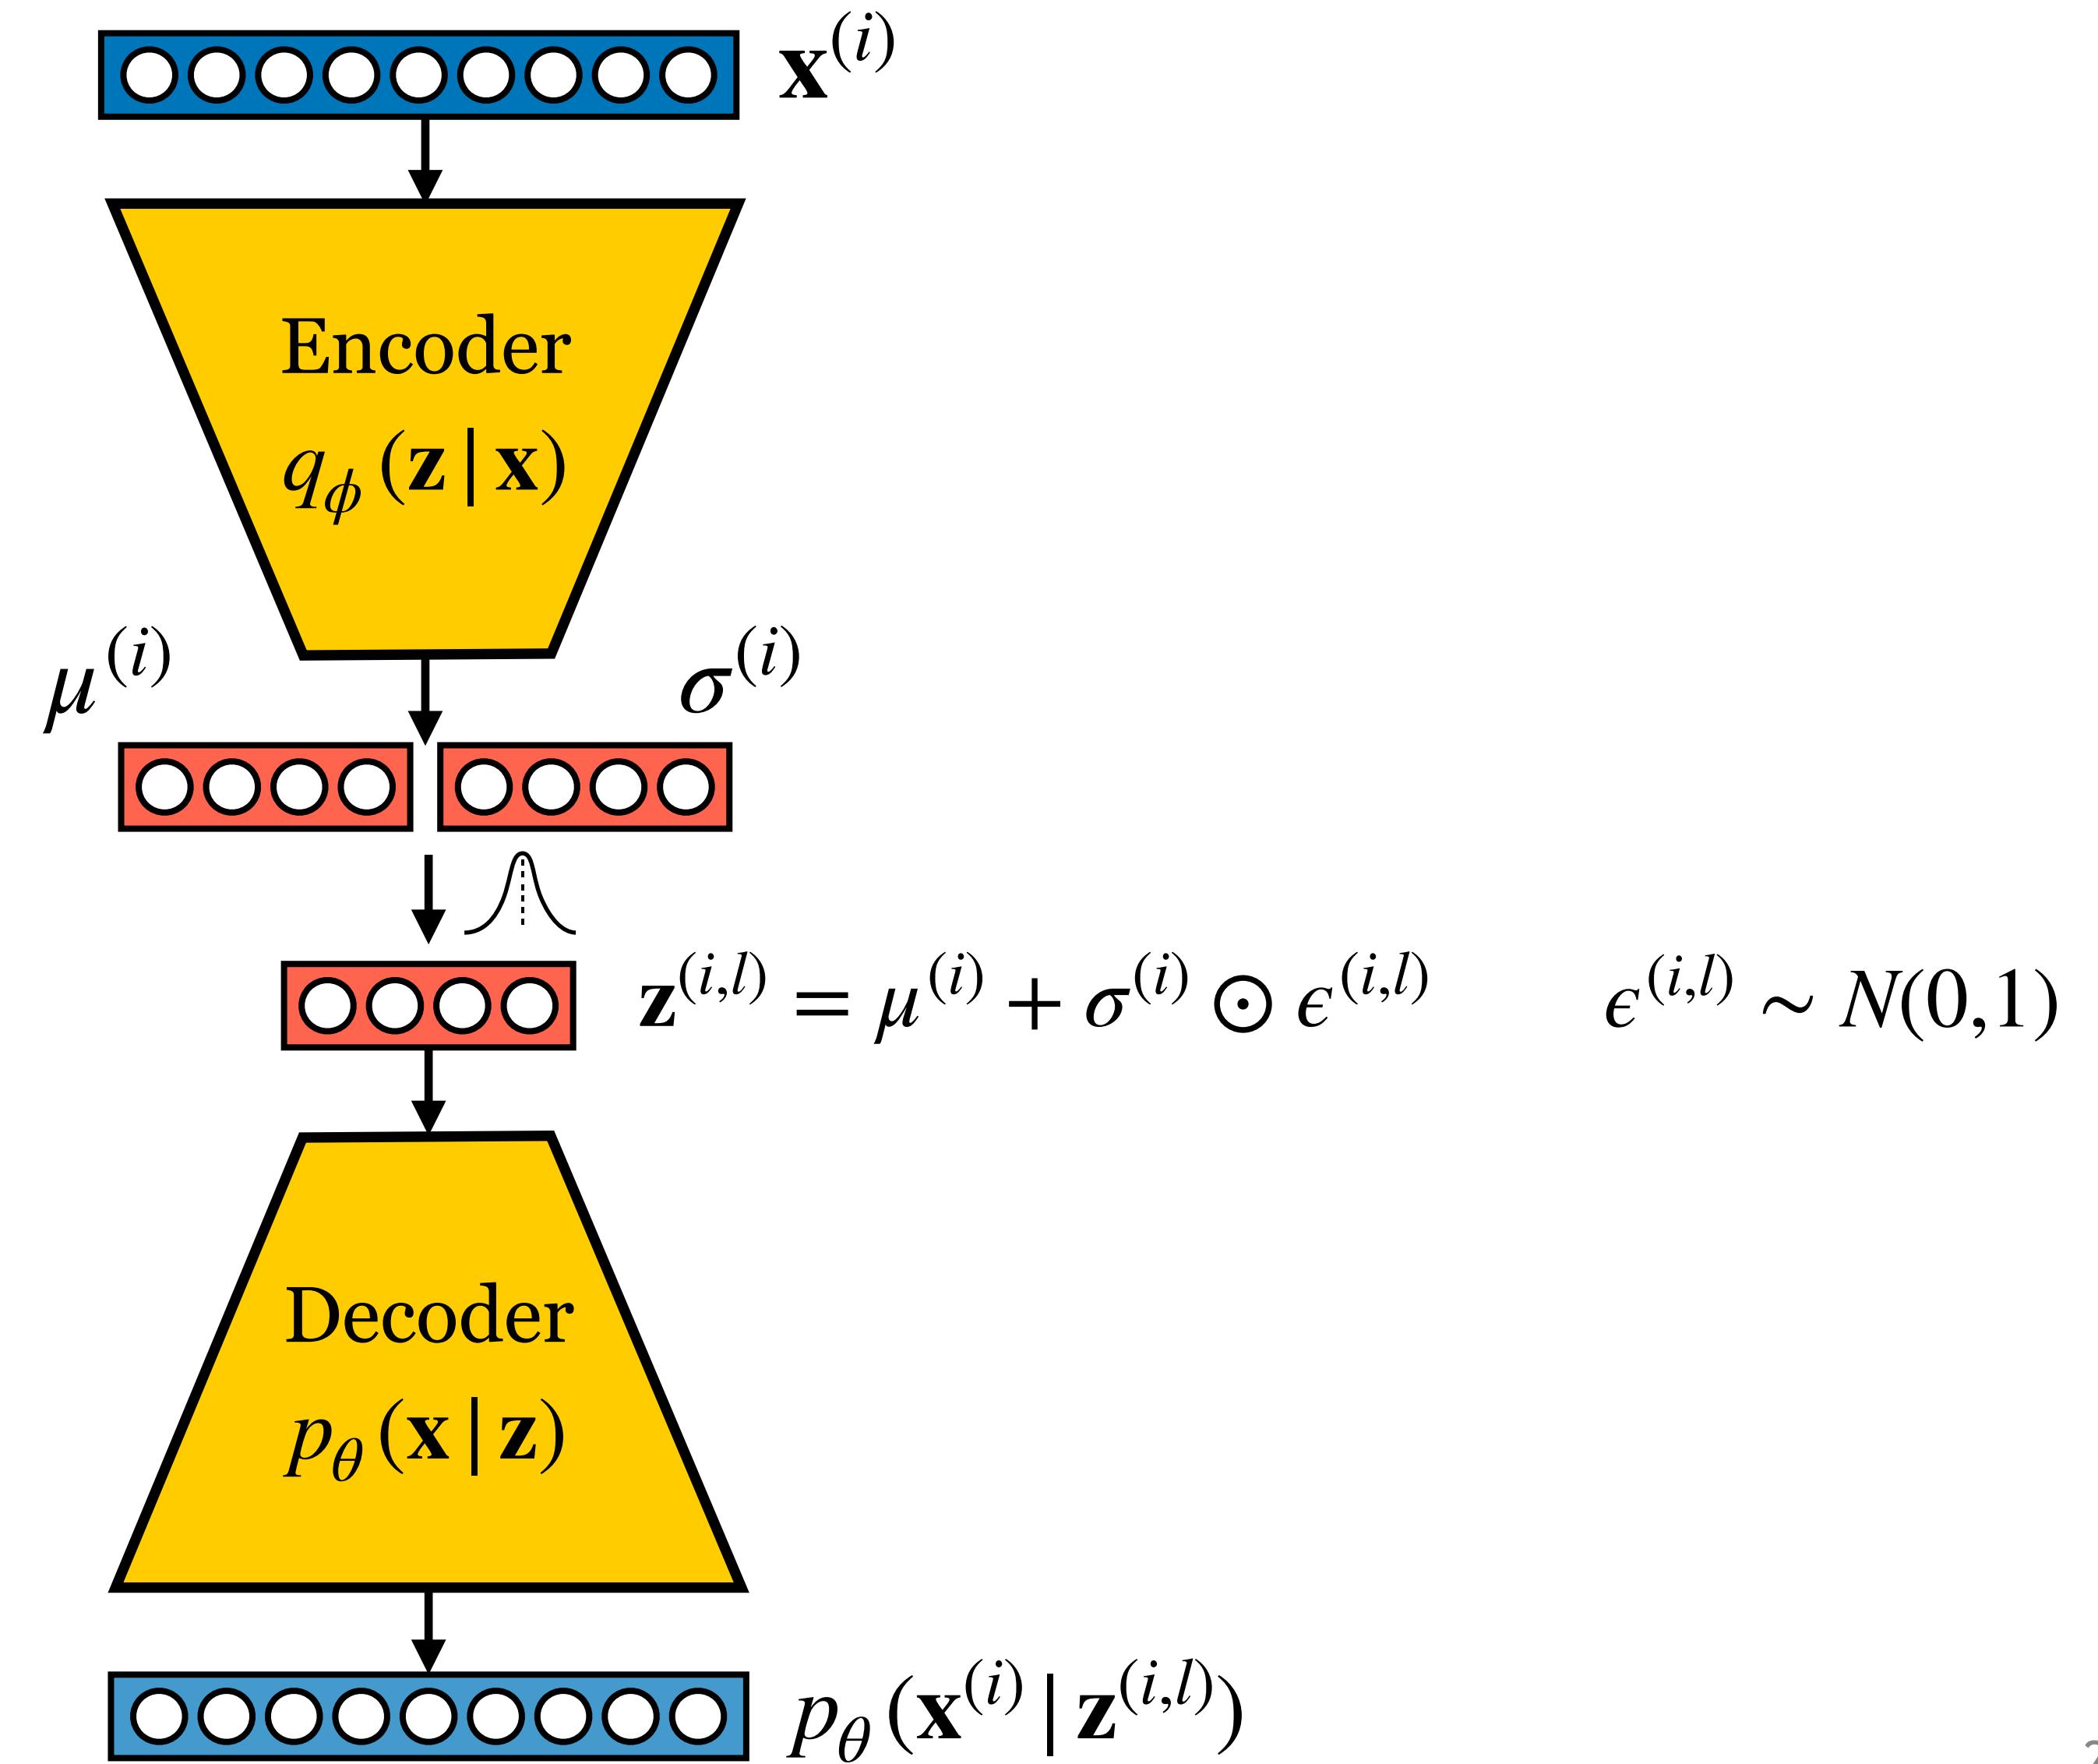
In this case, we get  $\frac{\partial \mathbf{z}}{\partial \mu} = 1$  and  $\frac{\partial \mathbf{z}}{\partial \sigma} = \epsilon$



# Reparametrization Trick



# Variational Autoencoders: Final Model



# Today

---

Autoencoders

Deep Generative Models

Variational Autoencoders

Architecture & Training

Variational Lower Bound

Experiments & Quantifying Performance

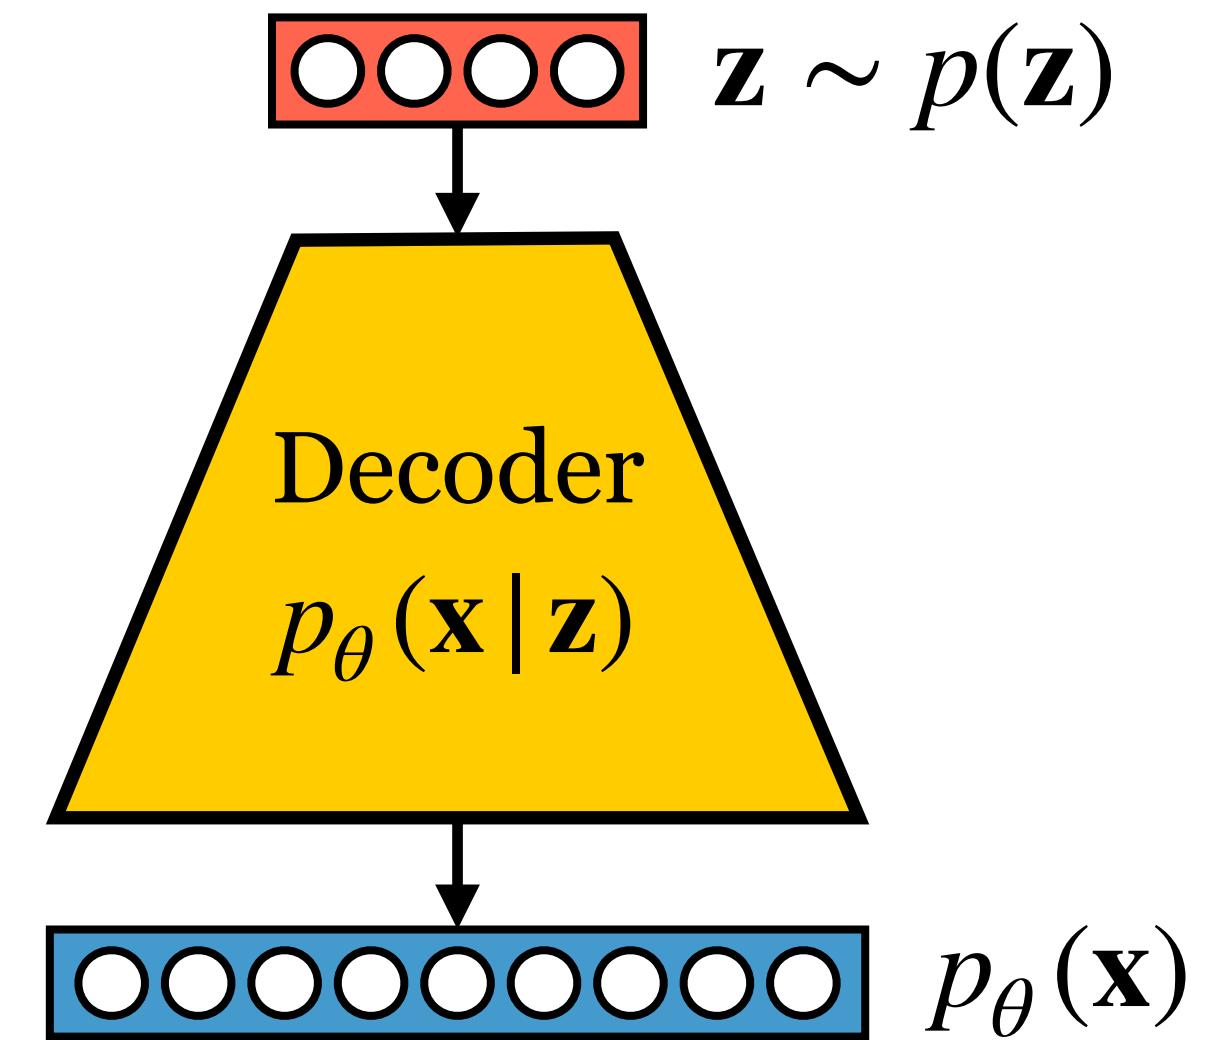
# Variational Lower Bound

Our posterior is only an approximation to  $p_\theta(\mathbf{z} | \mathbf{x})$ .

## Background on variational inference:

The problem to invert a directed graphical model is a general problem.

- Given: A model for  $p_\theta(\mathbf{x} | \mathbf{z})$
- What we want:  $p_\theta(\mathbf{z} | \mathbf{x})$



Suppose we are given an intractable probability distribution  $p$ .

*Variational techniques try to solve an optimization problem over a class of tractable distributions  $Q$  in order to find a  $q \in Q$  that is most similar to  $p$ . We will then query  $q$  (rather than  $p$ ) in order to get an approximate solution.*

# Variational Lower Bound

---

**Justification for our approach:**

- Mainly we want to maximize  $\log p_\theta(\mathbf{x})$ . We can write this log-likelihood as:

$$\log p_\theta(\mathbf{x}) = D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p_\theta(\mathbf{z} \mid \mathbf{x}) \right) + \mathcal{L}(\theta, \phi, \mathbf{x})$$

with  $\mathcal{L}(\theta, \phi, X)$  being the "**variational lower bound**" as the (negated) loss objective from before:

$$\mathcal{L}(\theta, \phi, X) = -D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p(\mathbf{z}) \right) + \mathbb{E}_{q_\phi(\mathbf{z} \mid \mathbf{x})} [\log p_\theta(\mathbf{x} \mid \mathbf{z})]$$

sometimes also called:  
**"evidence lower bound" (ELBO)**

# Variational Lower Bound

$$\log p_\theta(\mathbf{x}) = D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p_\theta(\mathbf{z} \mid \mathbf{x}) \right) + \mathcal{L}(\theta, \phi, \mathbf{x}) \quad \quad \mathcal{L}(\theta, \phi, X) = -D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p(\mathbf{z}) \right) + \mathbb{E}_{q_\phi(\mathbf{z} \mid \mathbf{x})} [\log p_\theta(\mathbf{x} \mid \mathbf{z})]$$

$$D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p_\theta(\mathbf{z} \mid \mathbf{x}) \right) - D_{KL} \left( q_\phi(\mathbf{z} \mid \mathbf{x}) \parallel p(\mathbf{z}) \right) + \mathbb{E}_{q_\phi(\mathbf{z} \mid \mathbf{x})} [\log p_\theta(\mathbf{x} \mid \mathbf{z})]$$

$$= \int q_\phi(\mathbf{z} \mid \mathbf{x}) \log \frac{q_\phi(\mathbf{z} \mid \mathbf{x})}{p_\theta(\mathbf{z} \mid \mathbf{x})} dz - \int q_\phi(\mathbf{z} \mid \mathbf{x}) \log \frac{q_\phi(\mathbf{z} \mid \mathbf{x})}{p(\mathbf{z})} dz + \int q_\phi(\mathbf{z} \mid \mathbf{x}) \log p_\theta(\mathbf{x} \mid \mathbf{z}) dz$$

$$= \int q_\phi(\mathbf{z} \mid \mathbf{x}) \left[ \log \cancel{q_\phi(\mathbf{z} \mid \mathbf{x})} - \log p_\theta(\mathbf{z} \mid \mathbf{x}) - \log \cancel{q_\phi(\mathbf{z} \mid \mathbf{x})} + \log p(\mathbf{z}) + \log p_\theta(\mathbf{x} \mid \mathbf{z}) \right] dz$$

$$= \int q_\phi(\mathbf{z} \mid \mathbf{x}) \left[ \log \frac{p_\theta(\mathbf{x} \mid \mathbf{z}) p(\mathbf{z})}{p_\theta(\mathbf{z} \mid \mathbf{x})} \right] dz = \int q_\phi(\mathbf{z} \mid \mathbf{x}) \log p_\theta(\mathbf{x}) dz = \log p_\theta(\mathbf{x}) \int q_\phi(\mathbf{z} \mid \mathbf{x}) dz = \log p_\theta(\mathbf{x}) \quad \square$$

# Variational Lower Bound

**Justification for our approach:**

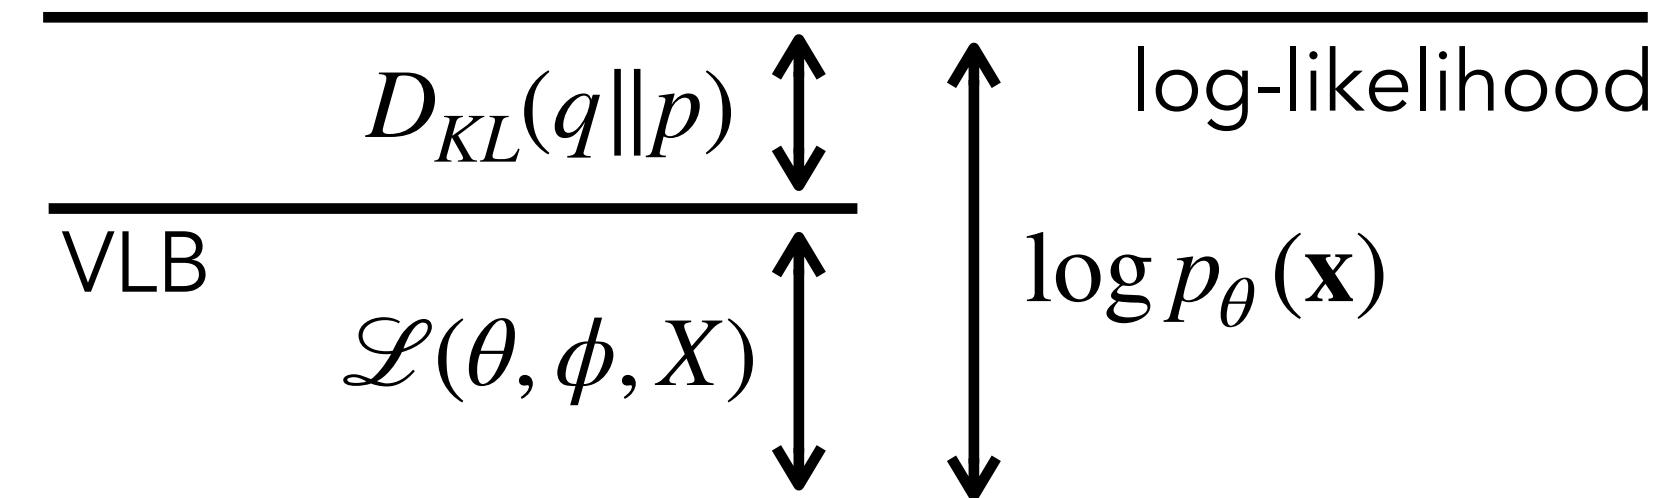
- Mainly we want to maximize  $\log p_\theta(\mathbf{x})$ . We can write this log-likelihood as:

$$\log p_\theta(\mathbf{x}) = D_{KL} \left( q_\phi(\mathbf{z} | \mathbf{x}) \| p_\theta(\mathbf{z} | \mathbf{x}) \right) + \mathcal{L}(\theta, \phi, \mathbf{x})$$

with  $\mathcal{L}(\theta, \phi, X)$  being the "**variational lower bound**" as the (negated) loss objective from before:

$$\mathcal{L}(\theta, \phi, X) = -D_{KL} \left( q_\phi(\mathbf{z} | \mathbf{x}) \| p(\mathbf{z}) \right) + \mathbb{E}_{q_\phi(\mathbf{z} | \mathbf{x})} [\log p_\theta(\mathbf{x} | \mathbf{z})]$$

- By optimizing this, we actually maximize a variational lower bound on  $\log p_\theta(\mathbf{x})$  because  $D_{KL}(q_\phi(\mathbf{z} | \mathbf{x}) || p_\theta(\mathbf{z} | \mathbf{x}))$  is always non-negative, i.e.,  $\log p_\theta(\mathbf{x}) \geq \mathcal{L}(\theta, \phi, X)$



# Today

---

Autoencoders

Deep Generative Models

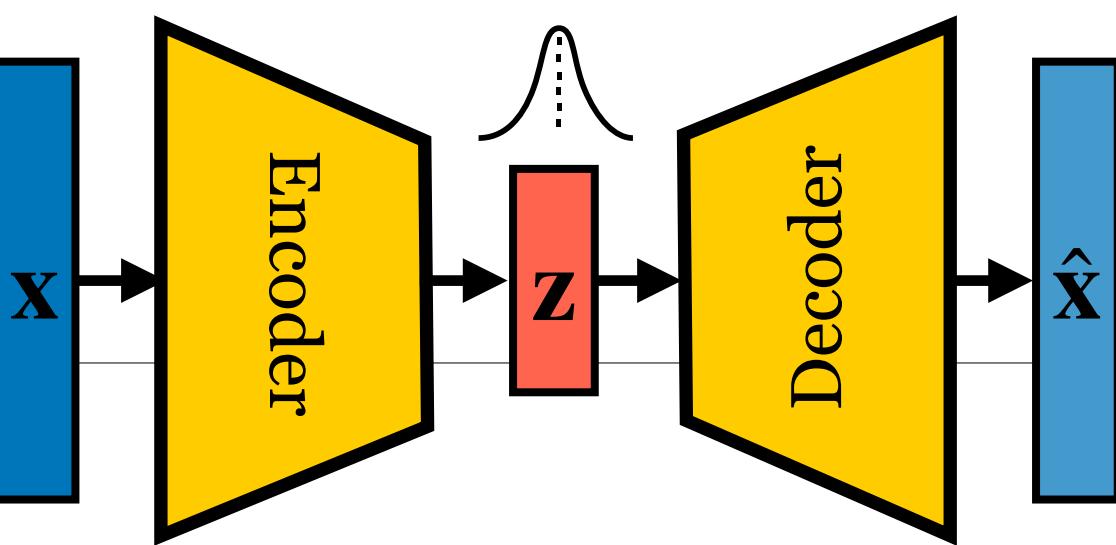
Variational Autoencoders

Architecture & Training

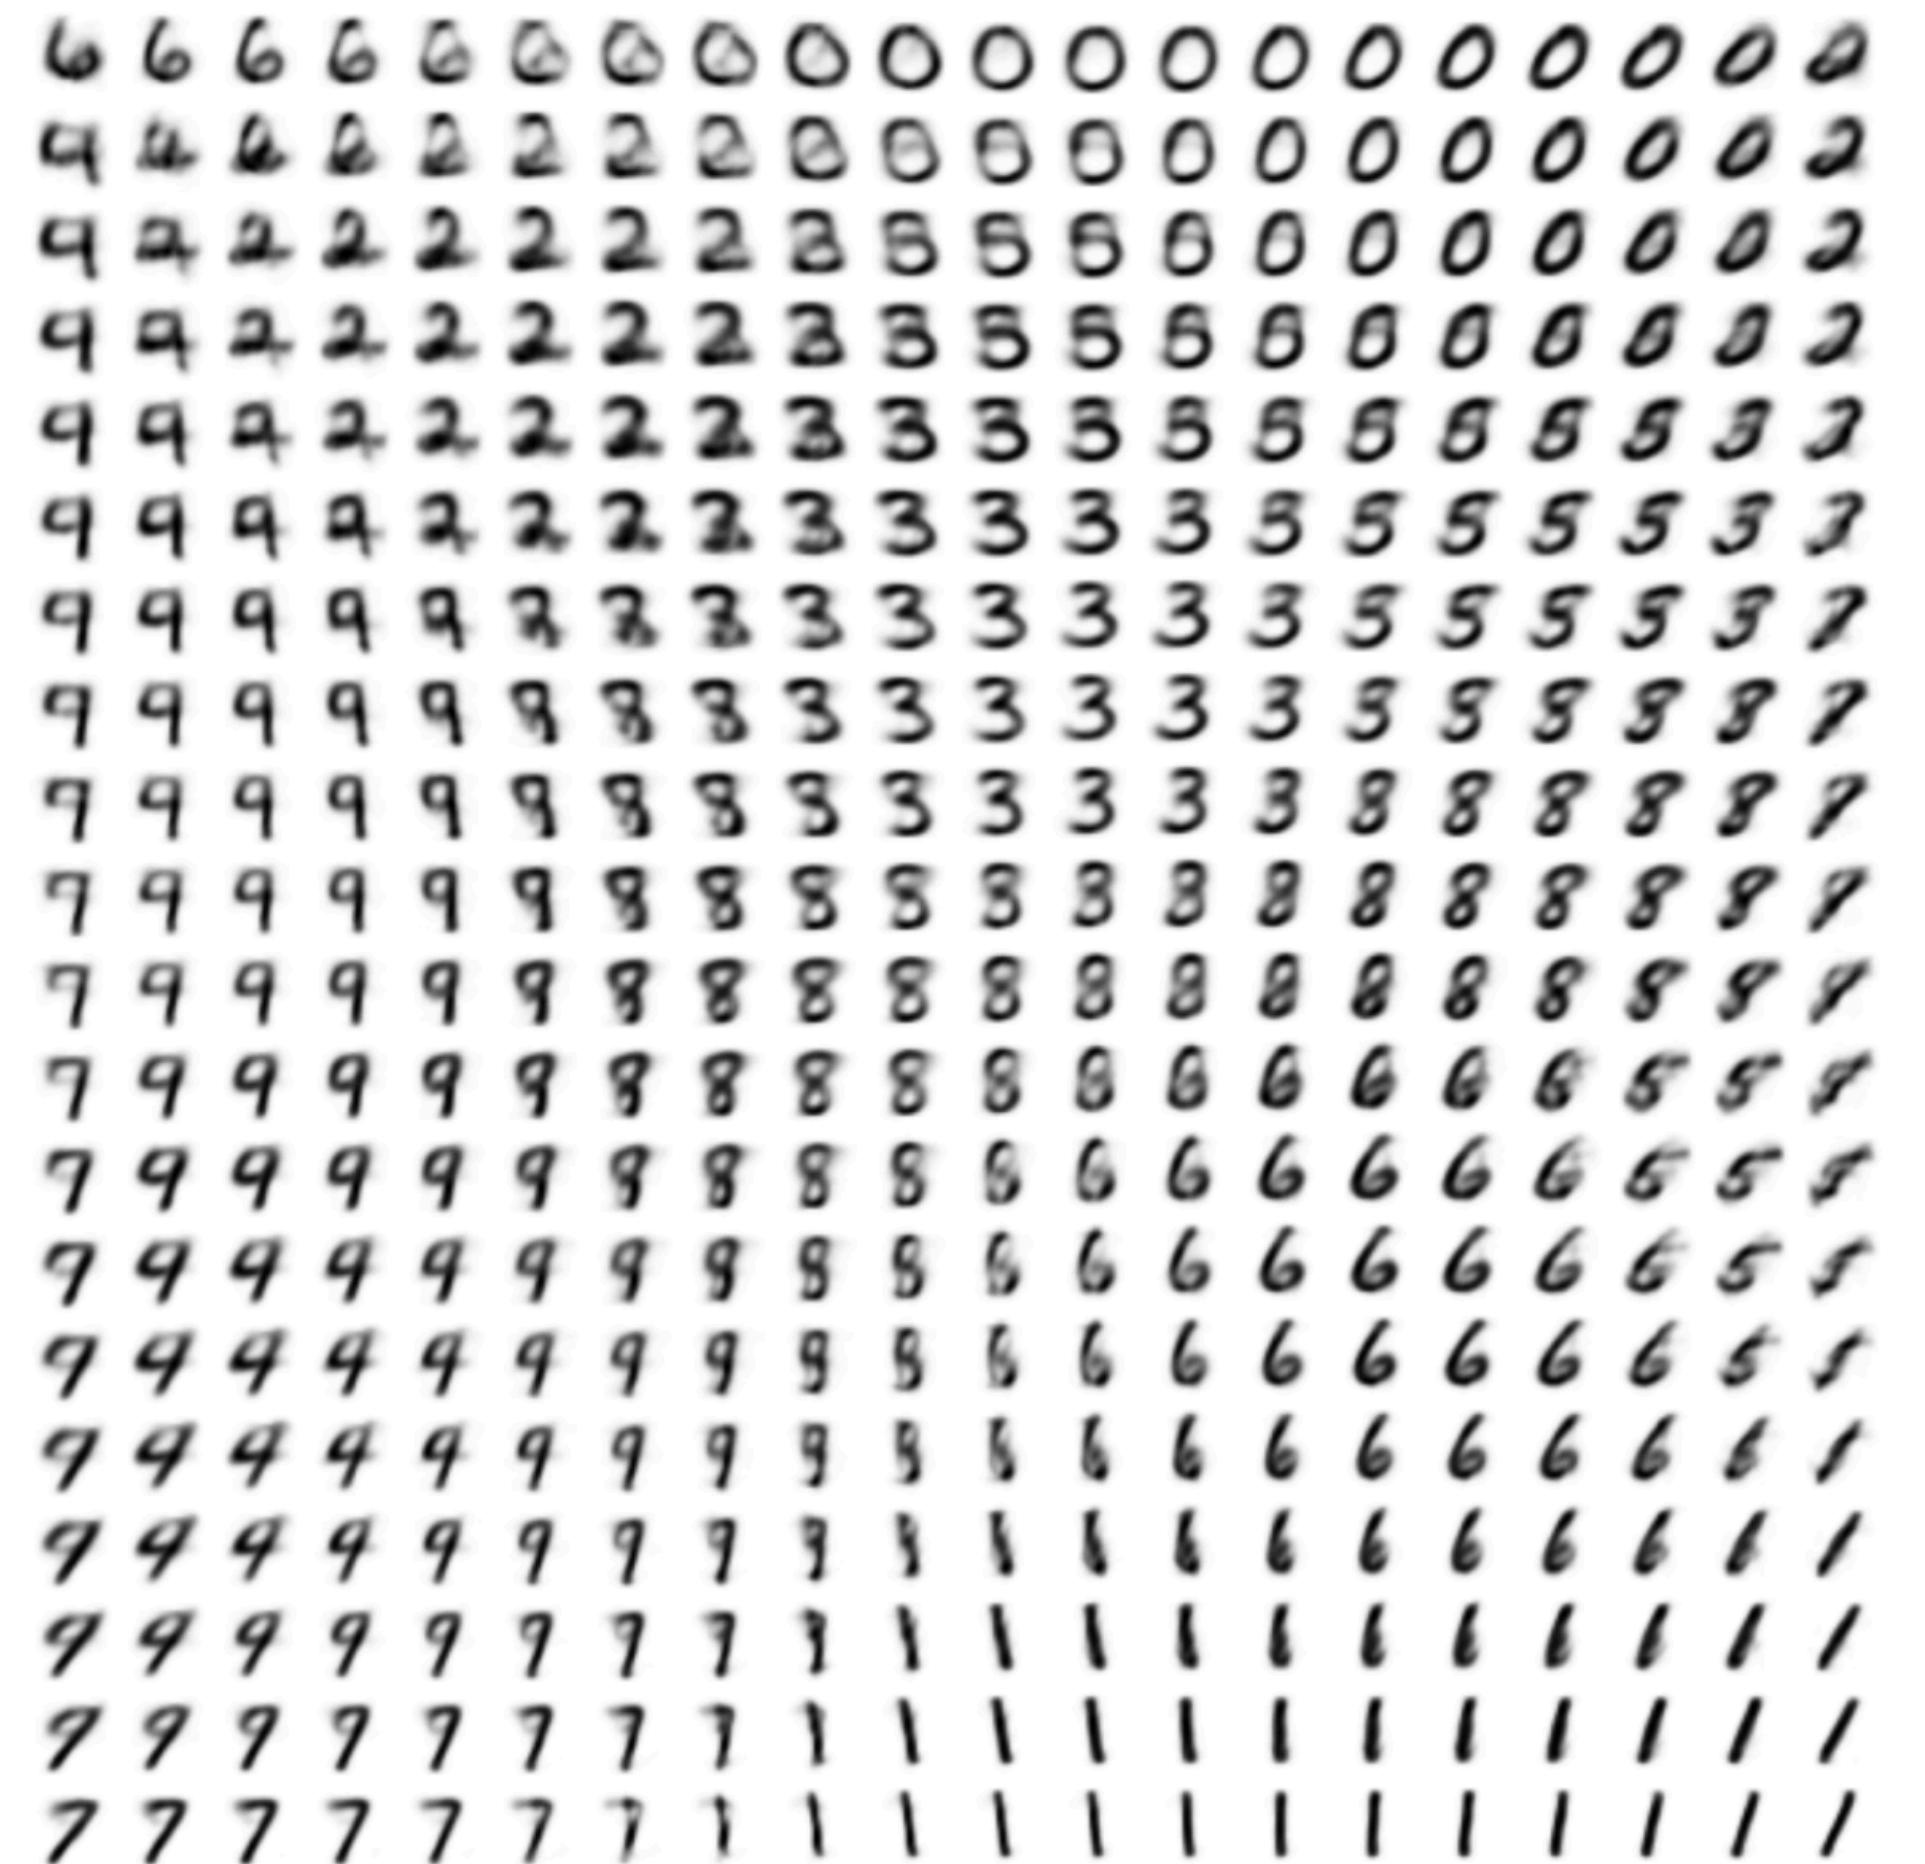
Variational Lower Bound

Experiments & Quantifying Performance

# Experiments

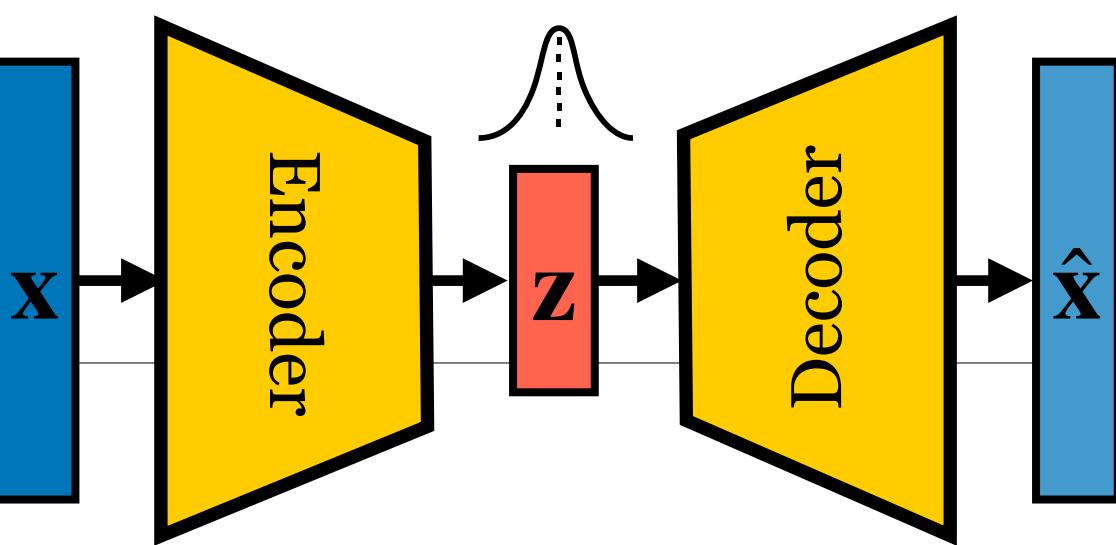


Latent manifold representation for two-dimensional latent space:



[Diederik P. Kingma & Max Welling,  
"Auto-Encoding Variational Bayes",  
arXiv preprint arXiv:1312.6114 (2013).]

# Experiments



Samples from generative model trained on MNIST:

8 6 1 7 8 1 4 8 2 8  
9 6 8 3 9 6 8 3 1 9  
3 3 9 1 3 6 8 1 7 9  
8 9 0 8 6 9 1 9 6 3  
8 2 3 3 3 3 1 3 8 6  
6 9 9 8 6 1 6 6 6 8  
9 5 2 6 6 5 1 8 9 9  
9 9 8 9 3 1 2 8 2 3  
0 4 6 1 2 3 2 0 8 8  
9 7 5 4 9 3 4 8 5 1

(a) 2-D latent space

6 1 6 5 1 0 7 6 7 2  
8 5 9 4 6 8 2 1 6 8  
6 1 0 3 2 2 8 8 1 3 8  
2 8 6 8 9 1 0 0 4 1  
5 1 9 3 0 1 8 3 5 9  
6 5 6 1 4 9 1 7 5 8  
1 3 4 3 9 8 3 2 7 0  
4 5 8 2 9 7 0 4 5 8  
6 9 9 4 8 7 2 8 2 3  
2 6 4 5 6 0 9 9 9 8

(b) 5-D latent space

2 8 3 8 3 0 5 7 3 8  
8 3 8 2 7 9 3 5 3 8  
3 5 9 9 4 3 9 5 1 6  
1 9 8 8 8 3 3 4 9 7  
2 7 3 6 4 3 0 2 0 3  
5 9 7 0 5 8 3 3 4 5  
6 9 4 3 6 2 8 5 7 2  
8 4 9 0 8 0 7 0 6 6  
7 4 3 6 3 0 3 6 0 1  
2 1 8 0 4 7 1 0 0 0

(c) 10-D latent space

8 2 0 8 9 2 3 9 0 0  
7 5 9 9 1 1 7 1 4 4  
8 9 6 2 0 8 2 8 2 9  
2 9 8 6 3 8 7 0 6 1  
5 4 7 9 8 9 8 9 1 0  
6 8 2 4 3 8 8 2 8 1  
2 5 8 2 4 6 1 3 8 8  
7 9 3 9 2 7 9 3 9 6  
4 5 2 4 3 9 0 1 8 4  
8 8 7 3 5 1 6 2 3 6

(d) 20-D latent space

[Diederik P. Kingma & Max Welling,  
"Auto-Encoding Variational Bayes",  
arXiv preprint arXiv:1312.6114 (2013).]

# Quantifying Performance

---

## **Quantitative measures for success of generative models.**

- **Test likelihood:** Measure their log-likelihood for a test dataset.
  - Not always possible to compute exactly and efficiently (e.g., not possible for GANs).
  - For VAEs: only possible to compute the lower bound on the log-likelihood.

# Quantifying Performance

---

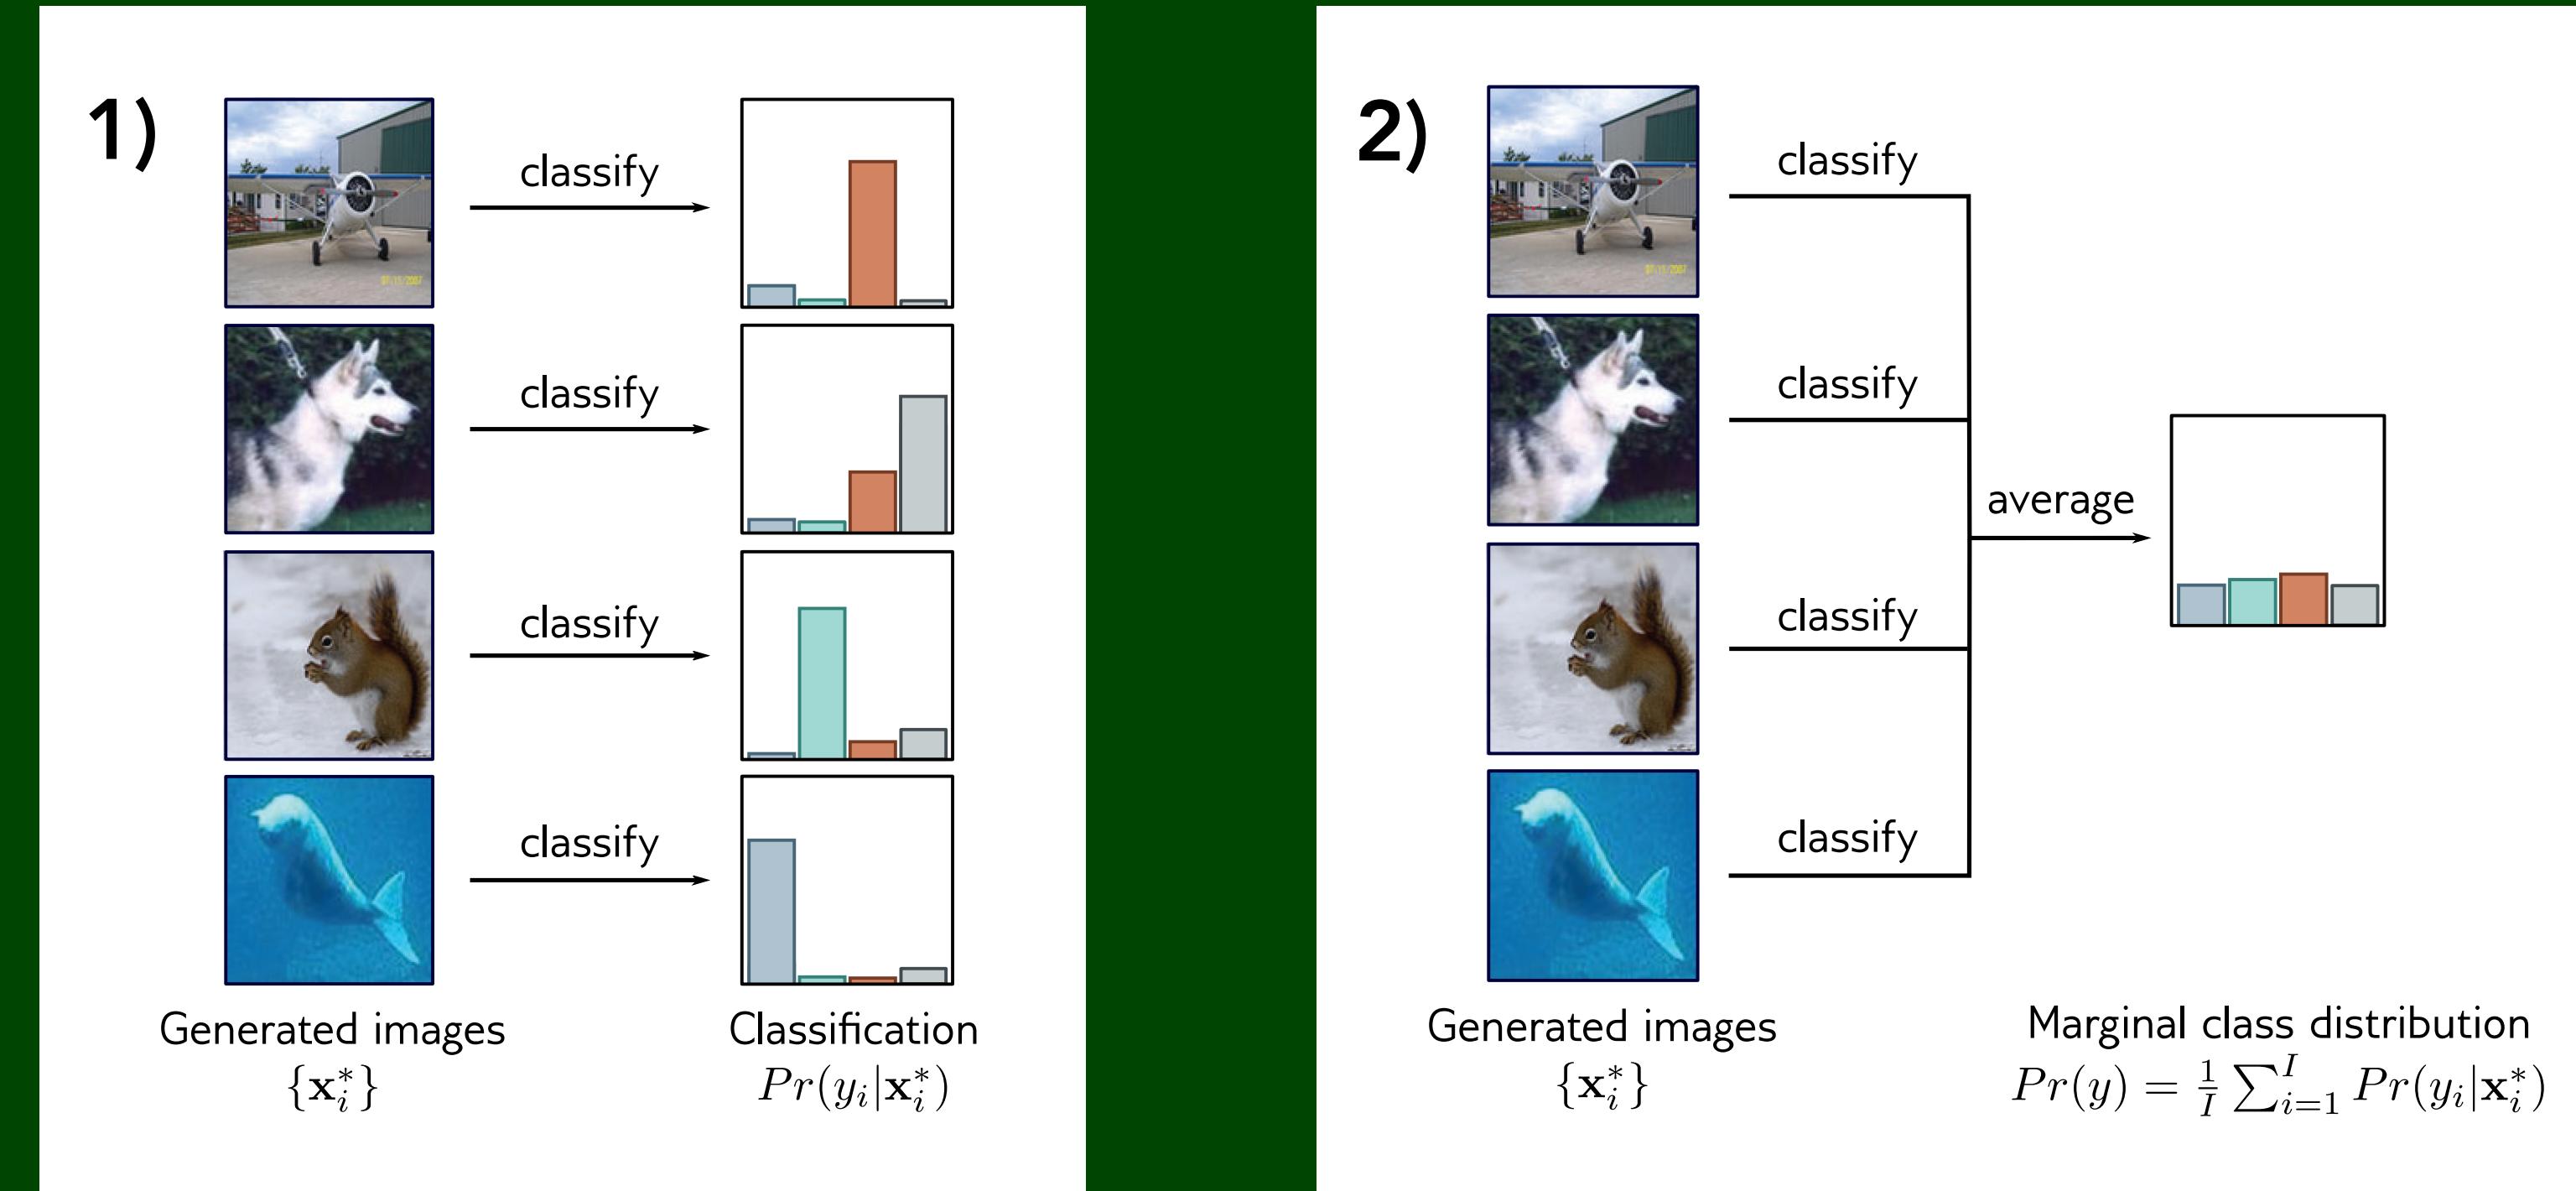
## Quantitative measures for success of generative models.

- **Test likelihood:** Measure their log-likelihood for a test dataset.
  - Not always possible to compute exactly and efficiently (e.g., not possible for GANs).
  - For VAEs: only possible to compute the lower bound on the log-likelihood.
- **Inception Score (IS):** Calculated using a pre-trained CNN, e.g., “Inception-v3”, on ImageNet.
  - Sensitive to the classifier, retraining it can give different results.
    - 1) Samples should look like only one of the 1000 ImageNet classes (sharpness).
    - 2) Model generates samples from all classes equally frequently (diversity).

# Quantifying Performance

## Quantitative measures for samples

- **Test likelihood:** Measure the probability of generated samples under the true distribution.
  - Not always possible to compute
  - For VAEs: only possible to sample from latent space
- **Inception Score (IS):** Calculate the similarity between generated samples and a reference set.
  - Sensitive to the classifier
  - 1) Samples should look like the training data
    - 2) Model generates samples that are diverse



**Inception Score:** 
$$IS = \exp \left( \frac{1}{I} \sum_{i=1}^I D_{KL} \left[ Pr(y_i | \mathbf{x}_i^*) \| Pr(y) \right] \right)$$

# Quantifying Performance

---

## Quantitative measures for success of generative models.

- **Test likelihood:** Measure their log-likelihood for a test dataset.
  - Not always possible to compute exactly and efficiently (e.g., not possible for GANs).
  - For VAEs: only possible to compute the lower bound on the log-likelihood.
- **Inception Score (IS):** Calculated using a pre-trained CNN, e.g., “Inception-v3”, on ImageNet.
  - Sensitive to the classifier, retraining it can give different results.
    - 1) Samples should look like only one of the 1000 ImageNet classes (sharpness).
    - 2) Model generates samples from all classes equally frequently (diversity).
- **Fréchet Inception Distance (FID):** Measures similarities between the feature representations (learned by a pre-trained classifier) for generated samples and a test dataset.
- ...

# Summary

---

- ▶ Variational autoencoders are encoder-decoder architectures.
- ▶ Training is performed by maximizing a variational lower bound on the log-likelihood.
- ▶ VAEs can be used as generative models or as feature extractors.
- ▶ Individual features of the latent space often encode specific parameters of the input.
  - ▶ e.g. facial expression, object rotation, etc.
  - ▶ By changing the latent variable, one can transform images (e.g., make a person smile).
- ▶ Often, encoders are CNNs and decoders use transposed convolutions.
- ▶ Some (indirect) measures to quantify performance exist (e.g., Inception score).
- ▶ VAEs are easier to train, but Generative Adversarial Networks are usually better generators.

[Ian Goodfellow et al. "Generative Adversarial Nets." NIPS 2014.]

# Today

---

Autoencoders

Deep Generative Models

Variational Autoencoders

Architecture & Training

Variational Lower Bound

Experiments & Quantifying Performance

# Questions?