

# Deep Learning: Generative Adversarial Networks

**Ozan Özdenizci**

Institute of Theoretical Computer Science

[ozan.ozdenizci@igi.tugraz.at](mailto:ozan.ozdenizci@igi.tugraz.at)

Deep Learning VO - WS 23/24

Lecture 10 - December 11th, 2023

# Today

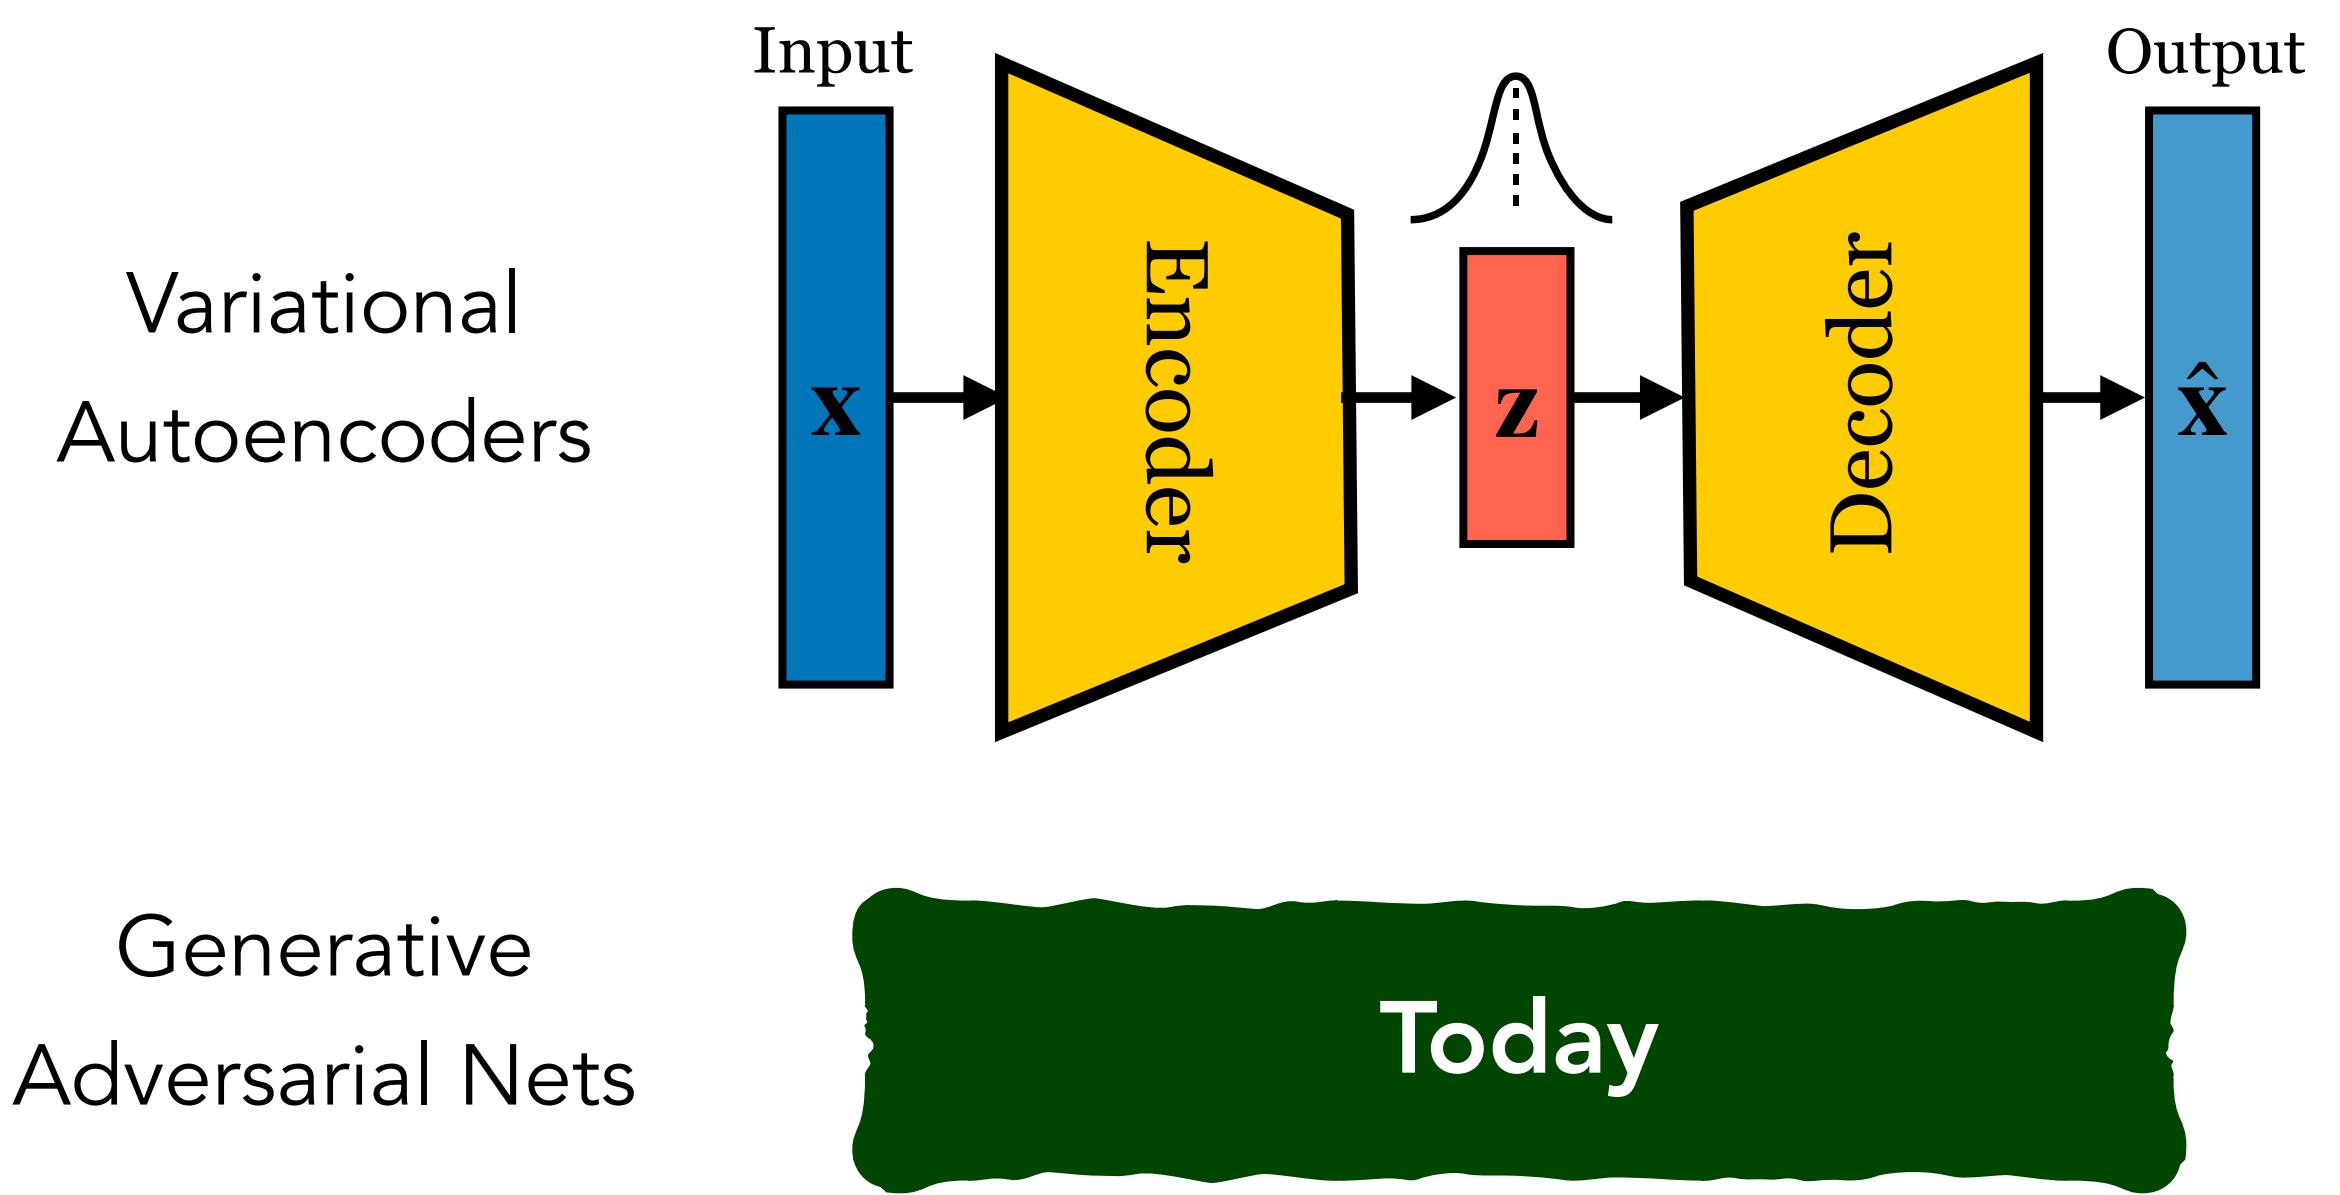
---

- Deep Generative Models
- Variational Autoencoders
- Generative Adversarial Networks
  - Architecture & Training
  - Mode Collapse
  - Further Extensions & Variants

Ian Goodfellow et al., "Generative Adversarial Nets", NIPS 2014.

# Recap: Deep Generative Models

- Given a list of data points  $X = \langle \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \rangle$  coming from a data distribution  $p(\mathbf{x})$ .
- Generative model:** a model  $p_\theta(\mathbf{x})$  for this distribution, from which we can sample.



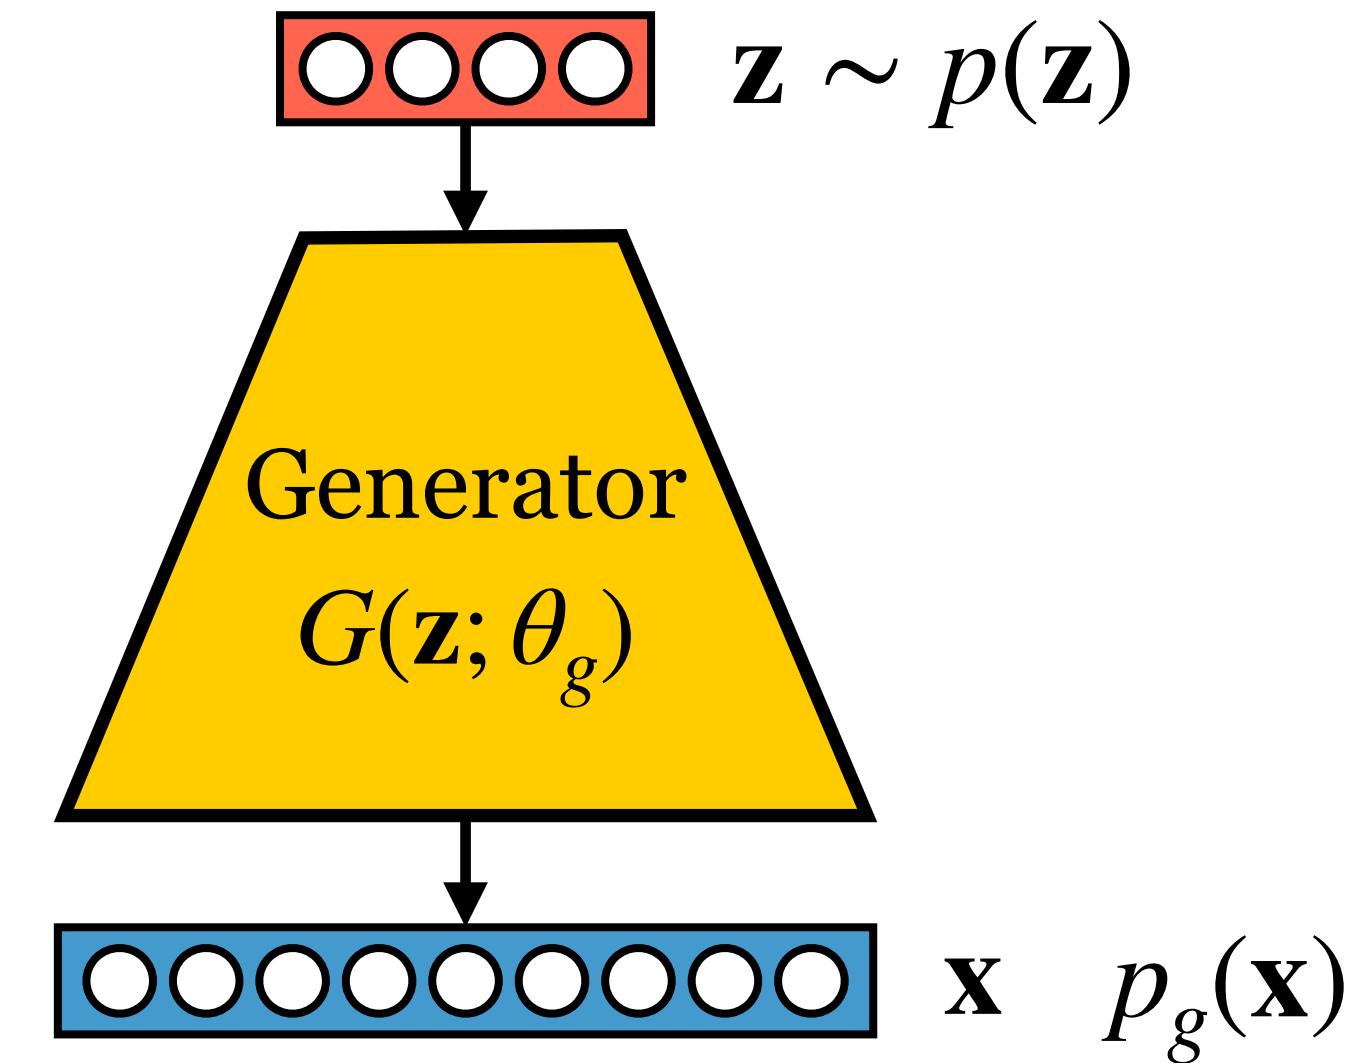
# Generative Stochastic Networks

Uses feed-forward networks to generate samples.

**Generator network:**  $G$

**Network input:** A random vector  $\mathbf{z}$  from some simple prior distribution  $p(\mathbf{z})$ .

**Network output:**  $G(\mathbf{z}; \theta_g)$  : A sample  $\mathbf{x}$  from a complex distribution  $p_g(\mathbf{x})$ .



# Generative Adversarial Networks (GANs)

Add a discriminator that tries to distinguish real data from generated data.

**Generator network:**  $G$

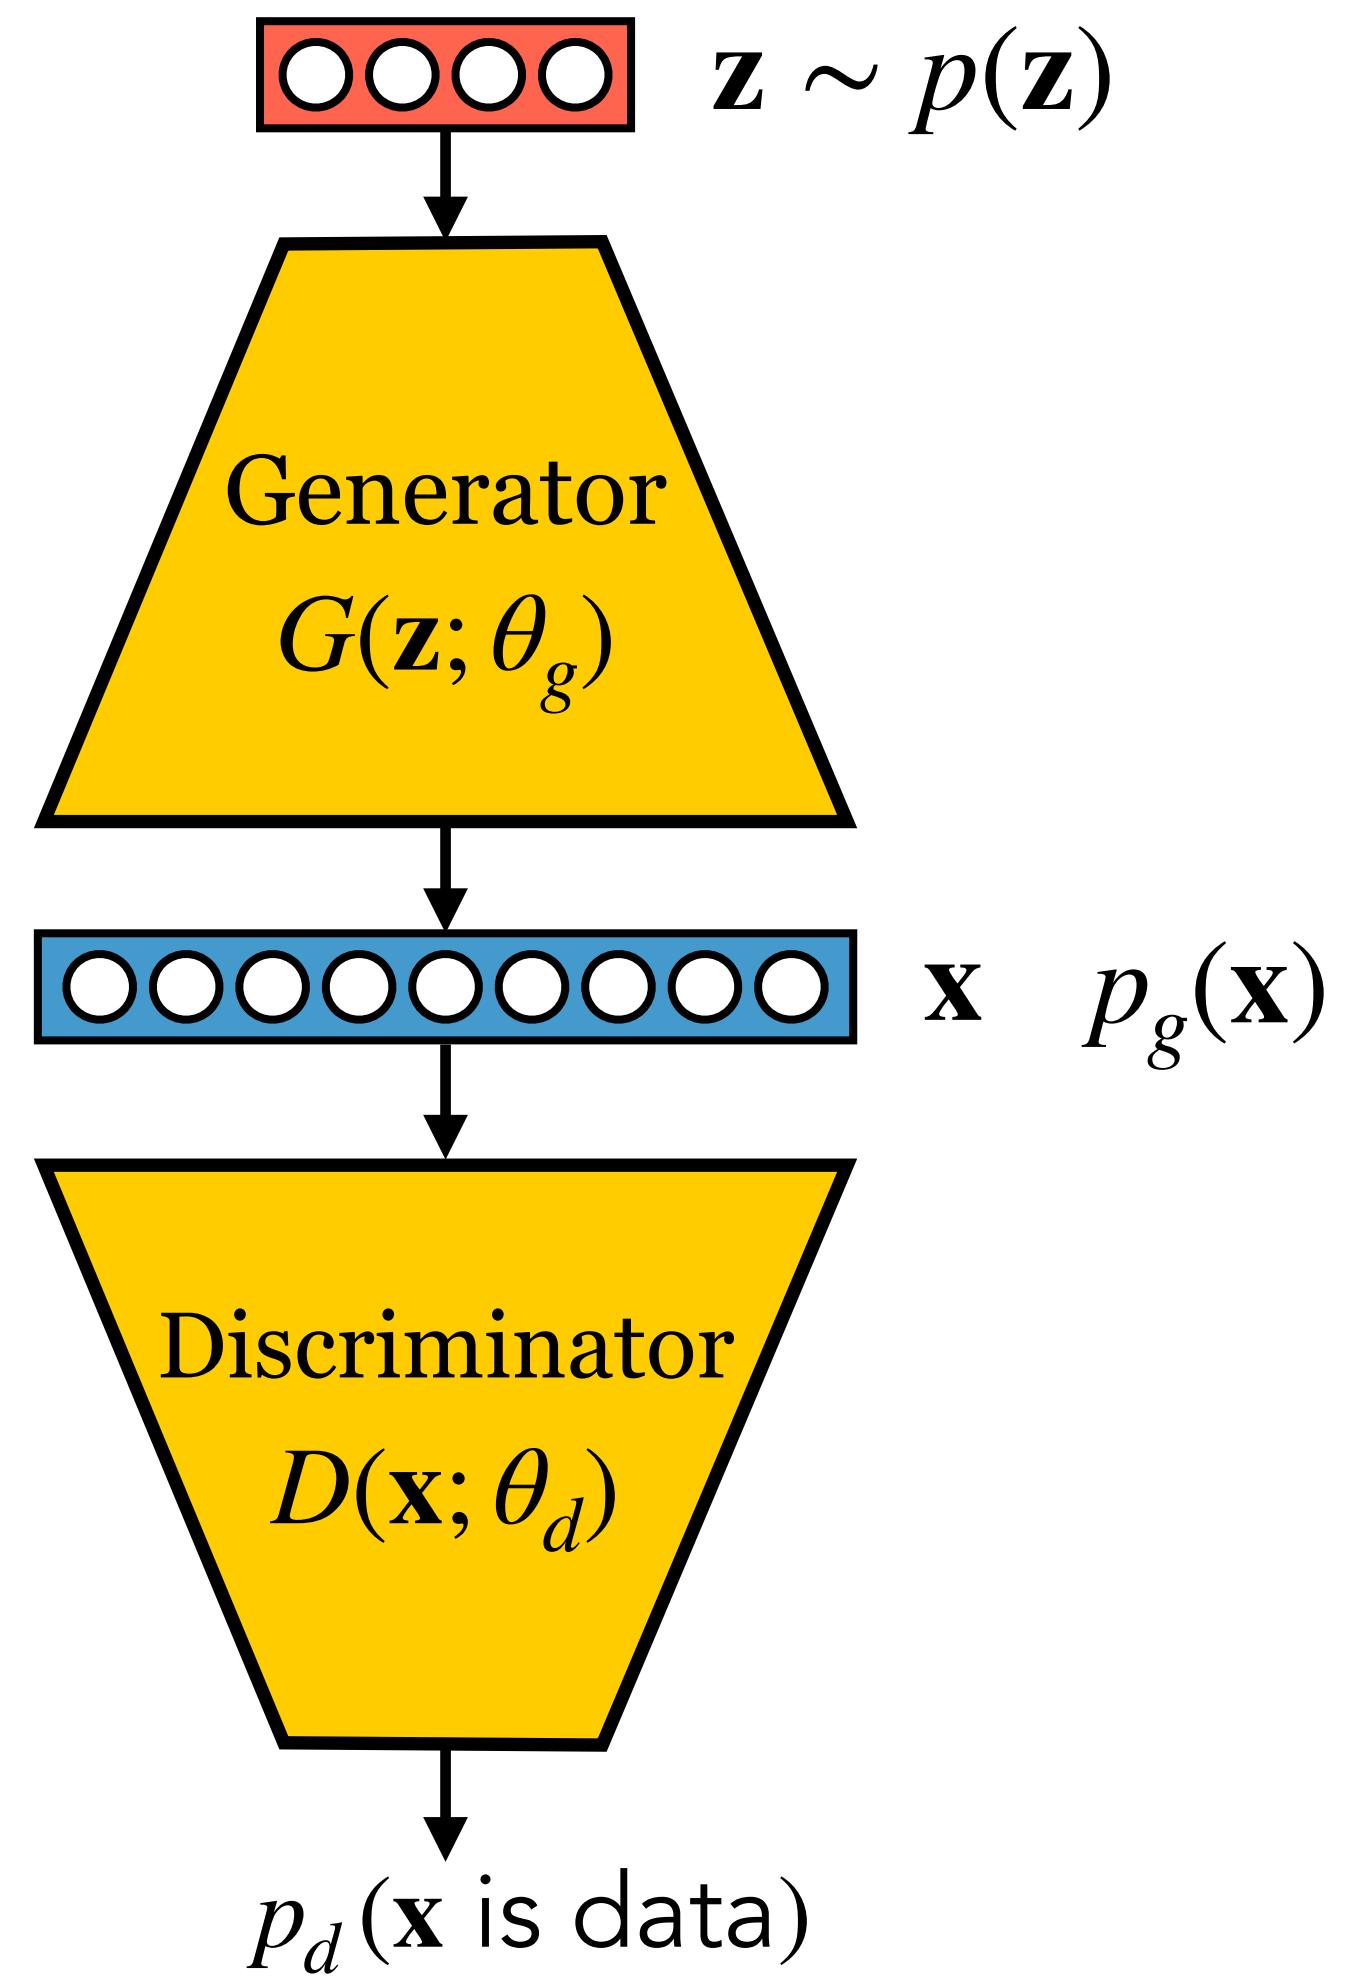
**Network input:** A random vector  $\mathbf{z}$  from some simple prior distribution  $p(\mathbf{z})$ .

**Network output:**  $G(\mathbf{z}; \theta_g)$  : A sample  $\mathbf{x}$  from a complex distribution  $p_g(\mathbf{x})$ .

**Discriminator network:**  $D$

**Network input:**  $\mathbf{x}$ , either coming from  $G$  or from  $p_{data}(\mathbf{x})$

**Network output:**  $D(\mathbf{x}; \theta_d)$  : probability  $p(\mathbf{x}$  is an actual data point)



# Training of GANs

**Discriminator network:**  $D$  is a classifier that performs binary classification.

$$\text{target} = \begin{cases} 1, & \text{if real data point} \\ 0, & \text{if generated} \end{cases}$$

**Generator network:**  $G$  is trained to minimize:

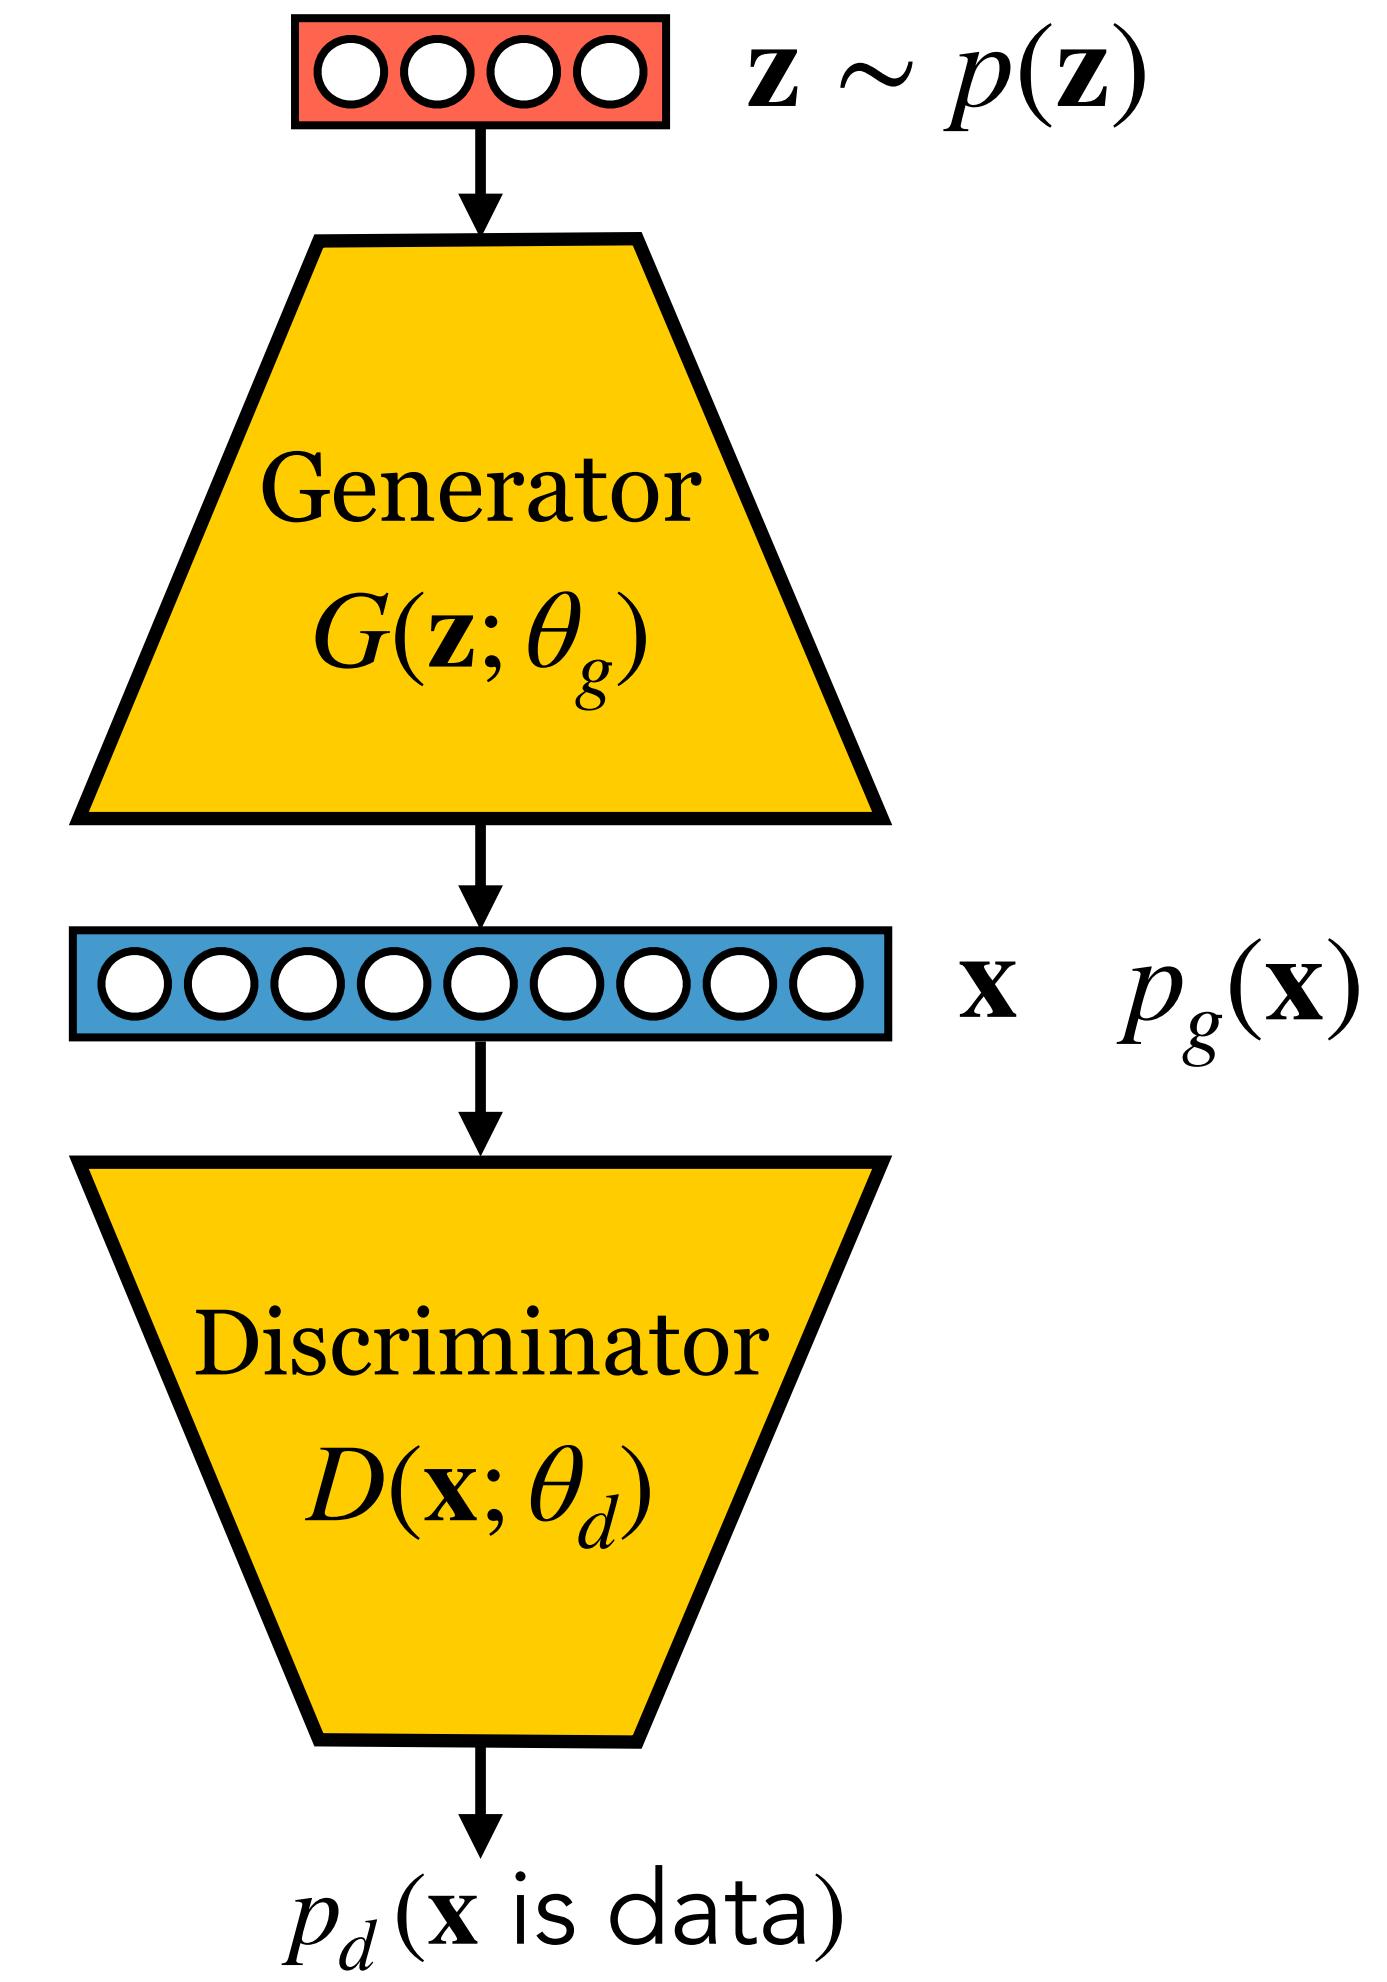
$$\log(1 - D(G(\mathbf{z})))$$

We adjust  $G$  such that  $D$  thinks the generated sample comes from *real* data.

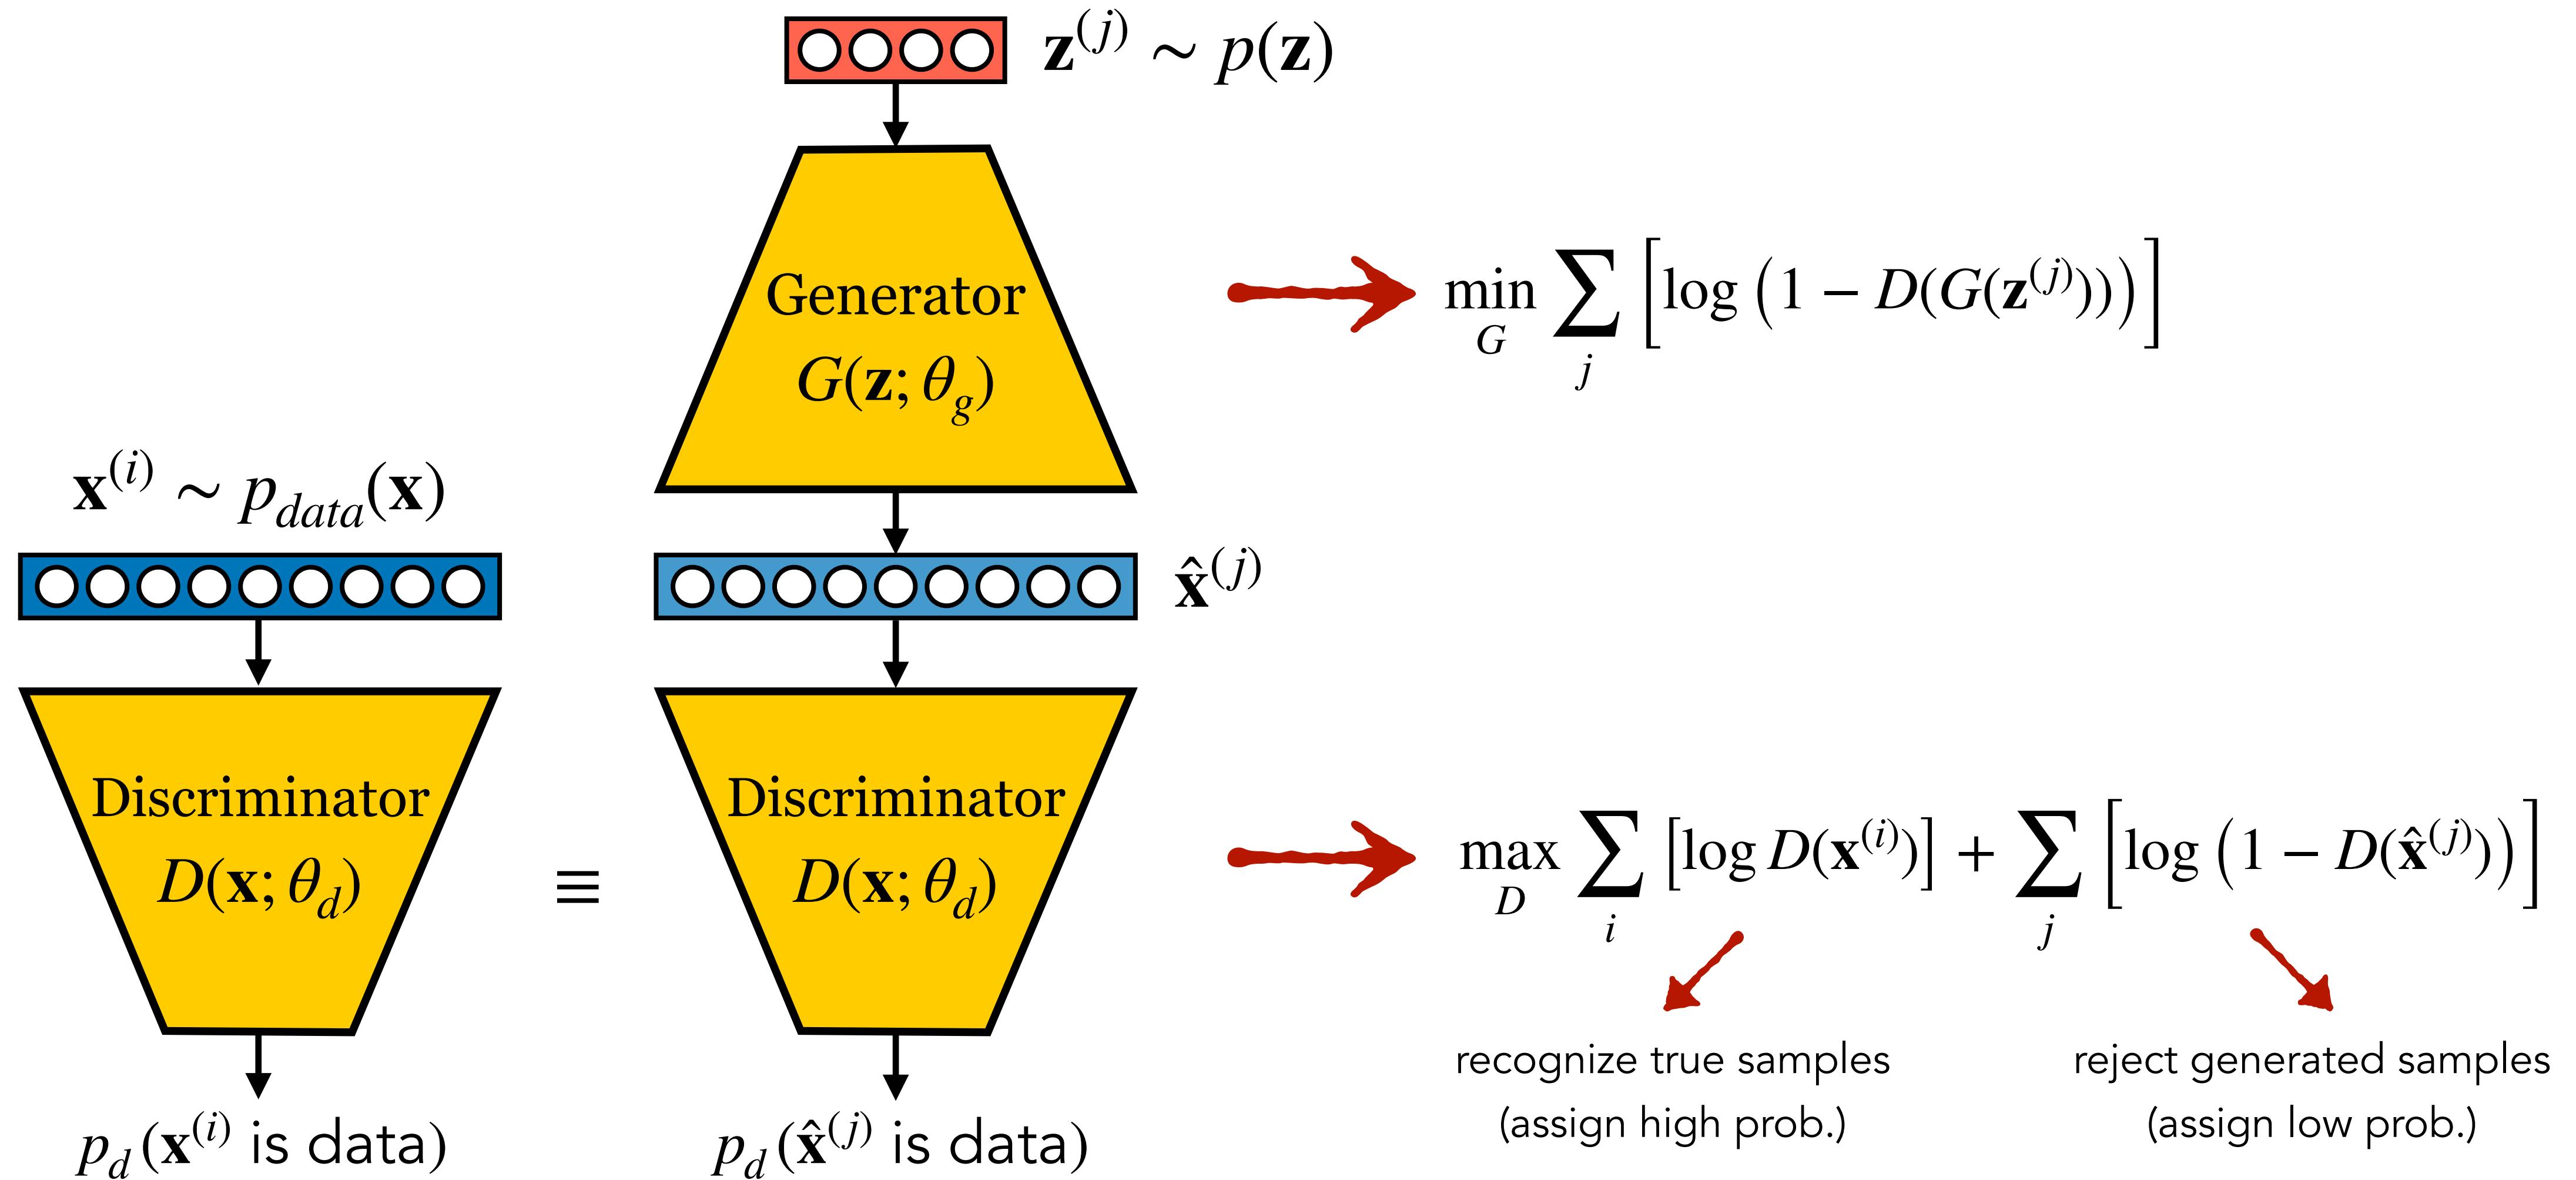
Two-player minimax game:  $G$  wants to minimize performance of  $D$ , and  $D$  wants to maximize its own performance.

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$

Global optimum (given enough capacity):  $p_g = p_{data}$



# Training of GANs



# Algorithm

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p(\mathbf{z})$
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{data}(\mathbf{x})$
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)})))] .$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p(\mathbf{z})$
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))) .$$

**end for**

# Algorithm

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$


**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p(\mathbf{z})$
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{data}(\mathbf{x})$
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)})))] .$$

**end for**

►  $D$  is optimized  $k$  times (typically  $k \in \{1, \dots, 5\}$ ),  
for each optimization step of  $G$ .

**end for**

# Algorithm

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$



**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p(\mathbf{z})$
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{data}(\mathbf{x})$
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)})))] .$$

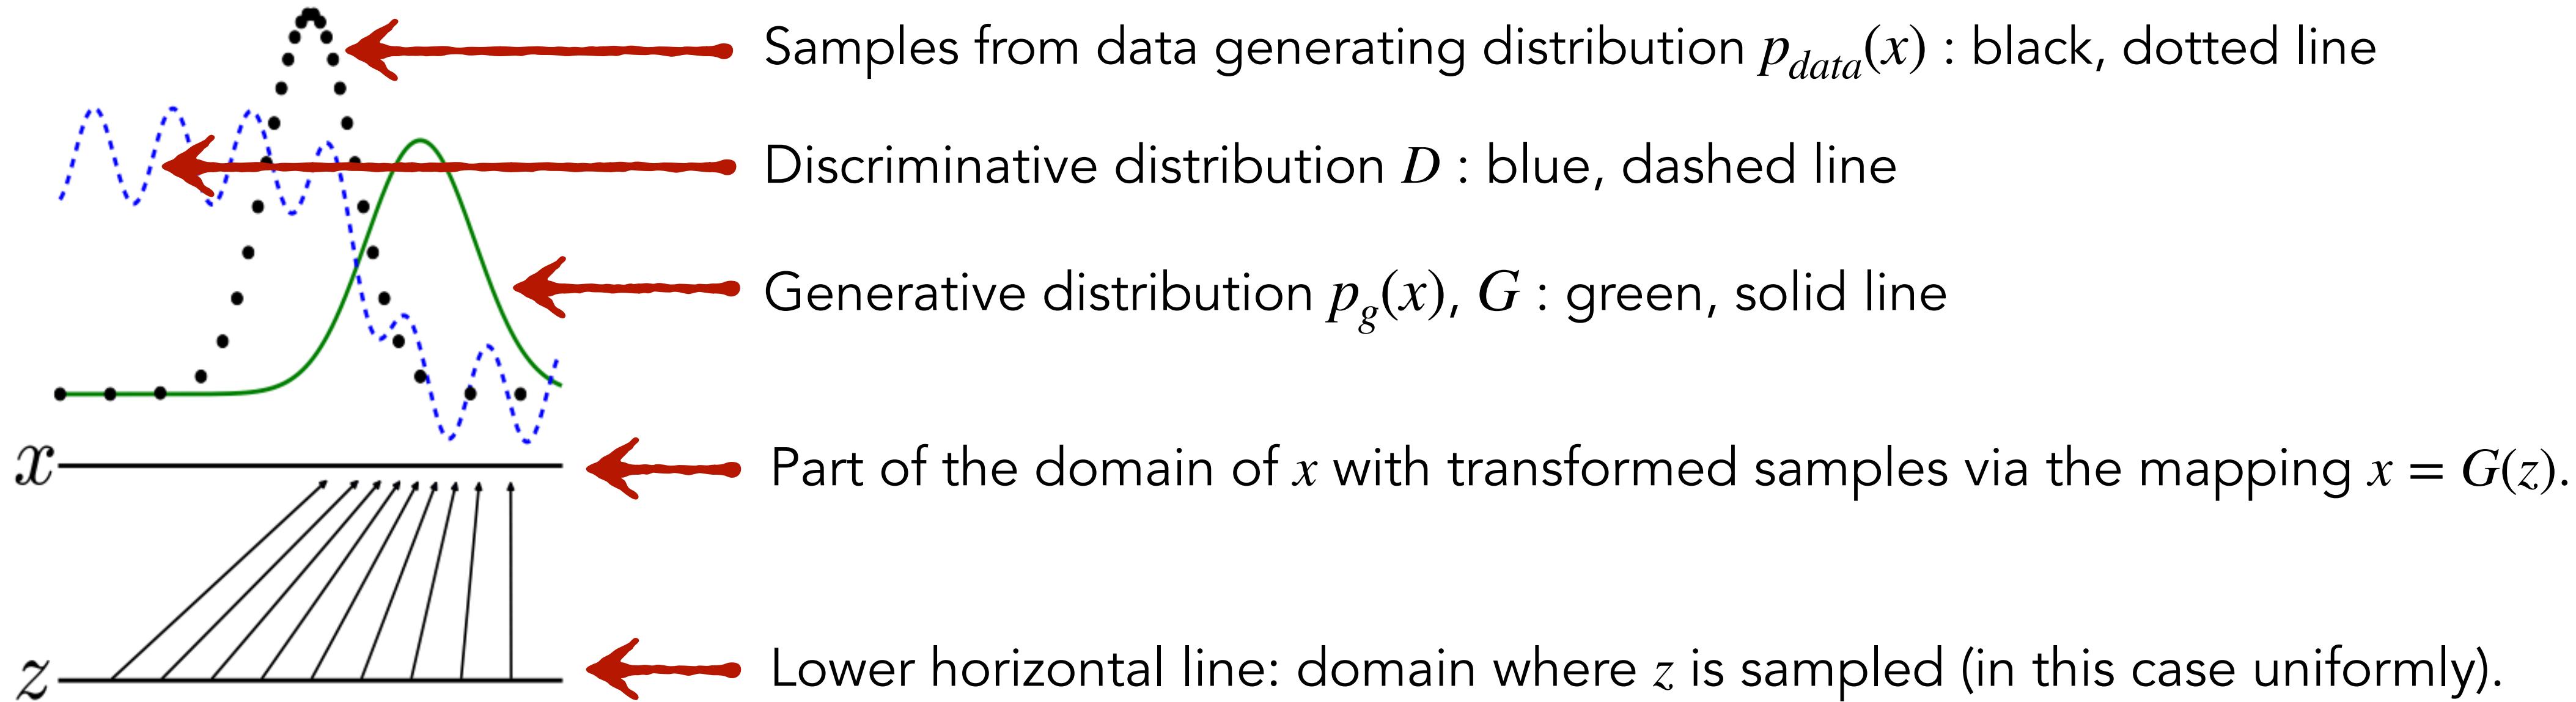
**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p(\mathbf{z})$
- Update the generator by descending its stochastic gradient:

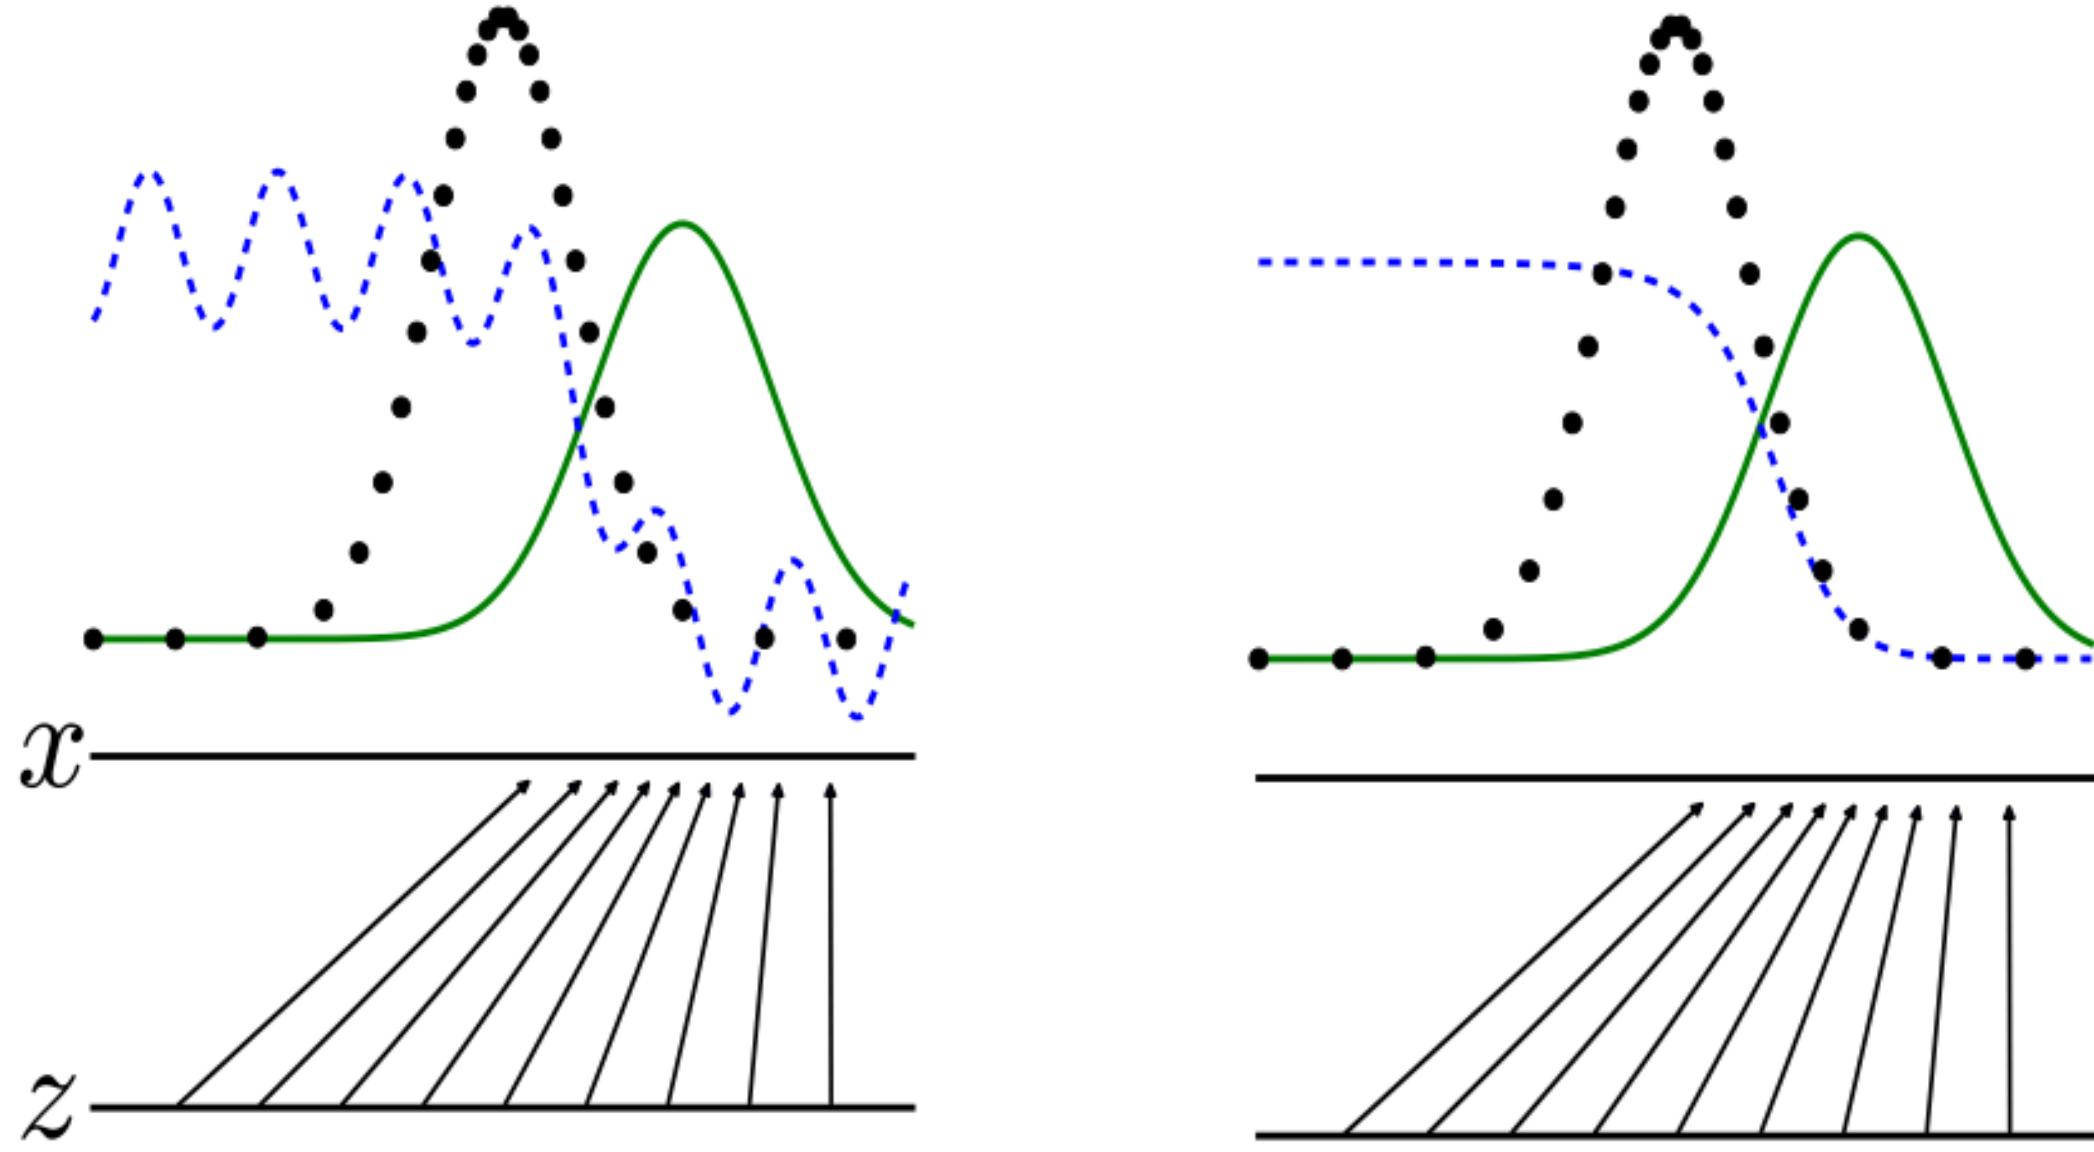
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))) .$$

**end for**

# Training of GANs: Basic Intuition

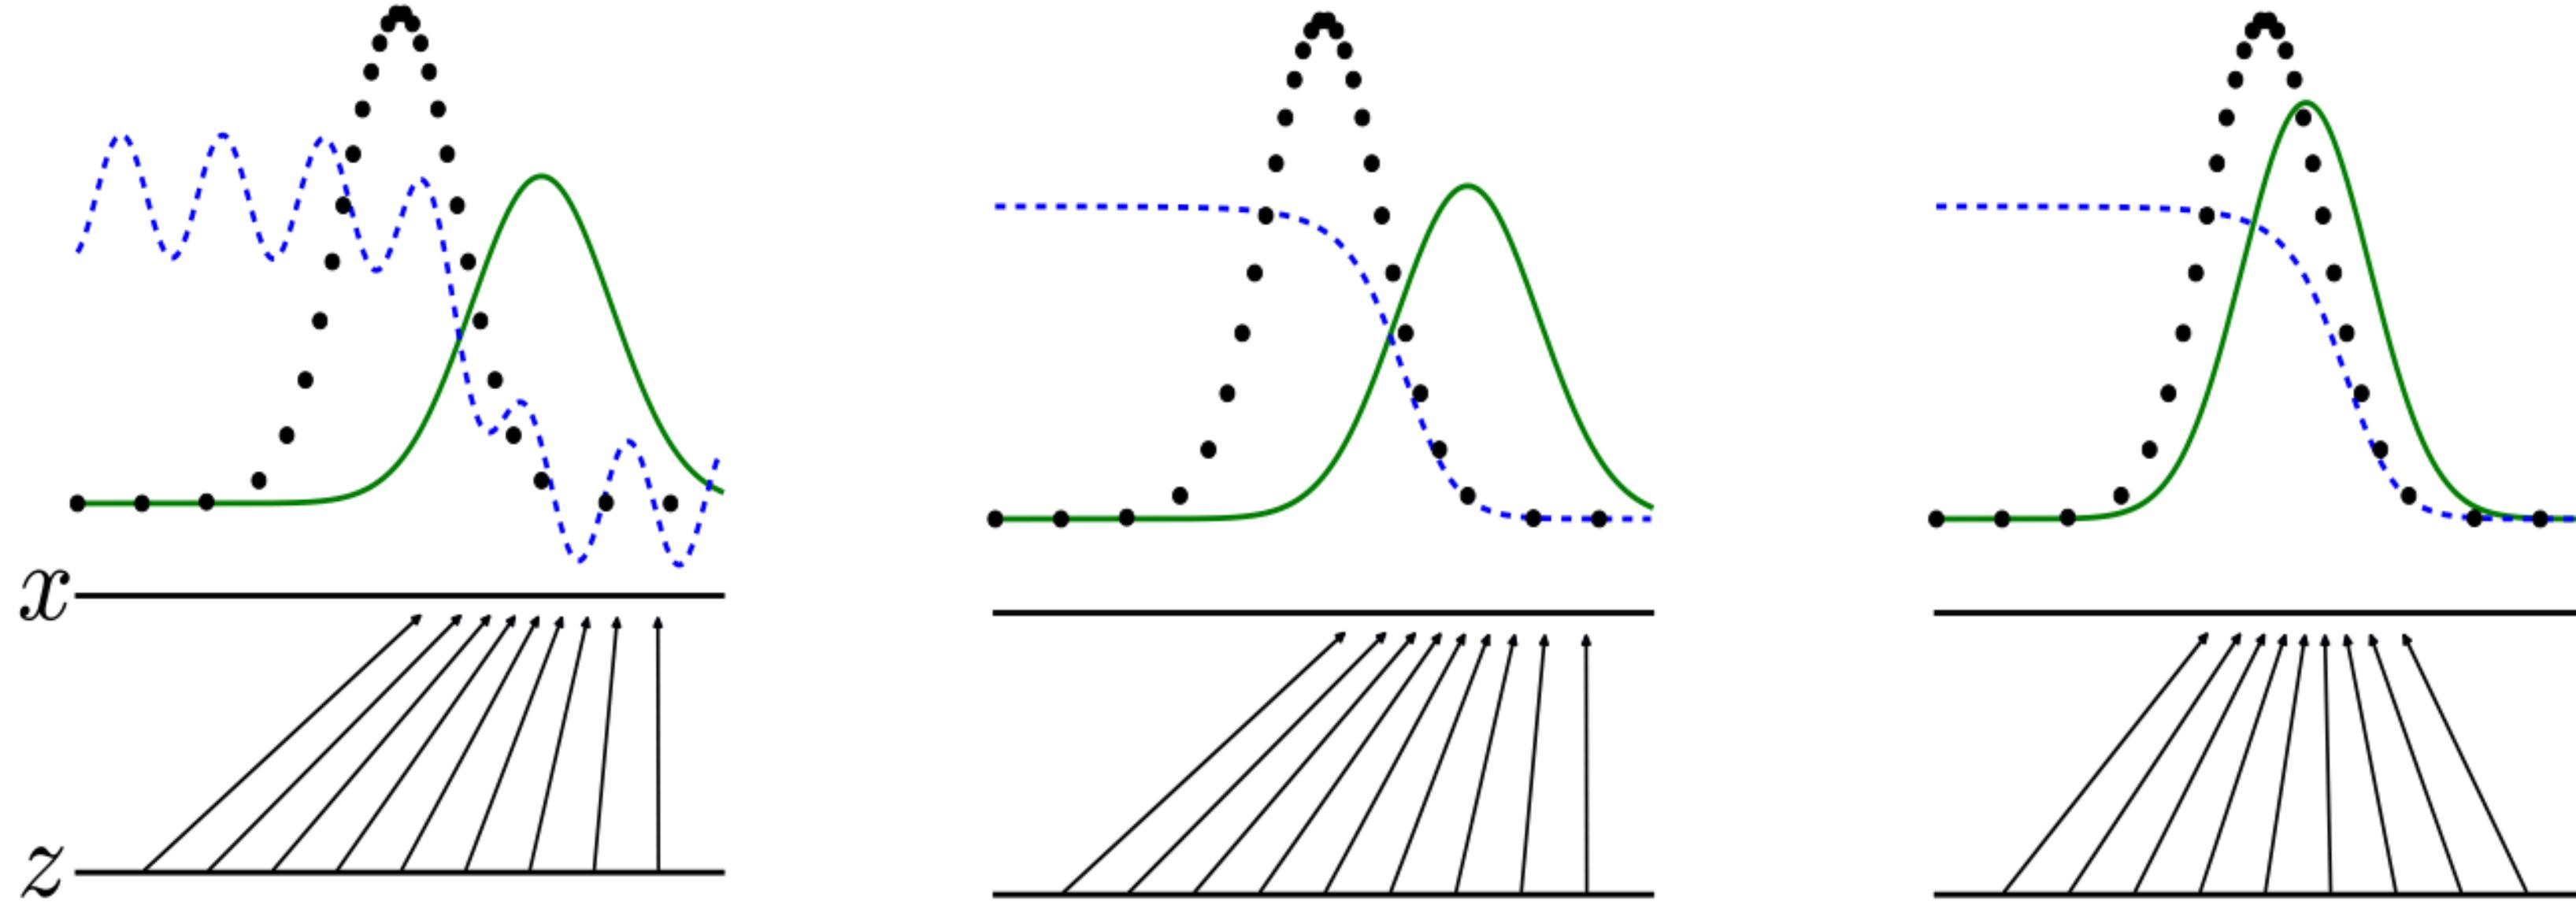


# Training of GANs: Basic Intuition



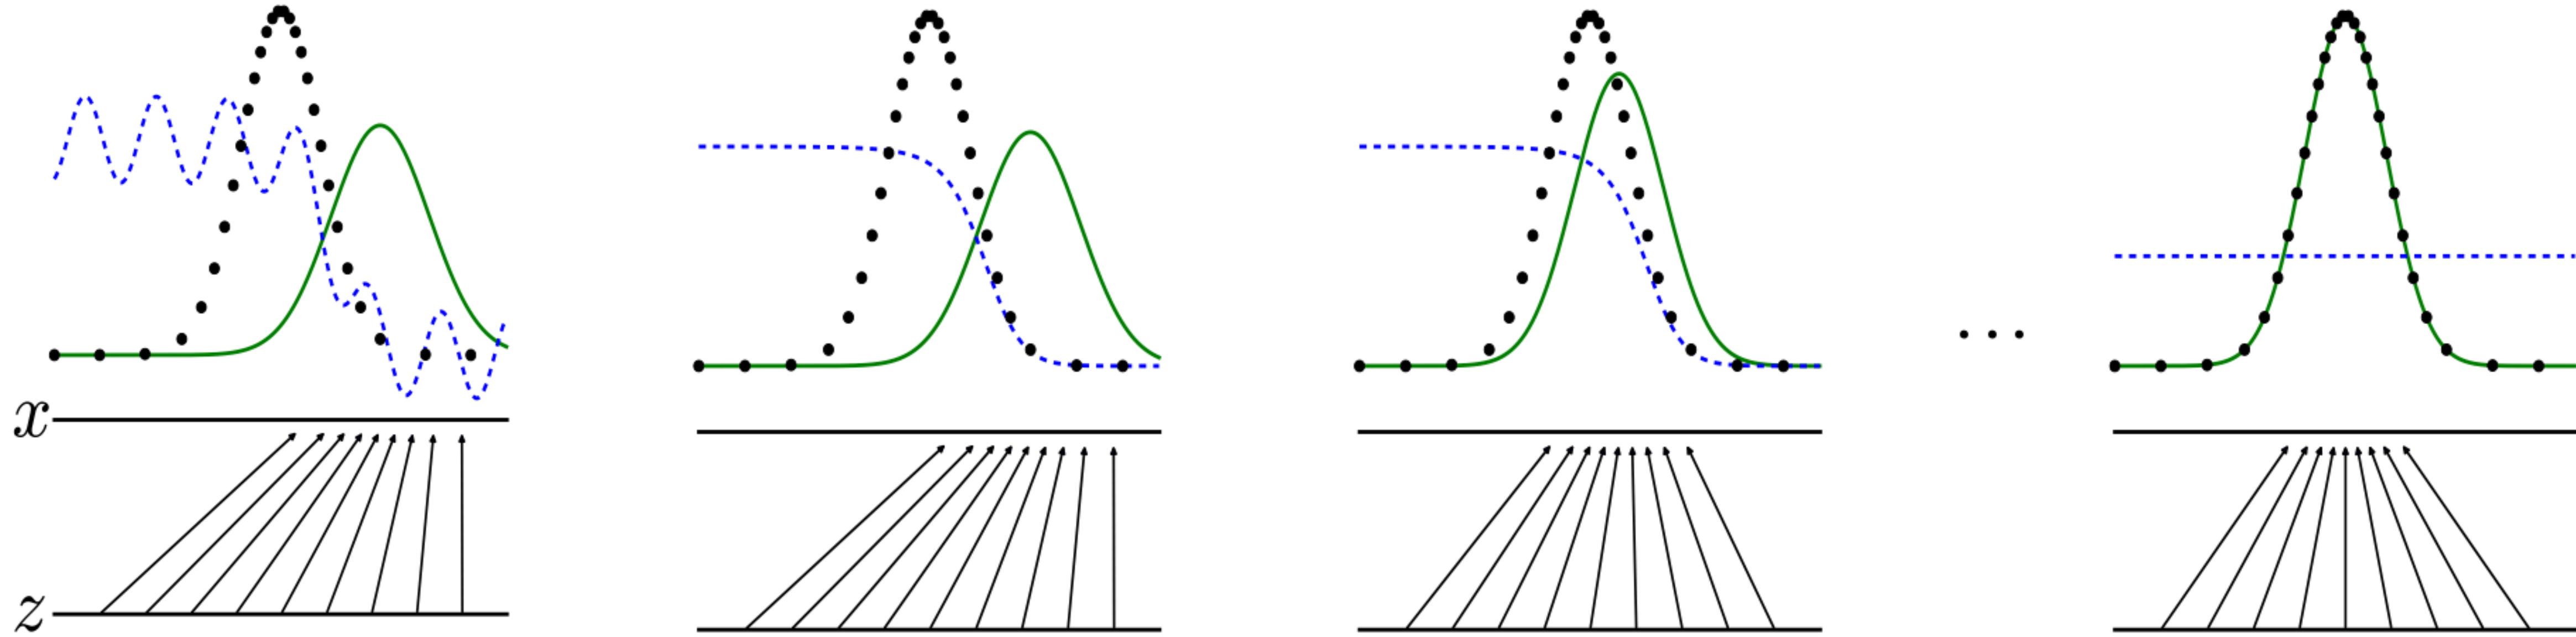
- ▶ In the inner loop,  $D$  is trained to discriminate real data, converging to  $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$ .

# Training of GANs: Basic Intuition



- ▶ In the inner loop,  $D$  is trained to discriminate real data, converging to  $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$ .
- ▶ After an update to  $G$ , gradient of  $D$  guides  $G$  to regions that are more likely to be classified as real data.

# Training of GANs: Basic Intuition



- ▶ In the inner loop,  $D$  is trained to discriminate real data, converging to  $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$ .
- ▶ After an update to  $G$ , gradient of  $D$  guides  $G$  to regions that are more likely to be classified as real data.
- ▶ Given enough capacity, they reach to  $p_g = p_{data}$  and  $D(x) = \frac{1}{2}$ .

# Today

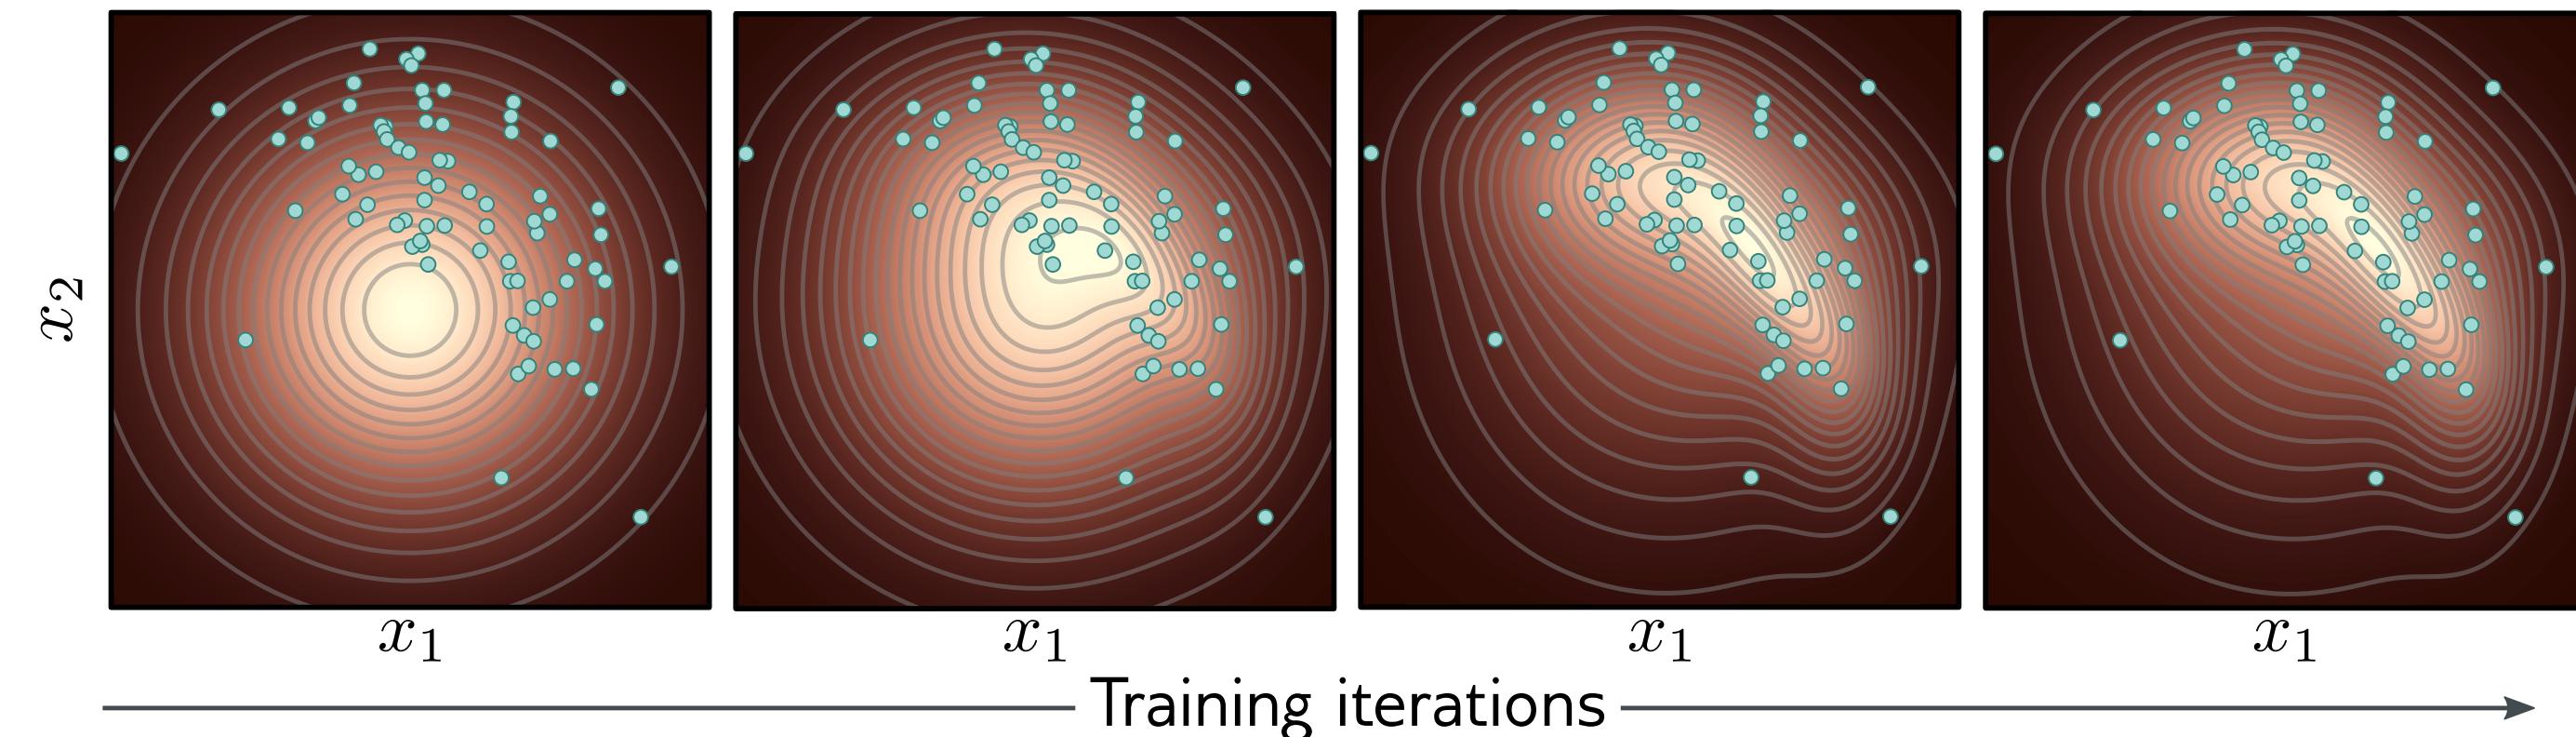
---

- Deep Generative Models
- Variational Autoencoders
- Generative Adversarial Networks
- Architecture & Training
- Mode Collapse
- Further Extensions & Variants

# Comparison to Variational Autoencoders

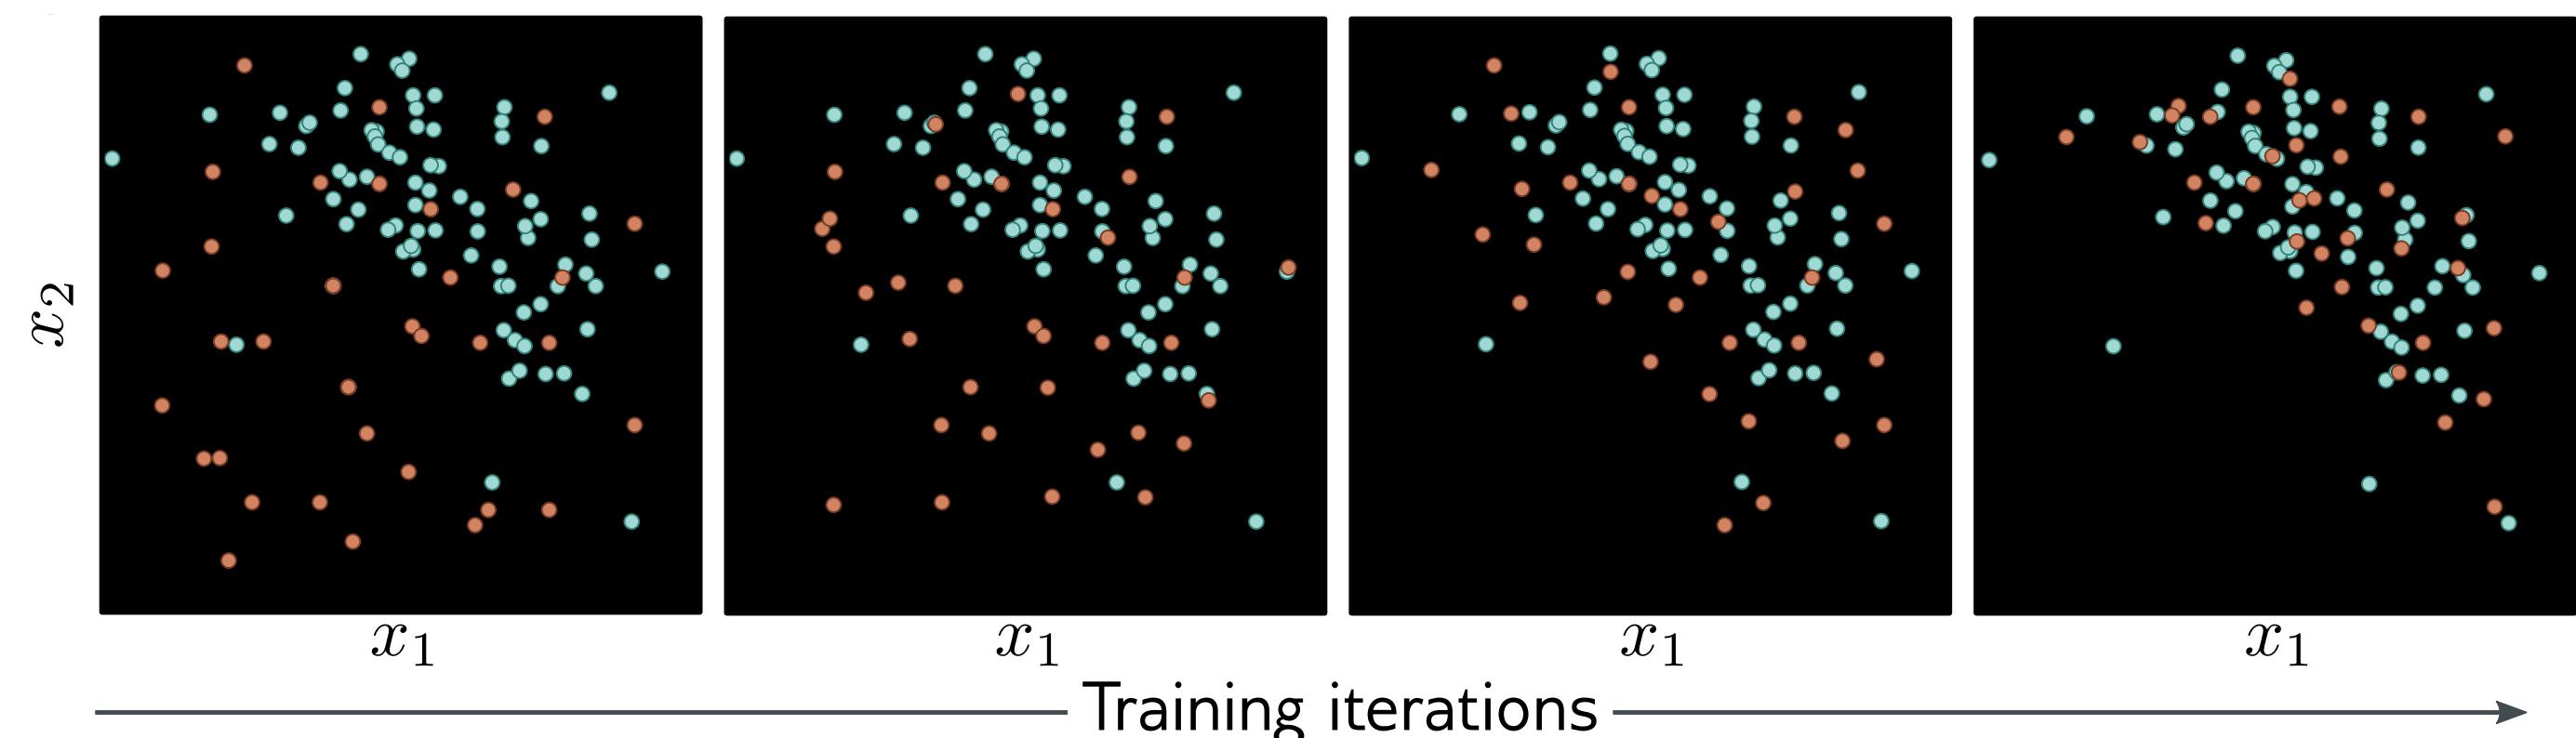
## Variational autoencoder:

- Aims to learn a probability distribution over the training data.
- Likelihood of **real examples** increase under this distribution (by maximizing VLB), which is then used to draw new samples.



## Generative adversarial network:

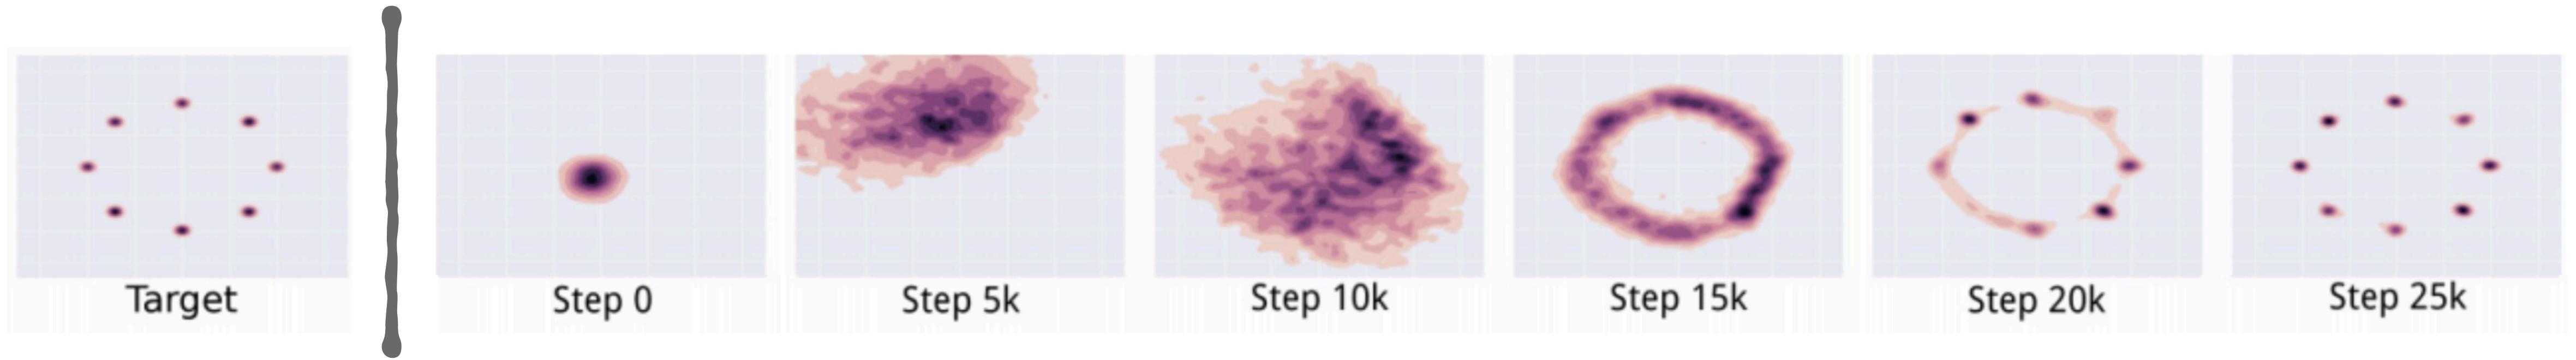
- Avoids the issue of explicitly designing a tractable density function.
- A mechanism for **generating samples**. Training encourages samples to become less distinguishable from **real examples**.



# Mode Coverage

An example of “good mode coverage”:

Here, the generative model is trained on data from a mixture of Gaussians.

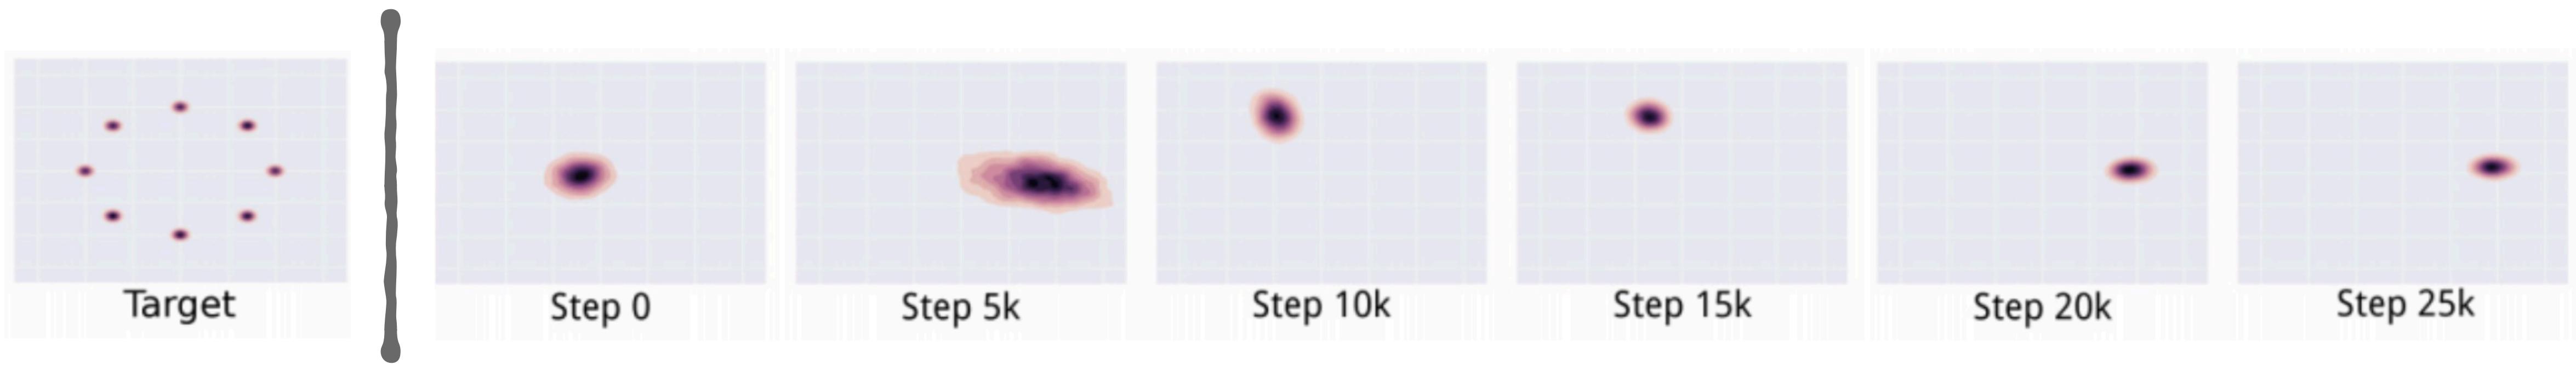


(Simple) GANs often struggle to achieve this.

# Mode-Collapse

**One main problem of GANs is the so-called “Mode-Collapse”:**

This happens when many latent configurations are mapped to single or similar outputs in the generator.

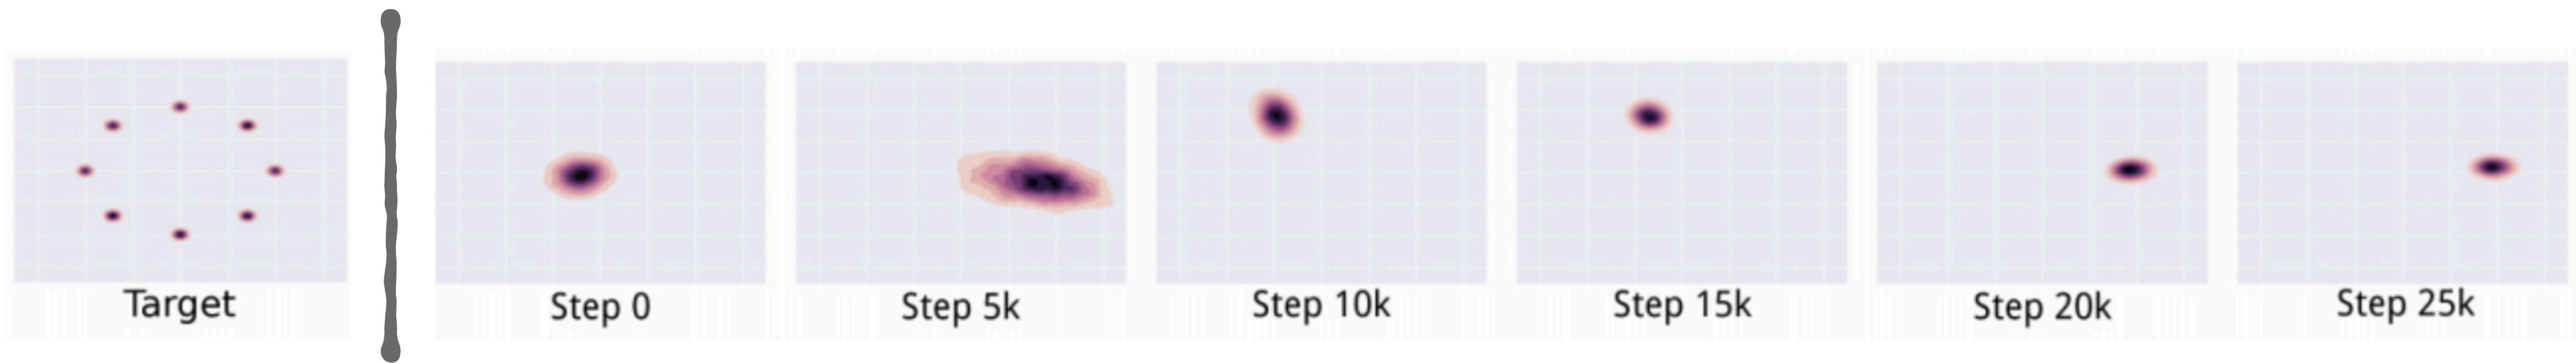


# Mode-Collapse

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$

**One main problem of GANs is the so-called “Mode-Collapse”:**

This happens when many latent configurations are mapped to single or similar outputs in the generator.



- The GAN switches between different modes during training when being trained on a mixture of Gaussians.
- When the generator sticks to one out of  $n$  modes (of equal probability mass), it is  $n$  times more likely that an  $\mathbf{x}$  from this mode comes from the generator (normalization of probabilities). Hence, the discriminator will predict that this  $\mathbf{x}$  is fake.
- Result: Generating high quality samples with very limited variability, covering only a part of  $p_{data}(\mathbf{x})$ .

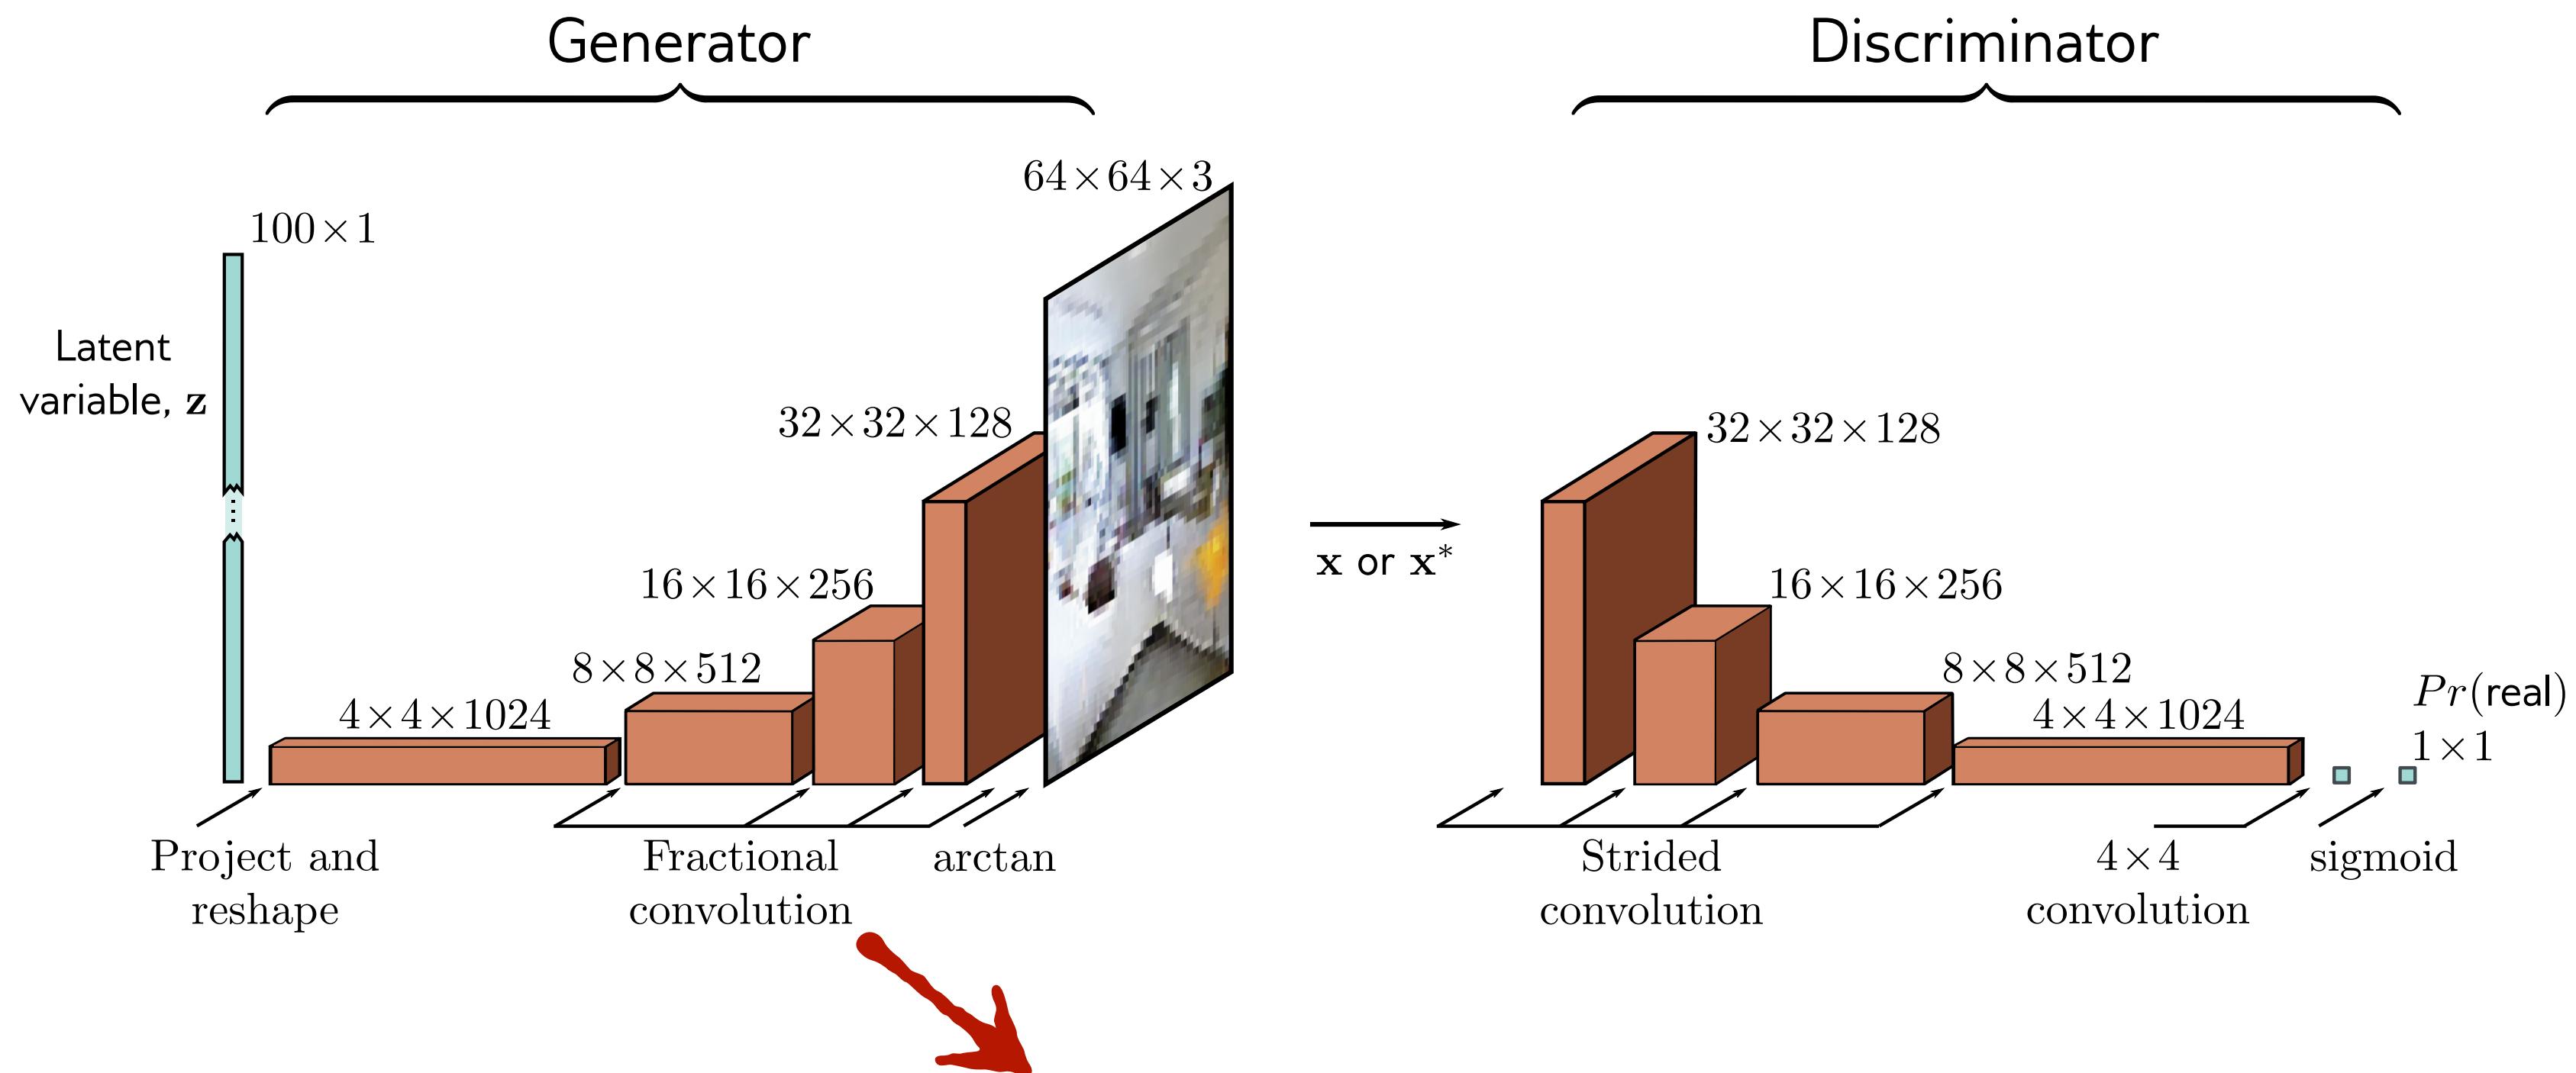
# Today

---

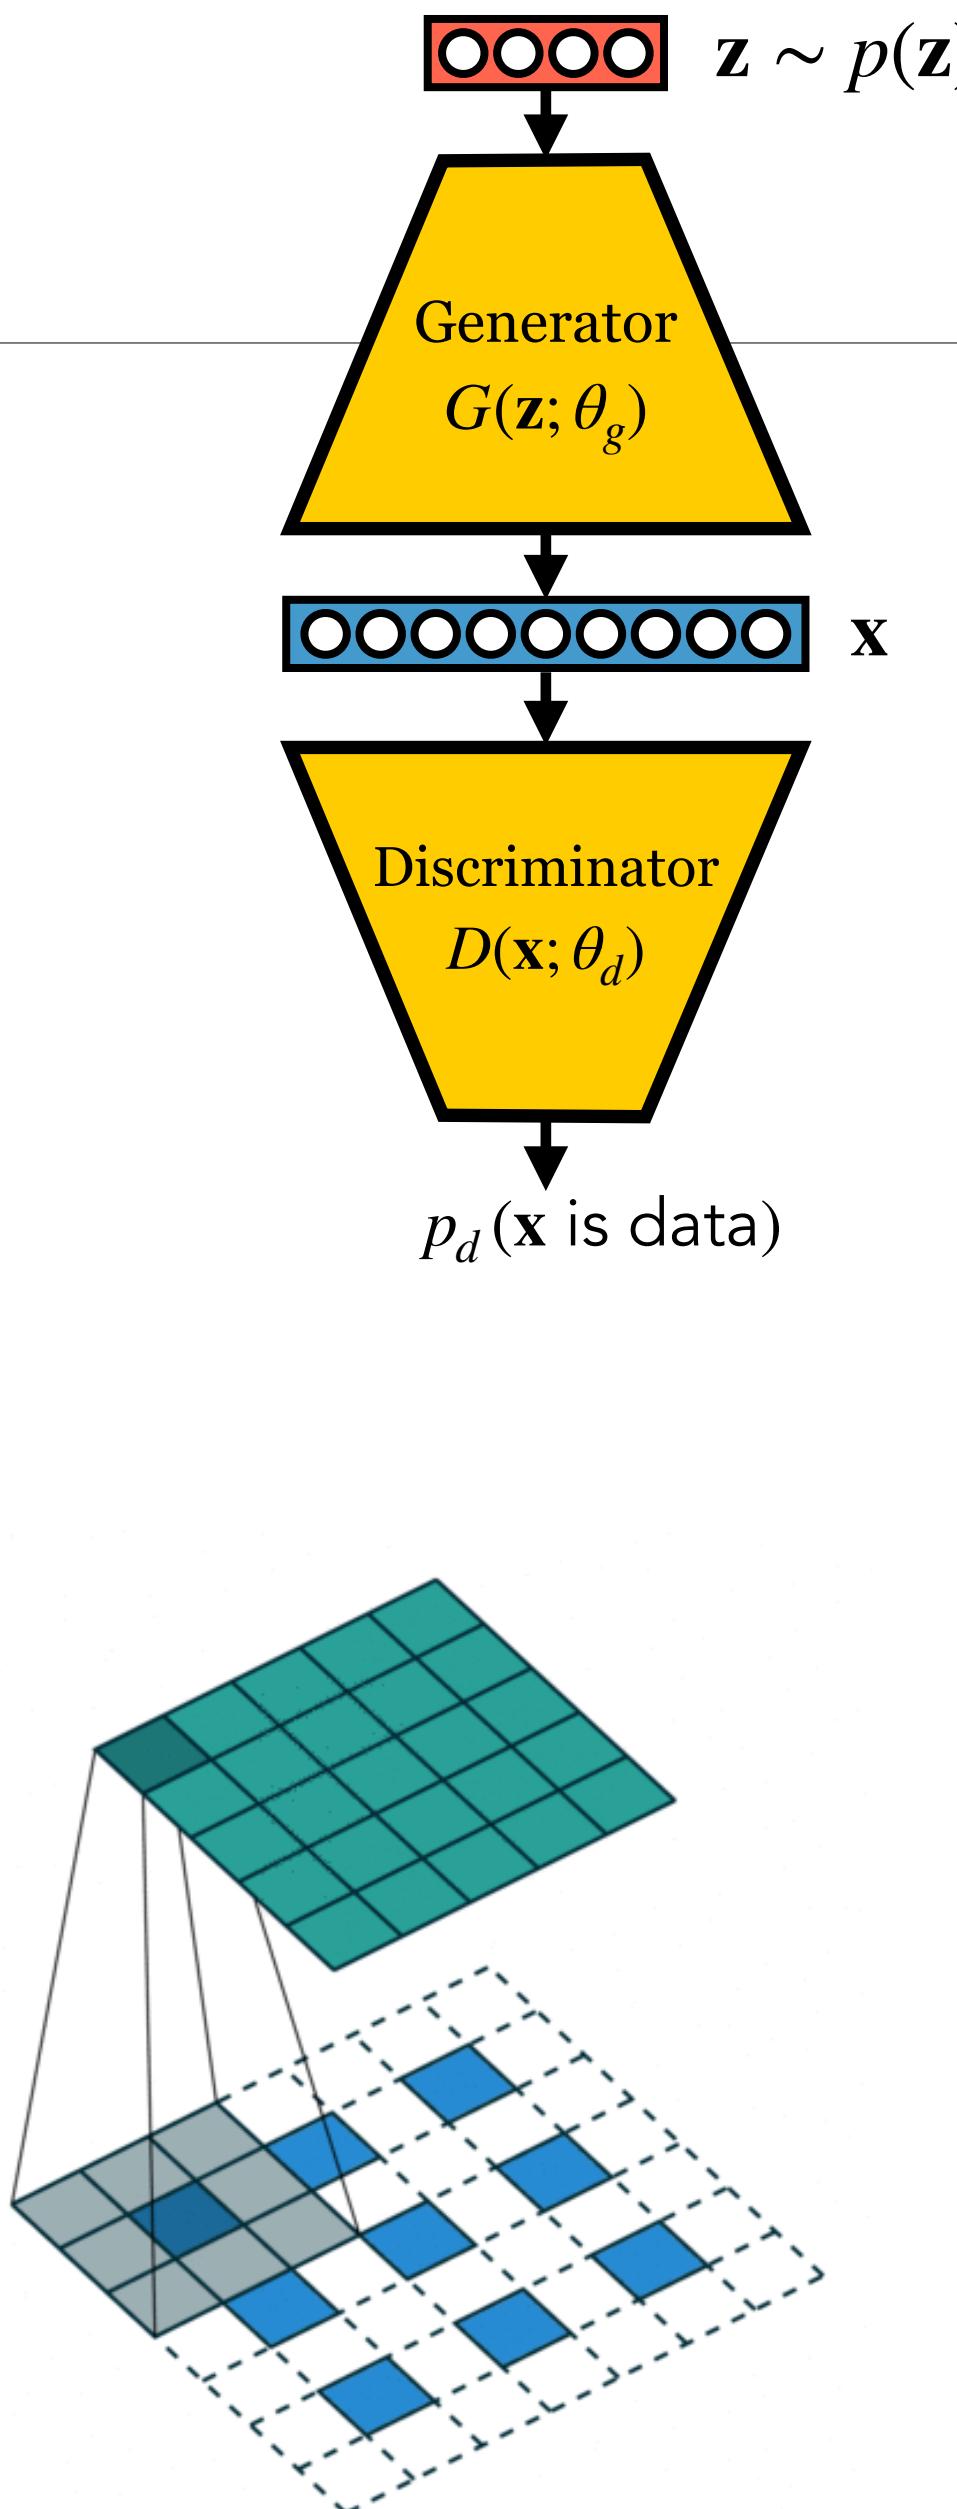
- Deep Generative Models
- Variational Autoencoders
- Generative Adversarial Networks
- Architecture & Training
- Mode Collapse
- Further Extensions & Variants

# Deep Convolutional GANs (DCGAN)

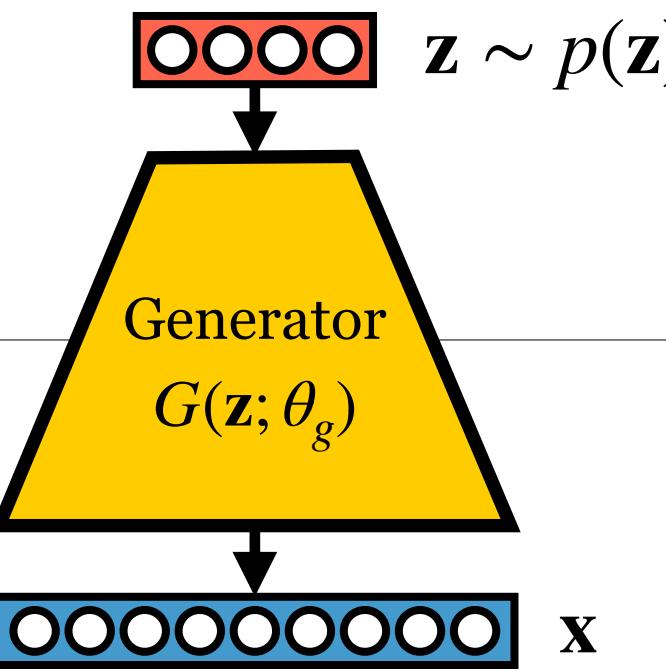
Uses a convolutional architecture without (de)-pooling and with Batch-Norm.



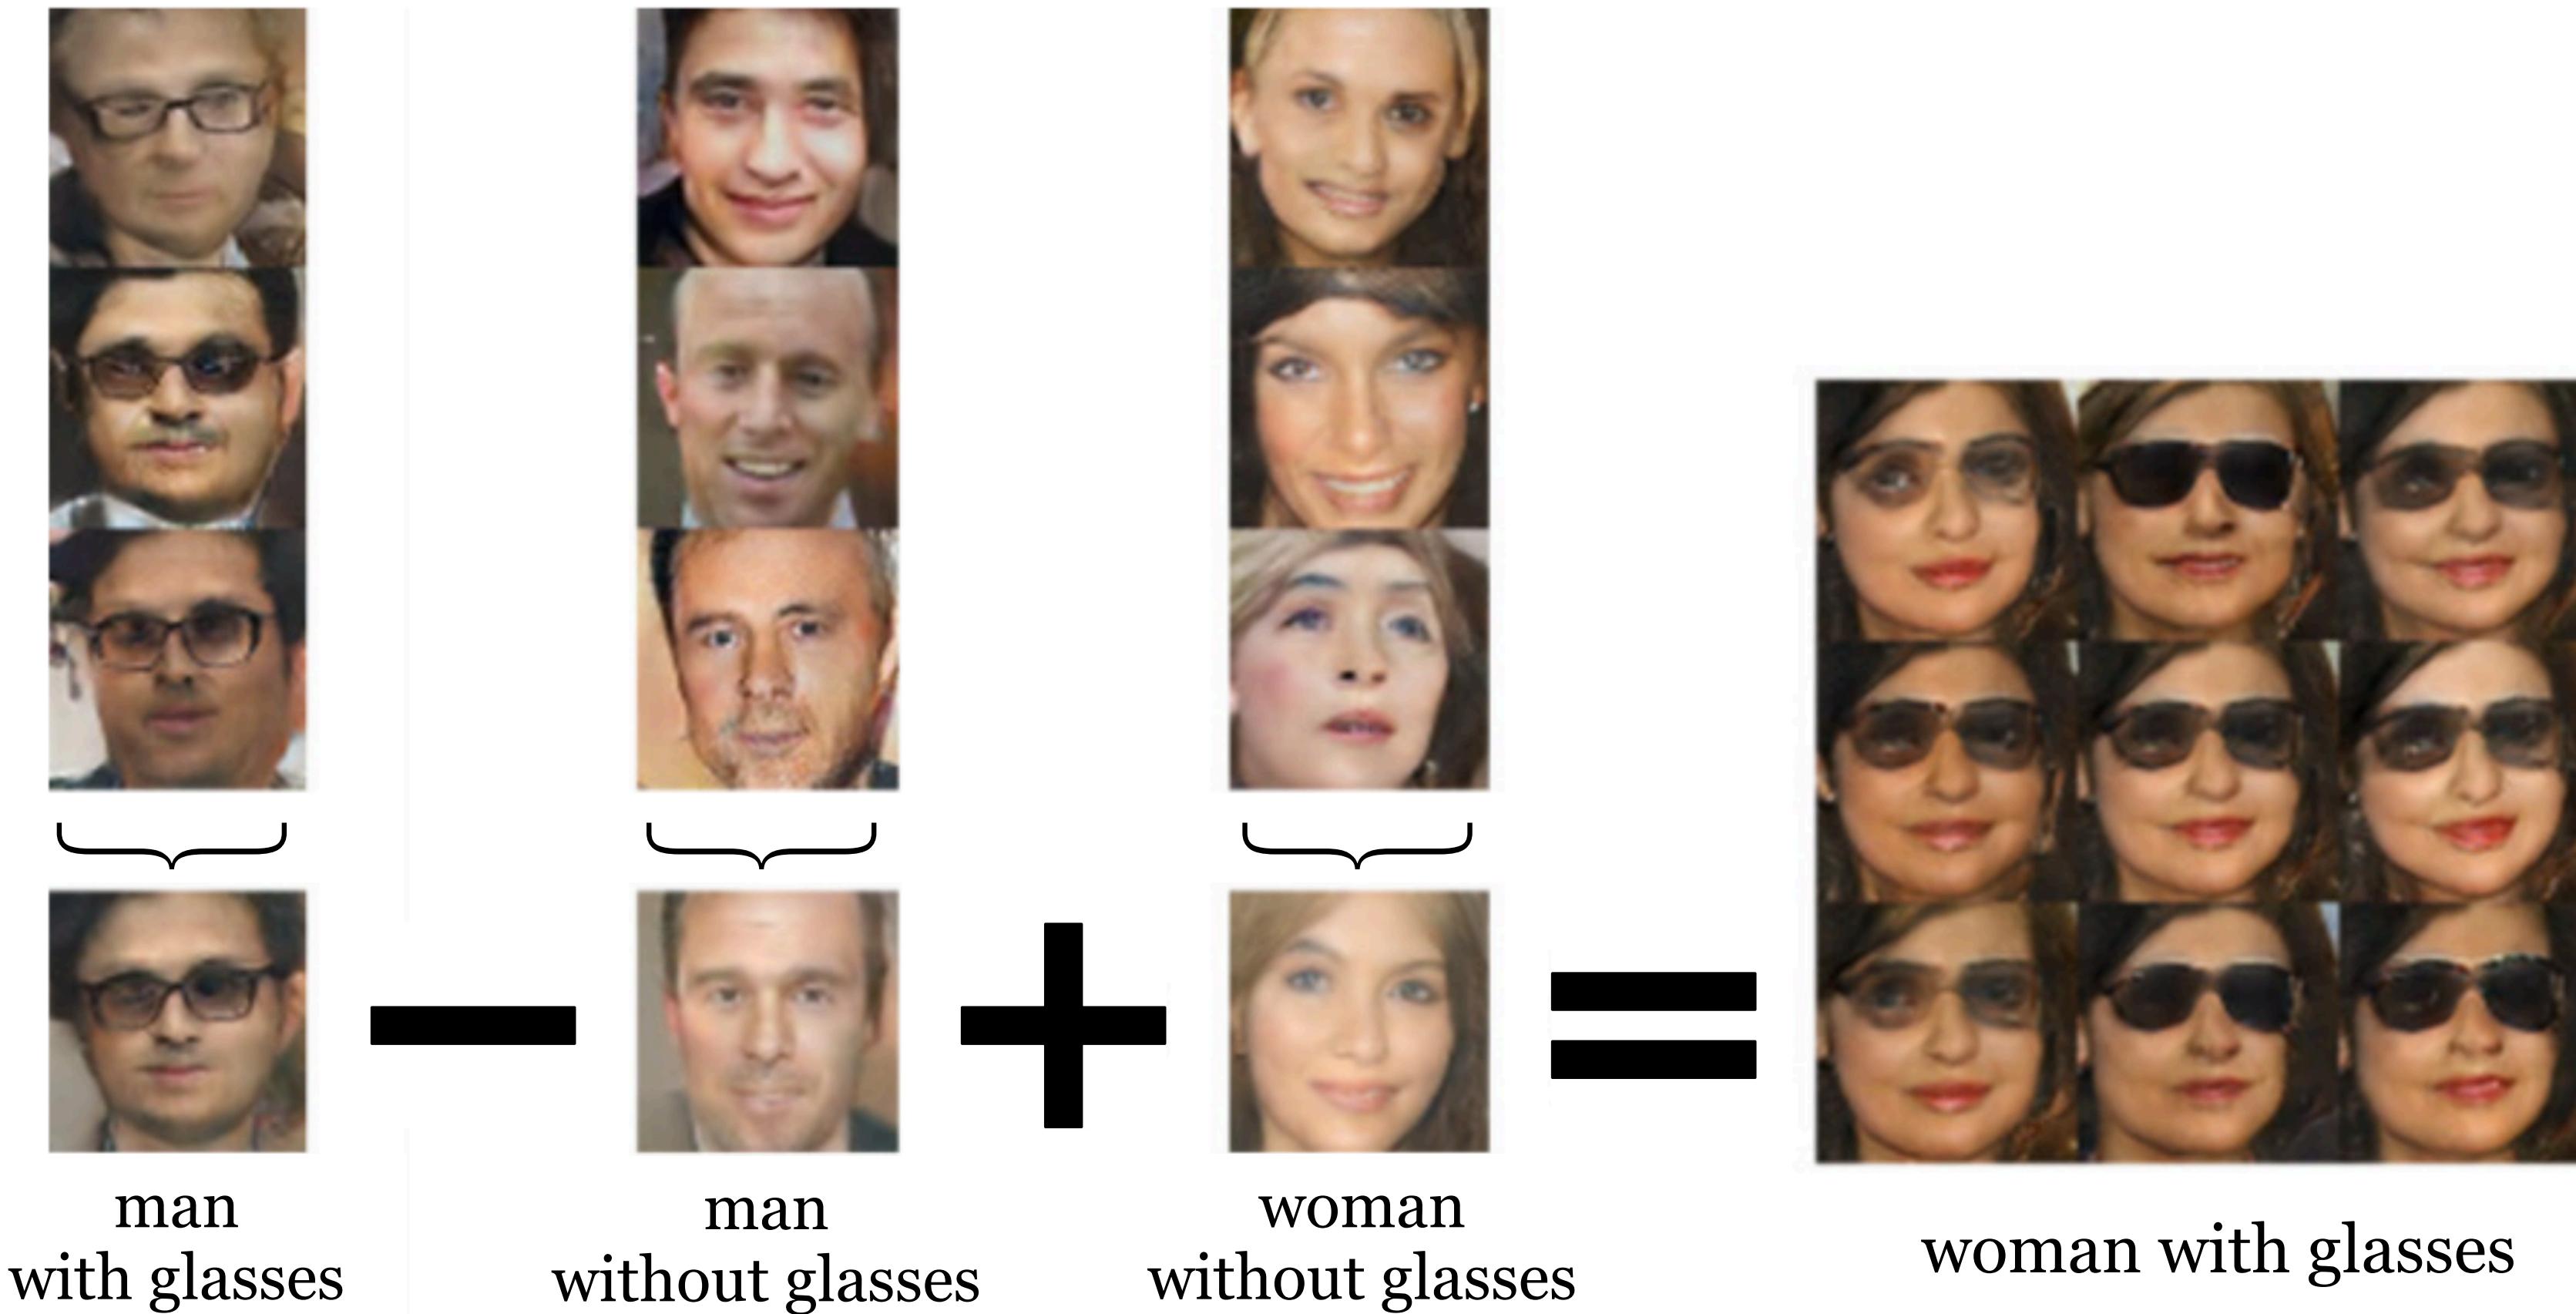
**Fractionally-strided convolutions:** Transposed convolutions with strides & adding empty pixels between the real ones.



# Vector Arithmetics in Latent Space



It turns out, semantic arithmetics can be performed in latent space of GANs.



# Conditional GANs

One can use an additional (condition) information as an extra input to the generator.

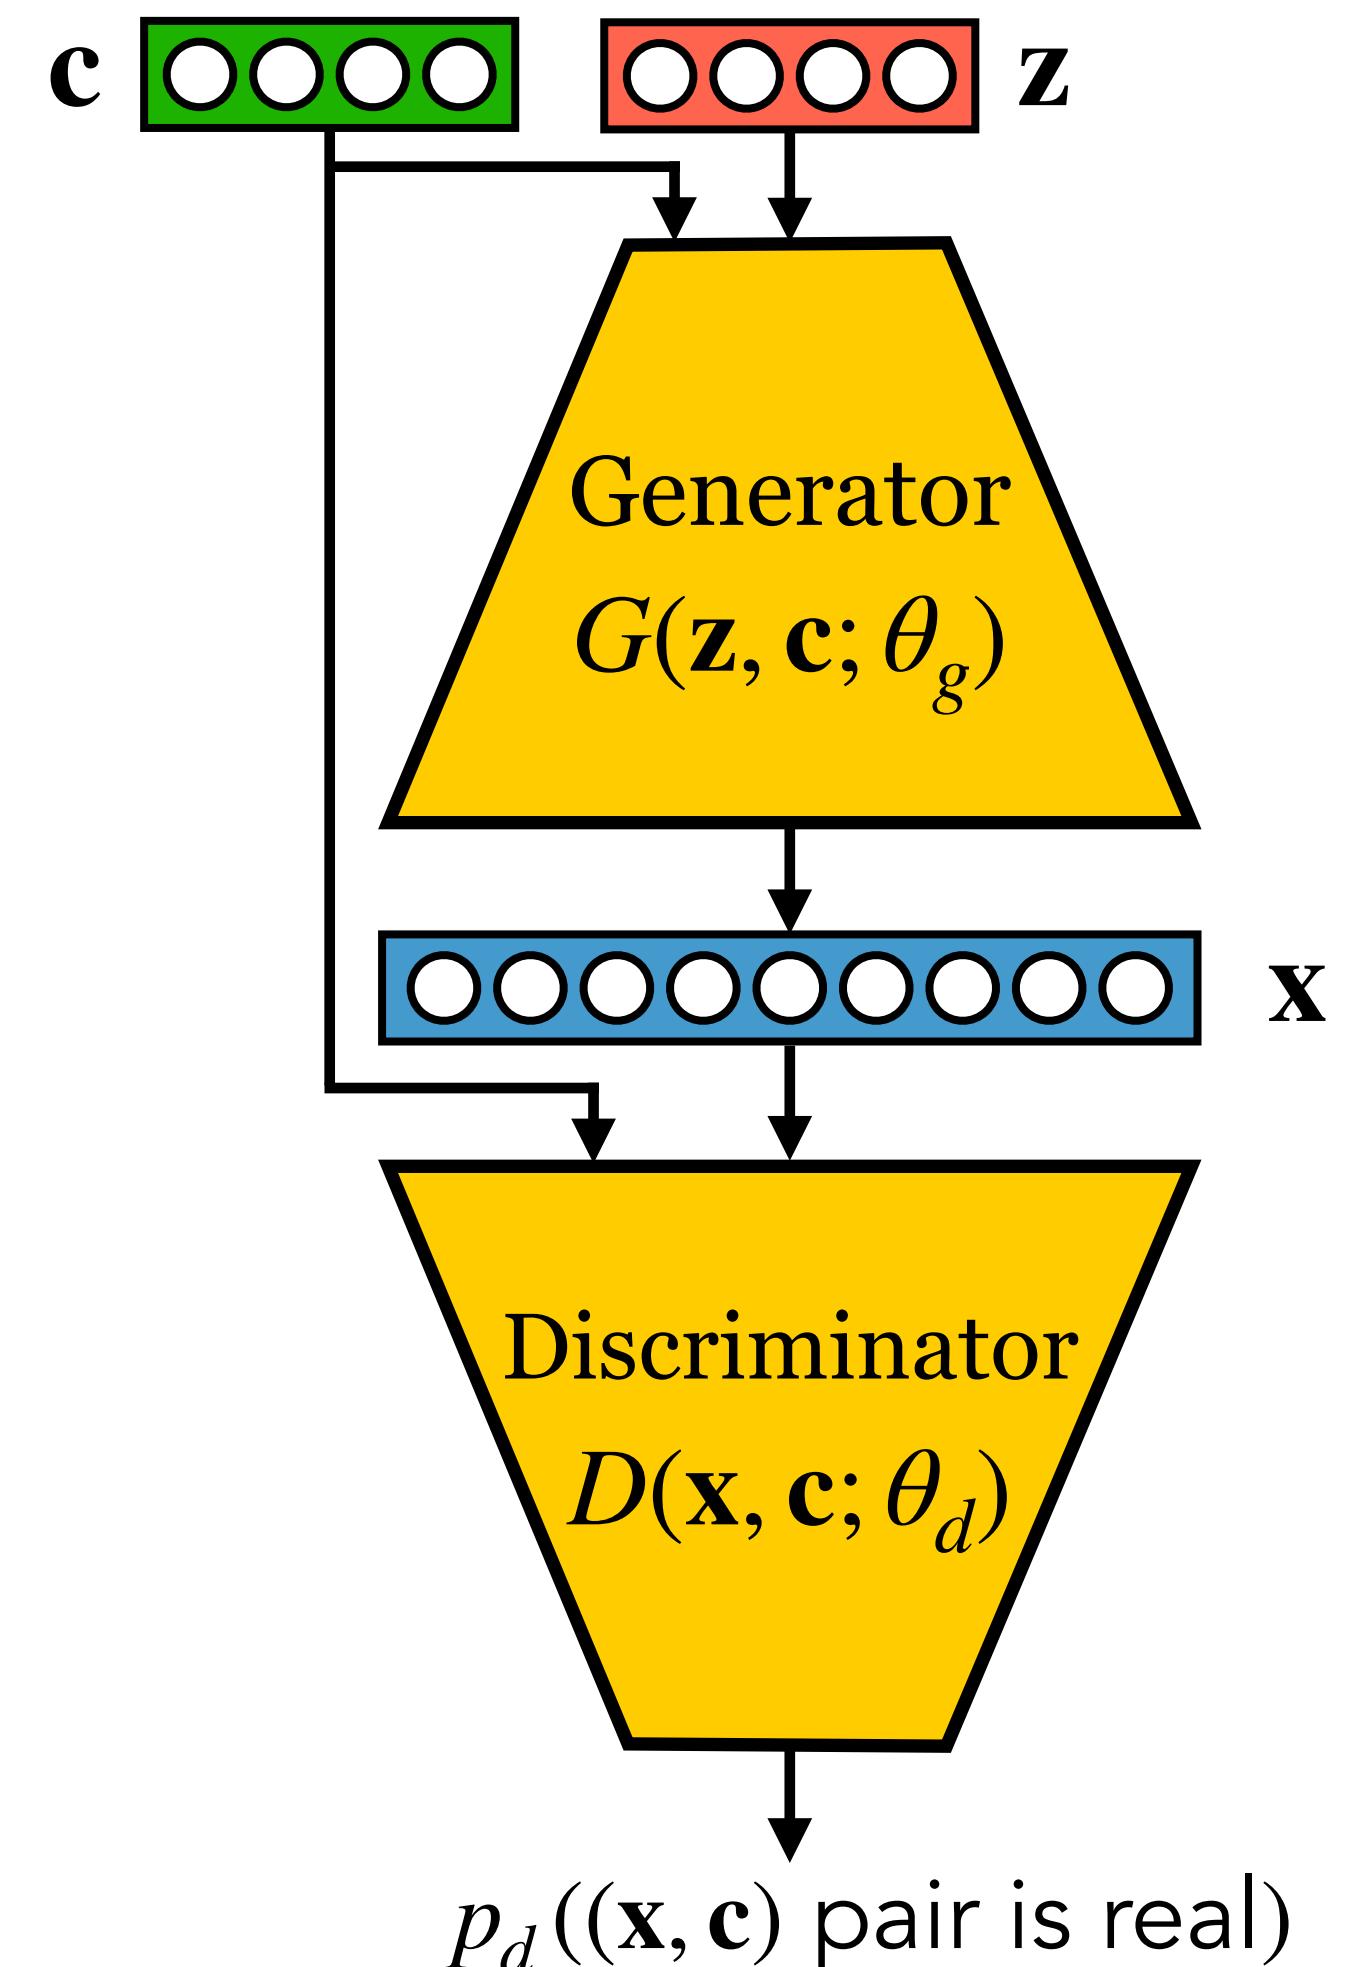
The discriminator also gets this information.

An example:

- A (standard) GAN trained on CIFAR-10 can generate samples from all categories, but we can not specify the category we want.
- **Conditional GAN:** We can generate cat images when given the condition “cat”.

The condition input can be:

- the class/label of the image to be generated,
- an attribute vector for an image (e.g., male or female face)
- a different image, ...



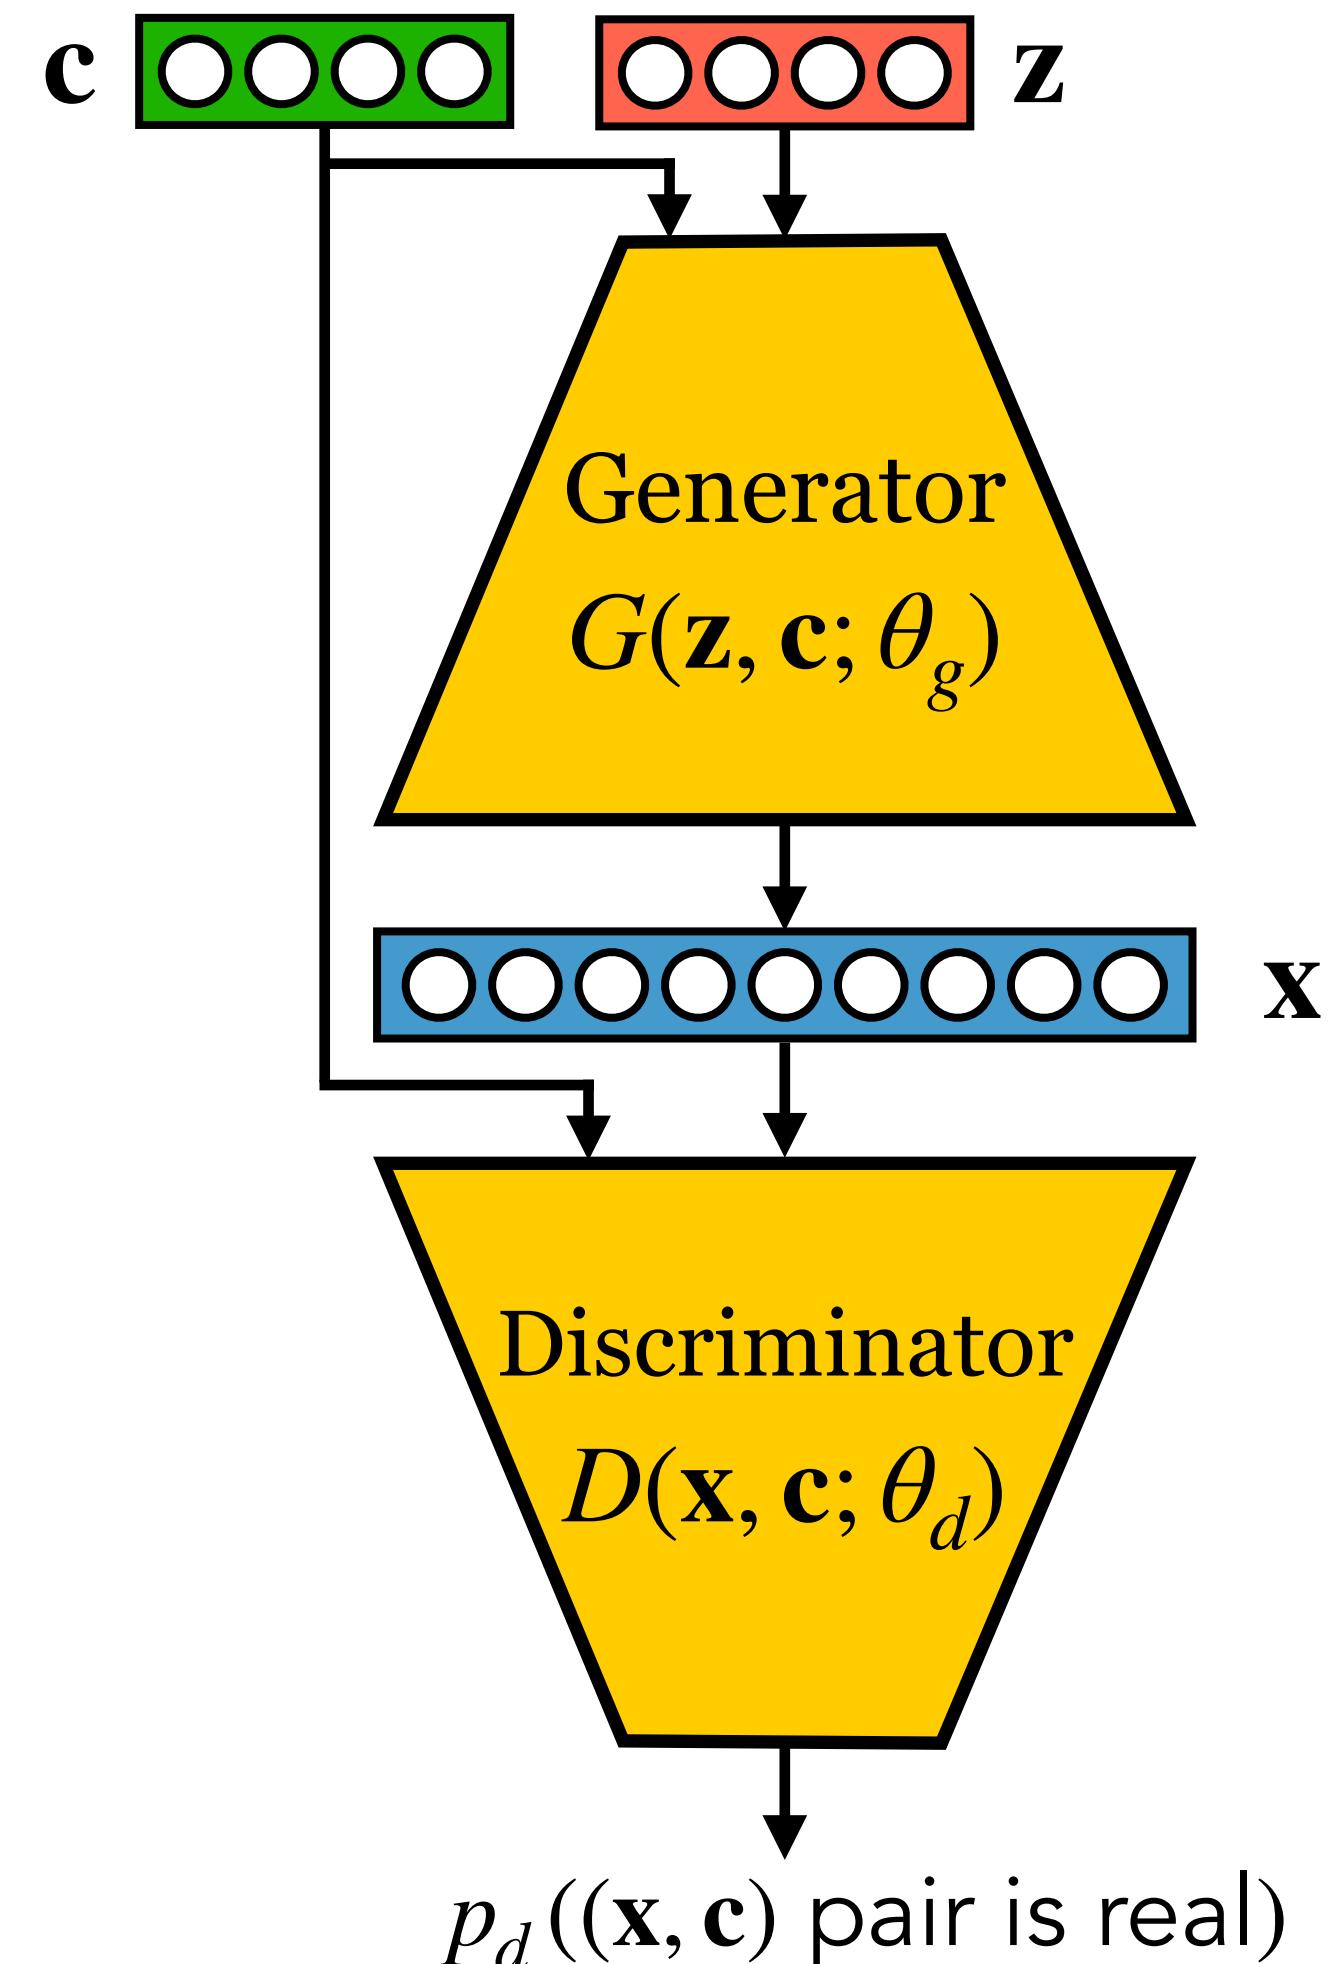
# Conditional GANs

One can use an additional (condition) information as an extra input to the generator.

The discriminator also gets this information.

This allows for many different applications of GANs:

- text-to-image translation
- image-to-image translation
- style transfer, photo colorization
- transforming photos from summer to winter, or day to night, ...



# Wasserstein GANs

---

**Original GANs can have problems with convergence, stability of learning, or mode collapse.**

- In the original GAN, the discriminator essentially minimizes the KL-divergence between the posterior distribution and the target distribution.
- Wasserstein GANs use the **Wasserstein Distance**.
  - This simply amounts to a different loss with better properties.
  - Discriminator provides a better and more stable learning signal for the generator.
- Here, one has to assure that the weights don't become too large to avoid instabilities.

# Wasserstein GANs

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [D(\mathbf{x})] - \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [D(G(\mathbf{z}))]$$

---

Original GANs can have problems with convergence, stability of learning, or mode collapse.

---

**Algorithm 1** WGAN, our proposed algorithm. All experiments in the paper used the default values  $\alpha = 0.00005$ ,  $c = 0.01$ ,  $m = 64$ ,  $n_{\text{critic}} = 5$ .

---

**Require:** :  $\alpha$ , the learning rate.  $c$ , the clipping parameter.  $m$ , the batch size.  
 $n_{\text{critic}}$ , the number of iterations of the critic per generator iteration.

**Require:** :  $w_0$ , initial critic parameters.  $\theta_0$ , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w [\frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSProp}(\theta, g_\theta)$ 
12: end while
```

---

# Example Applications: Image Generation

GANs can produce photorealistic images.

Progress on generating “faces”:



2014



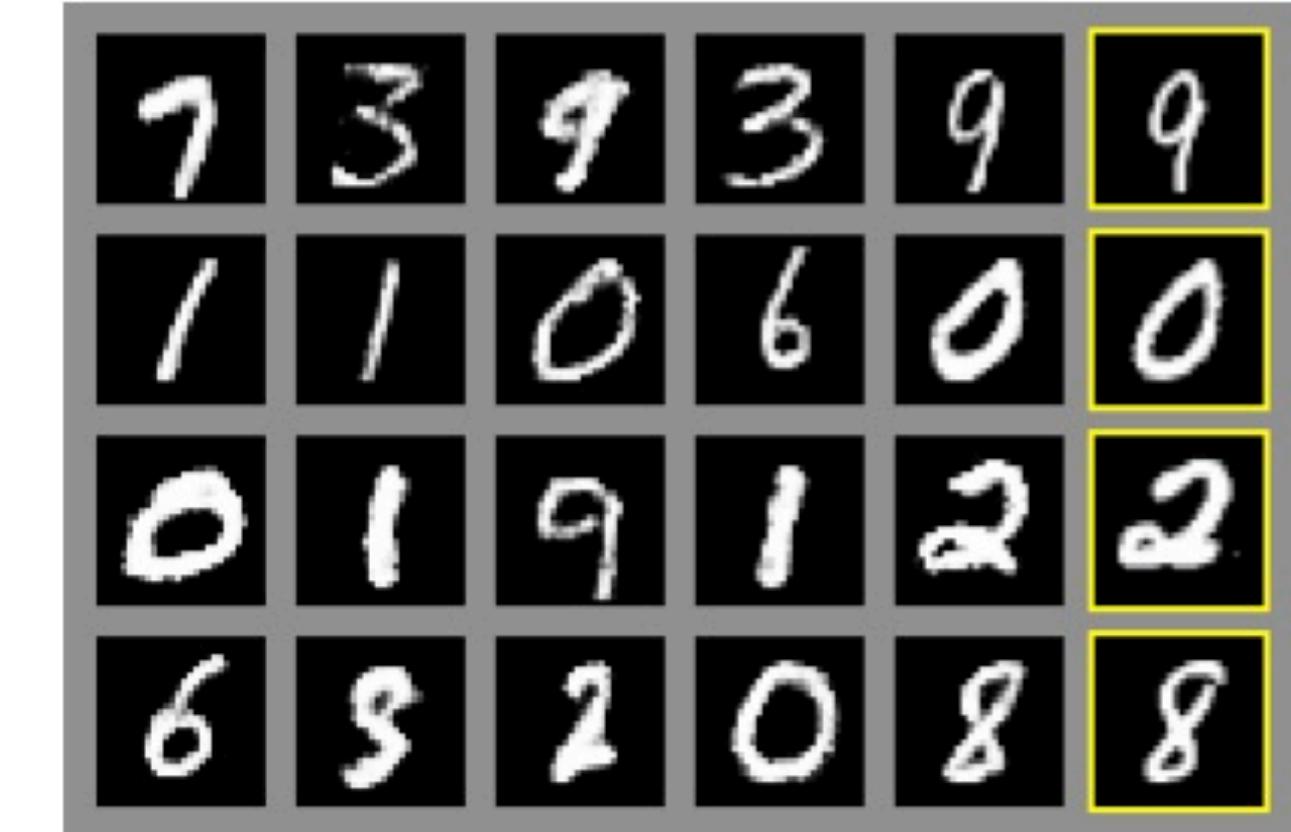
2015



2016



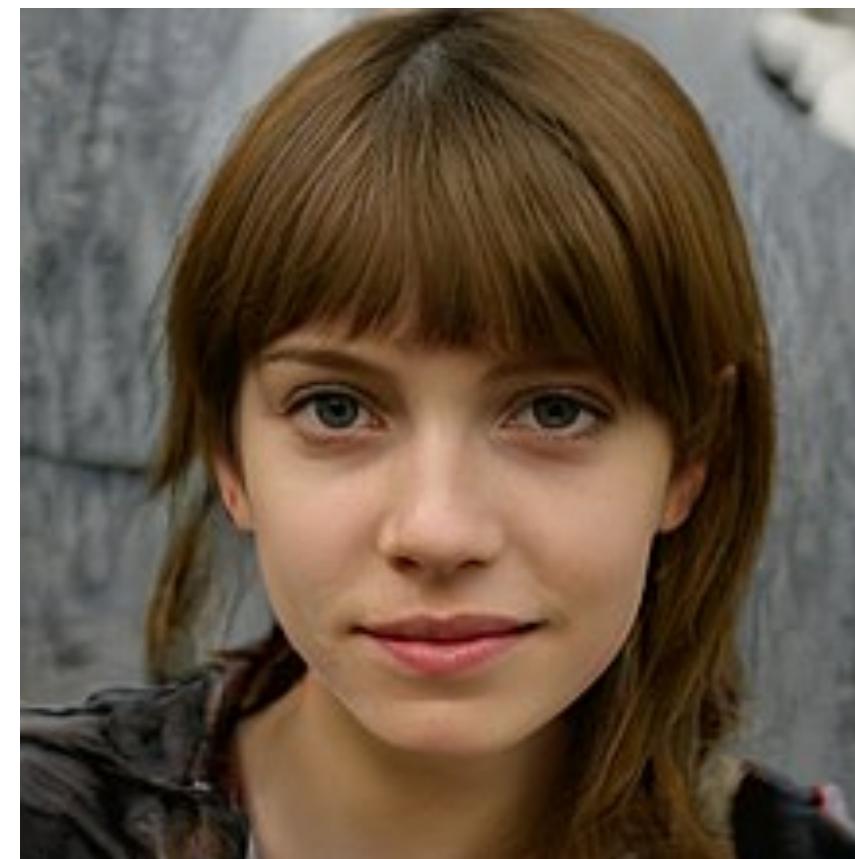
2017



[Goodfellow et al., “Generative Adversarial Nets”, NIPS 2014.]



An example generated by StyleGAN based on portraits:

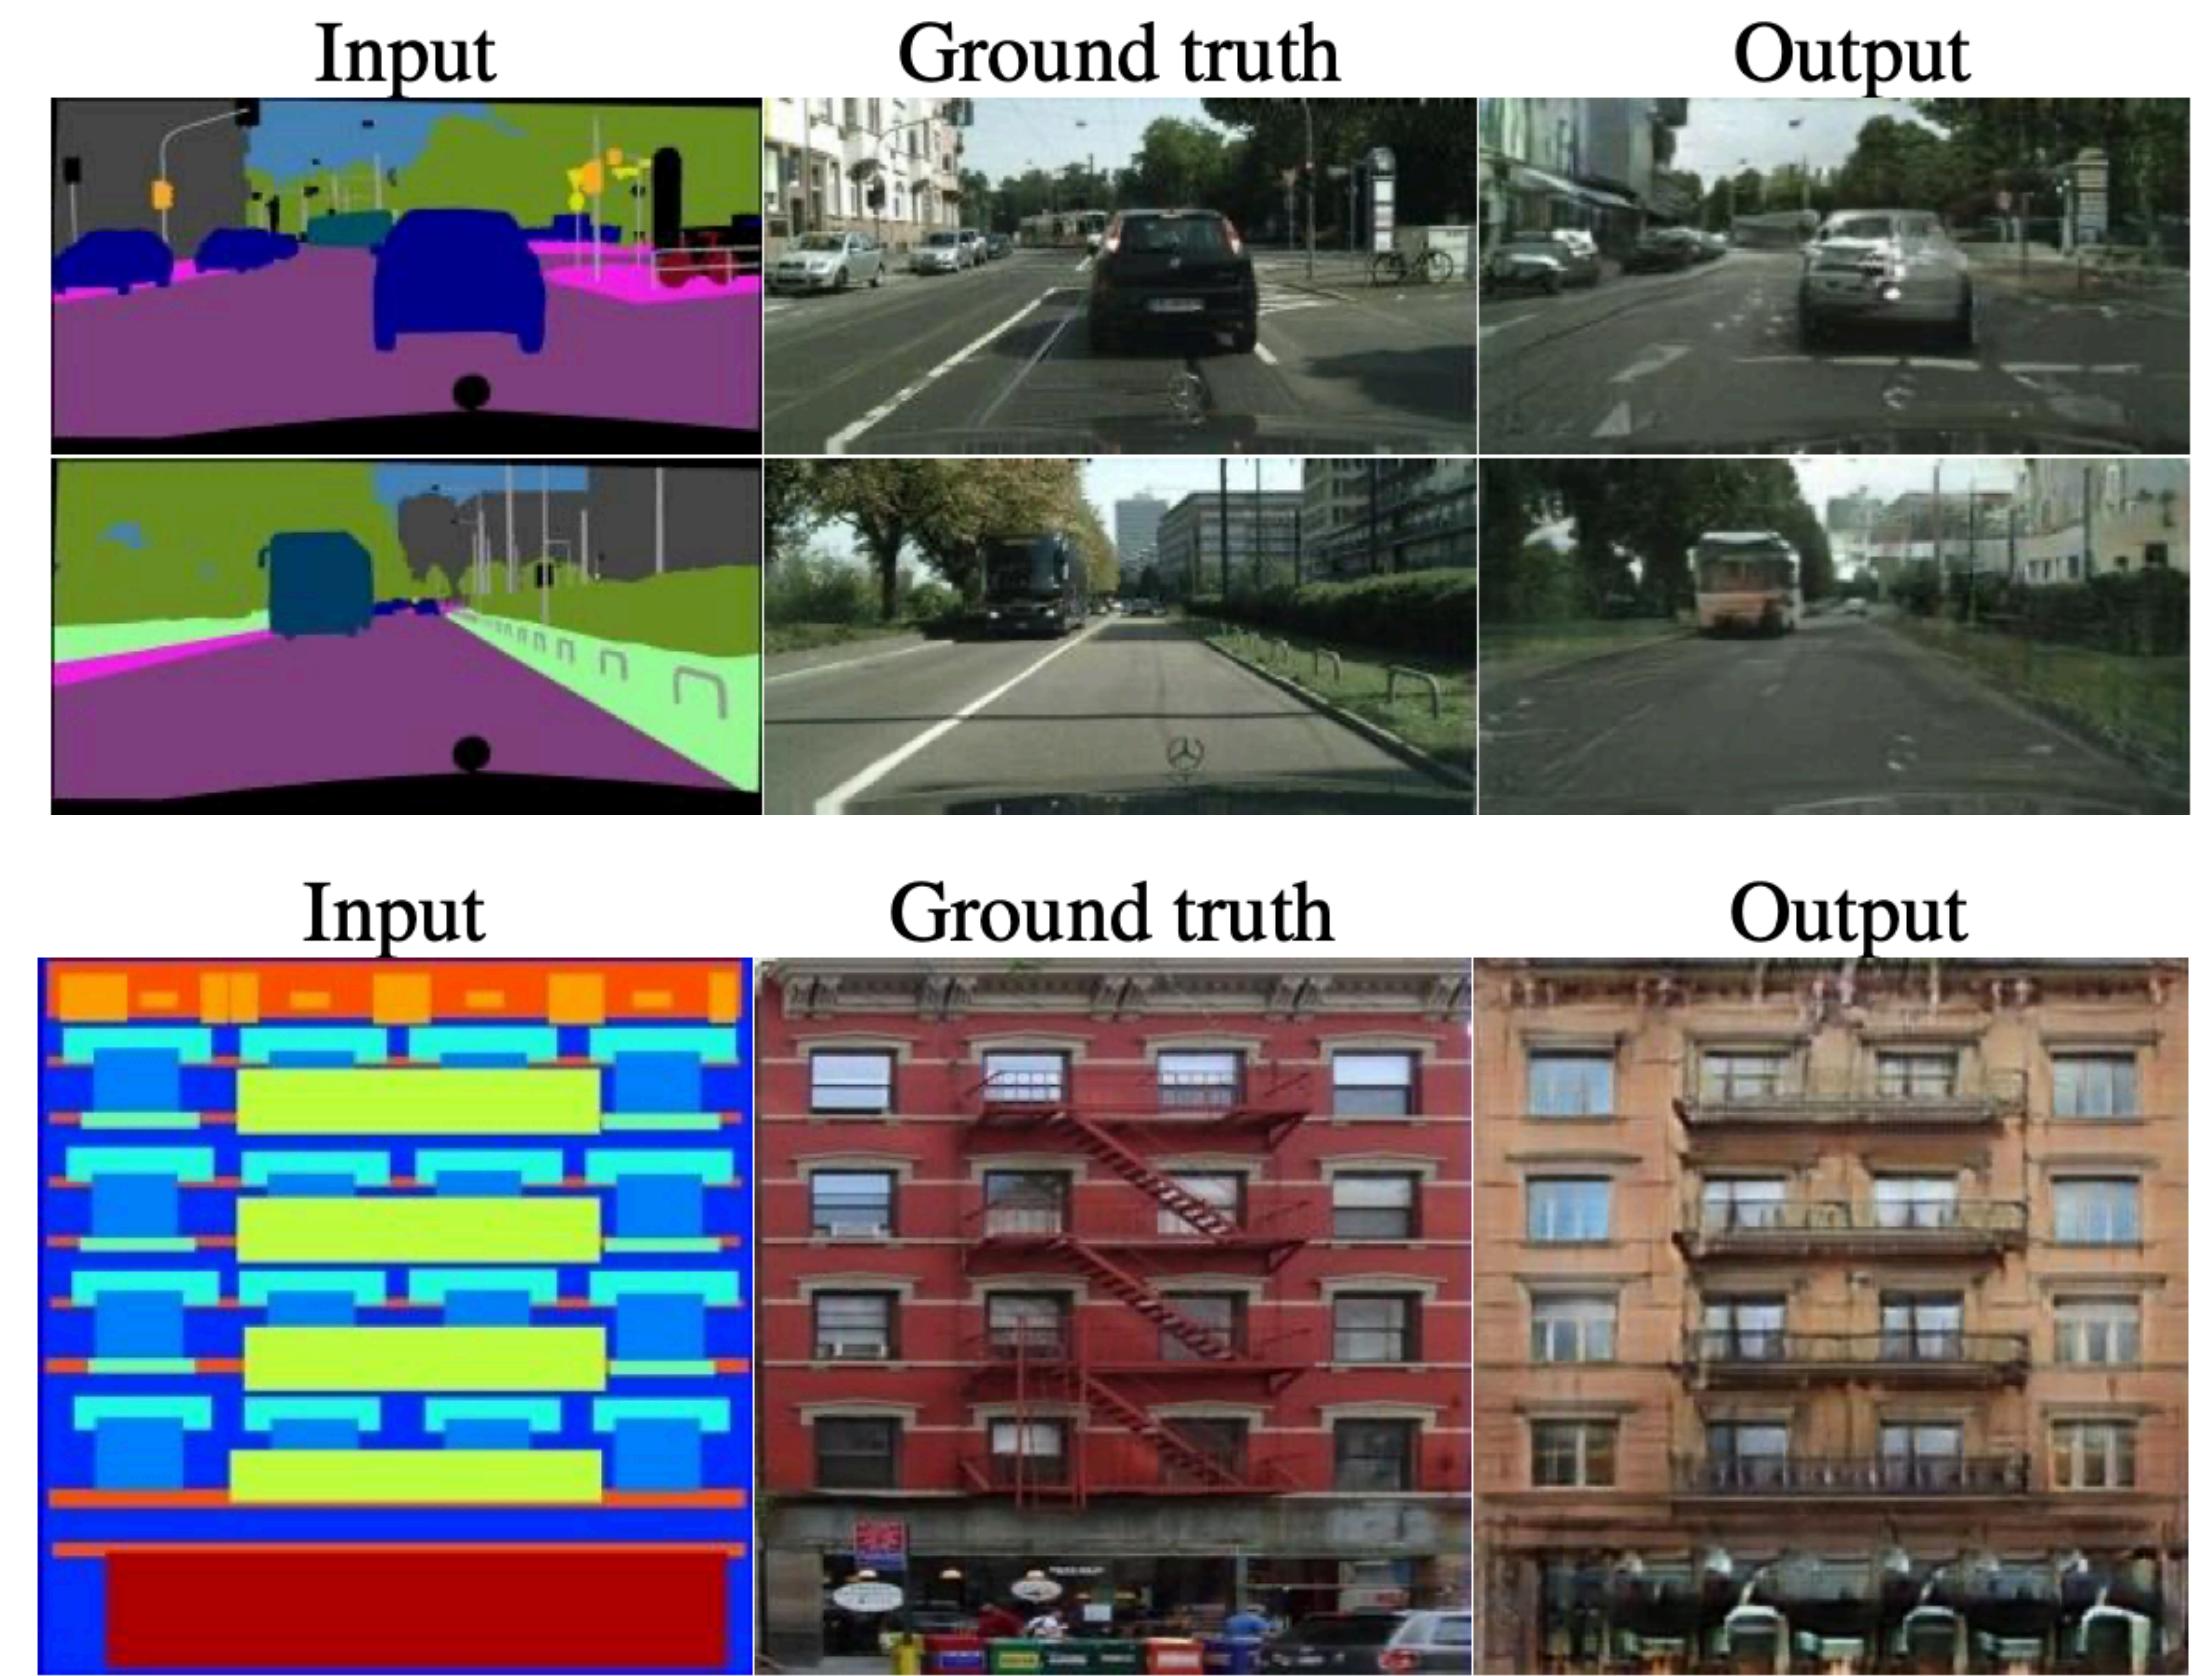


StyleGAN [Karras et al., CVPR 2019]



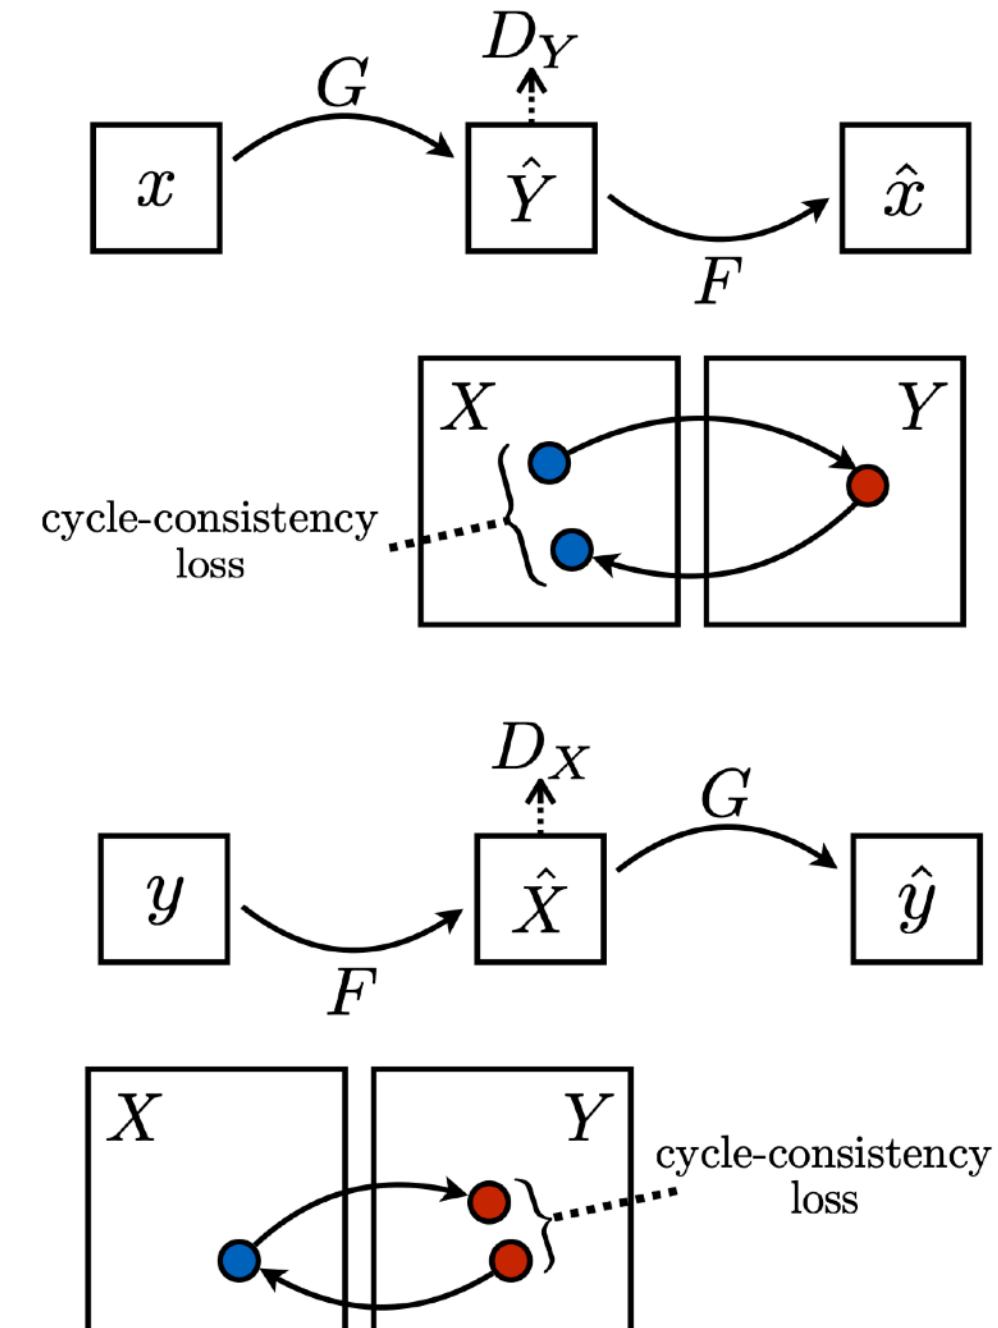
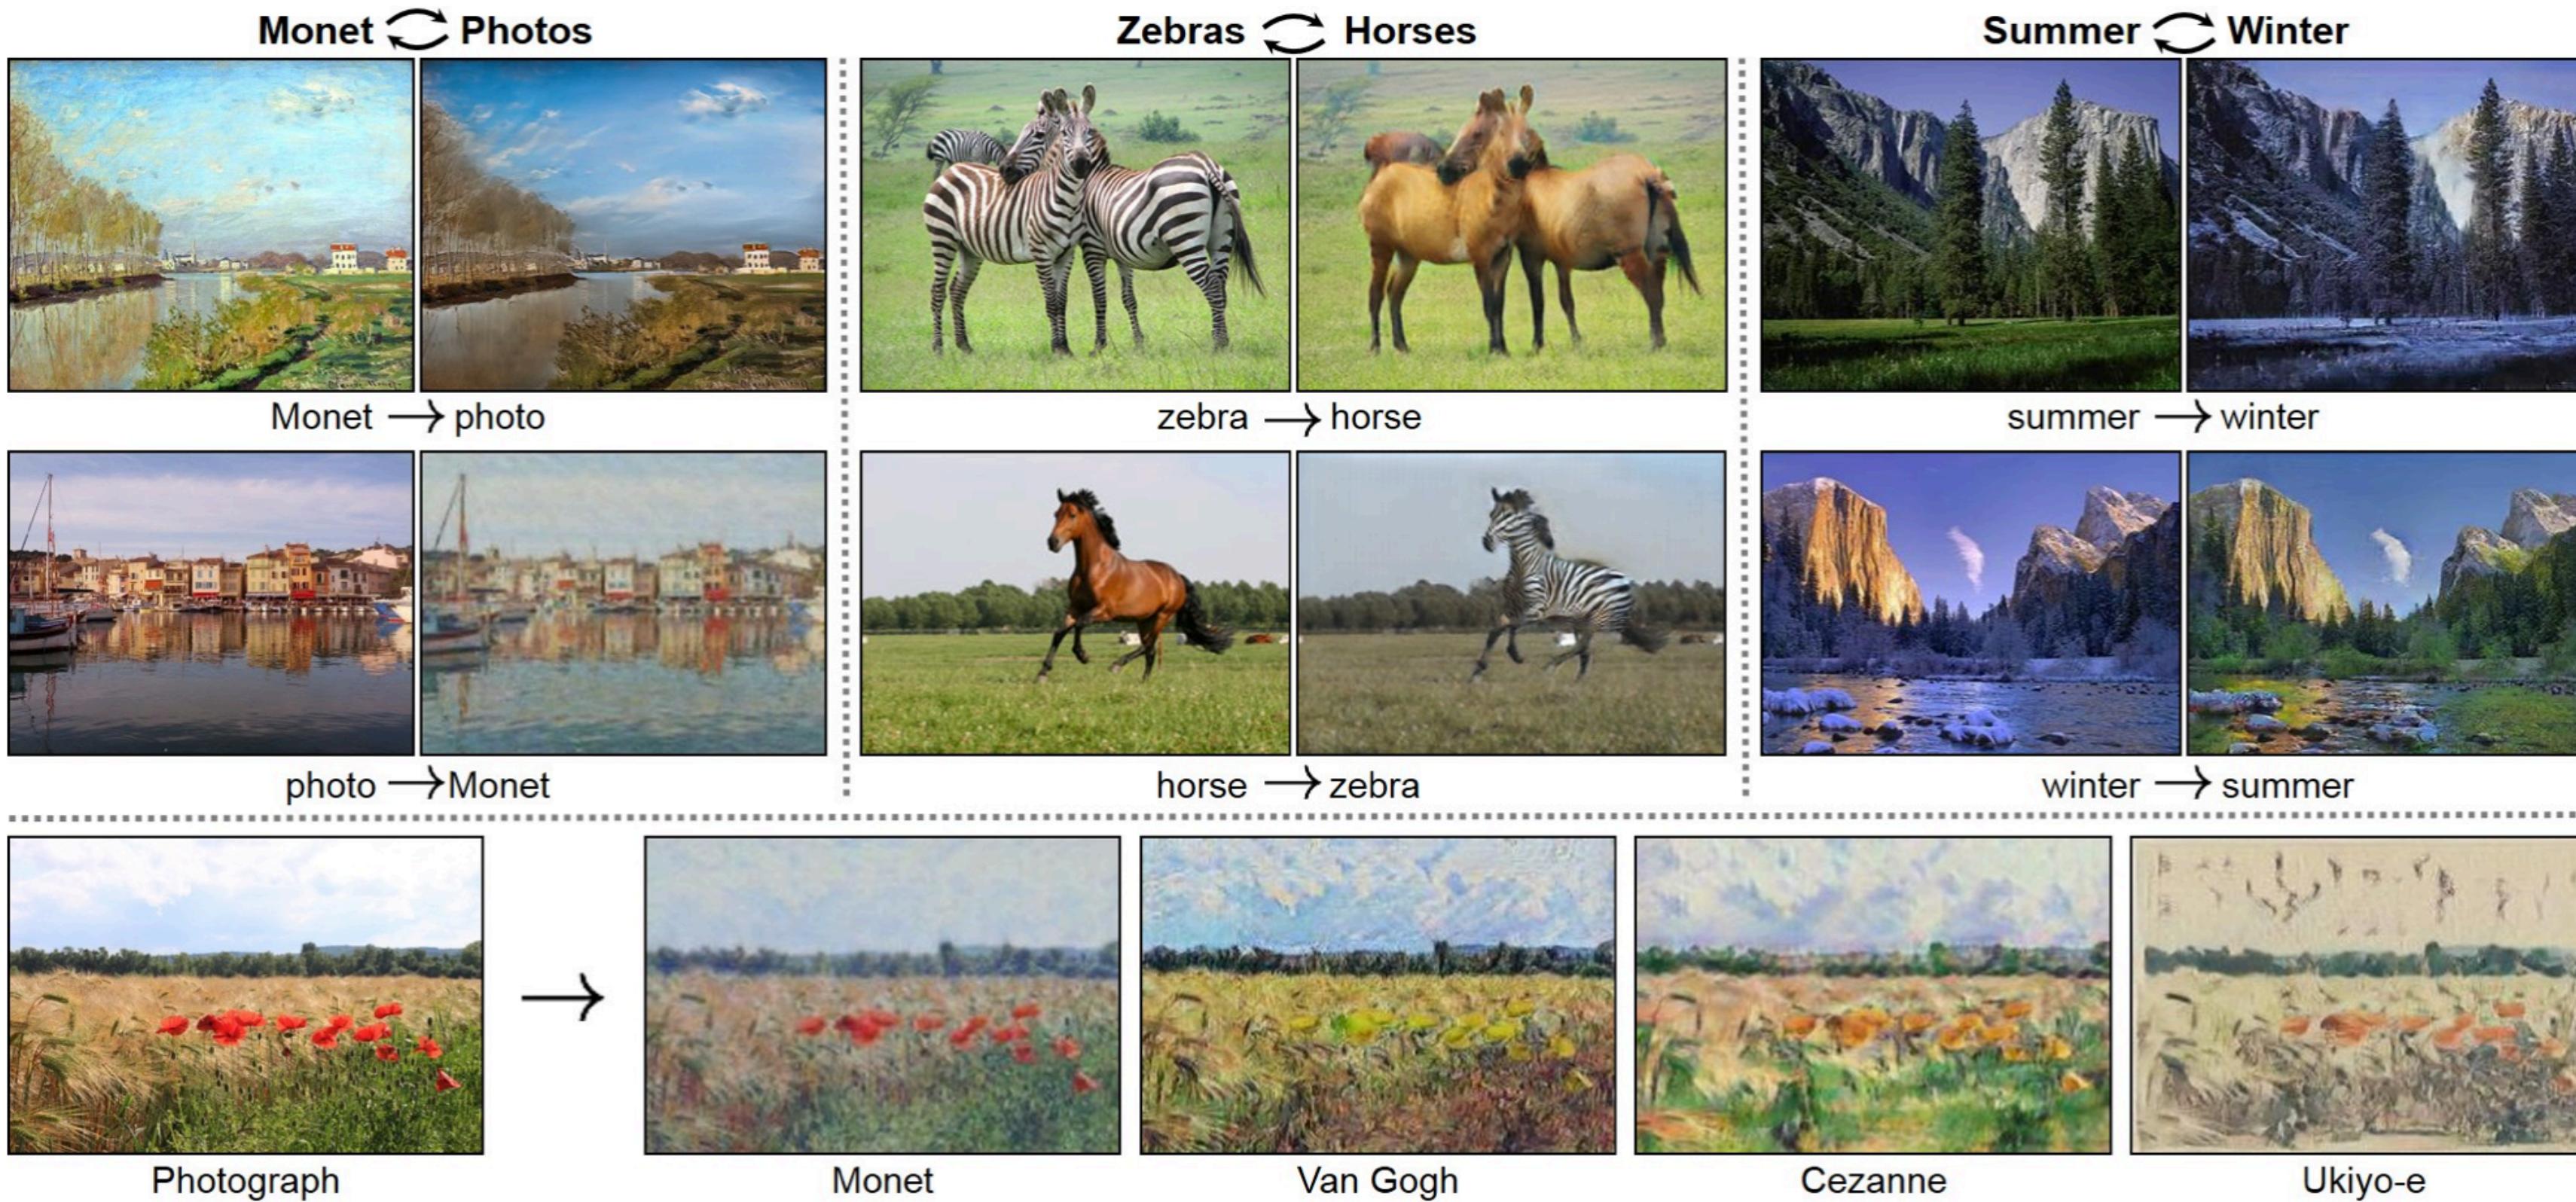
# Example Applications: Image-to-Image Translation

**Pix2Pix: Condition on one “style” of image, generate other style.**



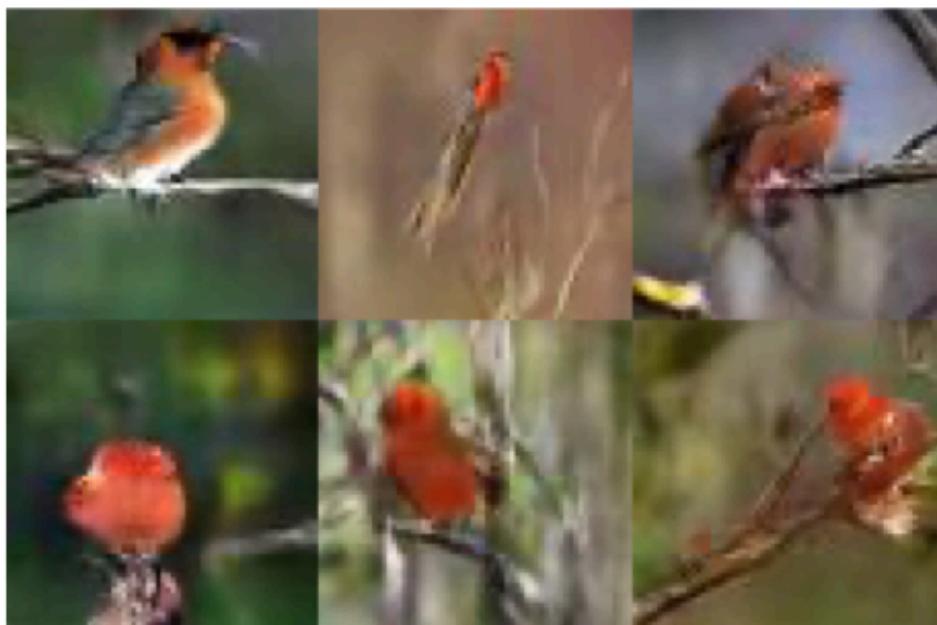
# Example Applications: Image-to-Image Translation

CycleGAN: Learns a forward and backward mapping between two domains, with an additional cycle-consistency loss to regularize this mapping.

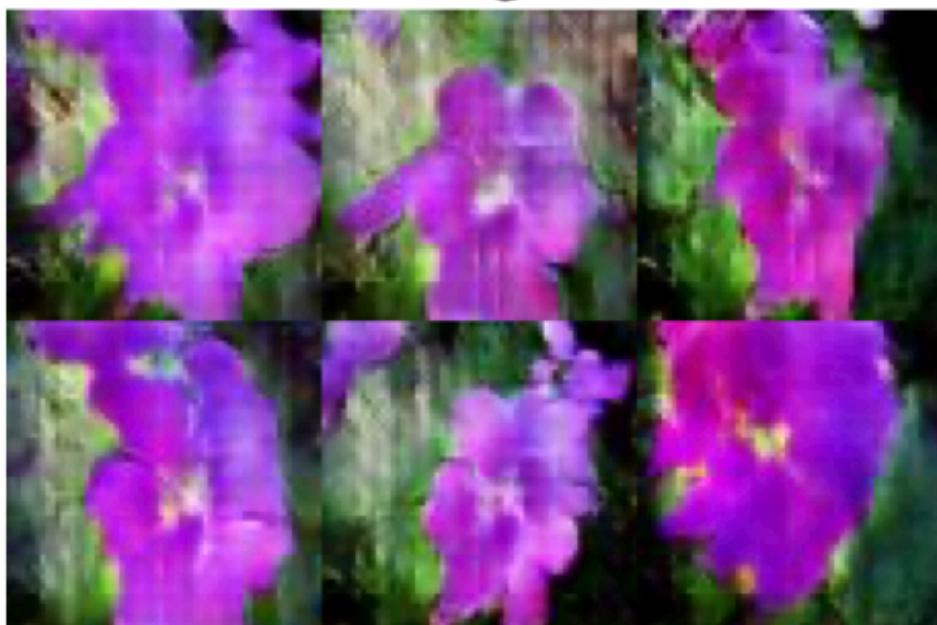


# Example Applications: Text-to-Image Synthesis

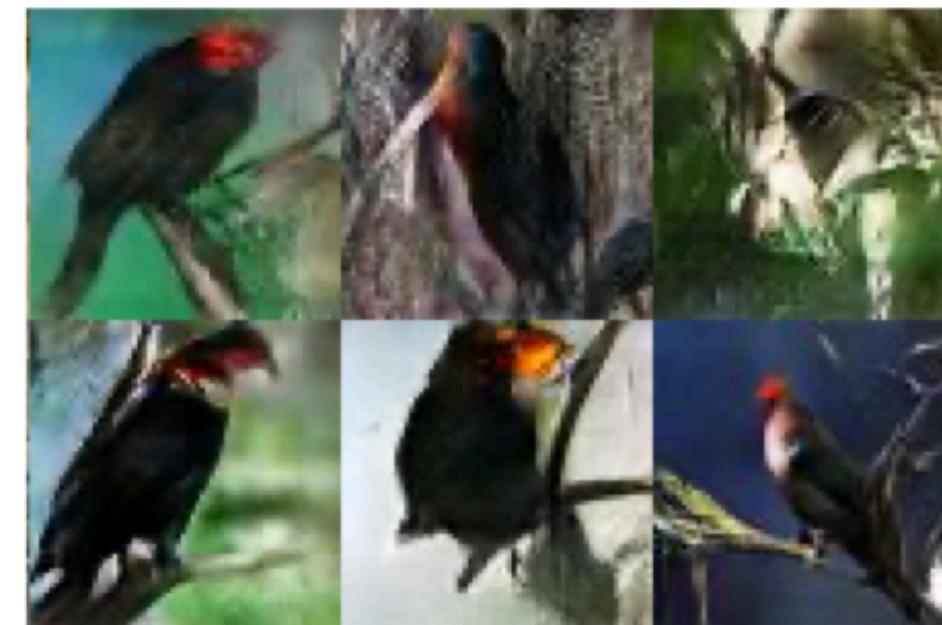
this small bird has a pink breast and crown, and black primaries and secondaries.



the flower has petals that are bright pinkish purple with white stigma



this magnificent fellow is almost all black with a red crest, and white cheek patch.



this white and yellow flower have thin white petals and a round yellow stamen



**Text descriptions (content)**

**Images (style)**

The bird has a **yellow breast** with **grey** features and a small beak.



This is a large **white** bird with **black** wings and a **red** head.



A small bird with a **black head and wings** and features grey wings.

This bird has a **white breast**, brown and white coloring on its head and wings, and a thin pointy beak.



A small bird with **white base** and **black stripes** throughout its belly, head, and feathers.

A small sized bird that has a cream belly and a short pointed bill.

This bird is **completely red**.



# Summary

---

- ▶ GANs are powerful generative models.
- ▶ They can be trained purely by gradient descent, no Monte Carlo sampling necessary.
  - ▶ Generator aims to fool the discriminator, and discriminator wants to maximize its own performance.
  - ▶ Relatively harder to train. Needs a lot of tuning to balance the generator and discriminator.
- ▶ Several improvements over the original model available.
  - ▶ Conditional GANs, DCGANs, Wasserstein GANs, ...
  - ▶ Various applications: image-to-image translation, text-to-image synthesis, ...
- ▶ Over the last decade, GANs became very popular with state-of-the-art results in image generation.
- ▶ Today, diffusion models have become the dominating model in image generation.

[J. Ho et al. "Denoising Diffusion Probabilistic Models." NeurIPS 2020.]

[P. Dhariwal & A. Nichol. "Diffusion Models Beat GANs on Image Synthesis." NeurIPS 2021.]

# Today

---

Deep Generative Models

Variational Autoencoders

Generative Adversarial Networks

Architecture & Training

Mode Collapse

Further Extensions & Variants

# Questions?